

Technicolor DDS User Manual

Contents

Overview.....	4
Middleware concepts and DDS.....	5
Need for a middleware.....	5
Publish-subscribe methodology.....	5
DDS in a nutshell.....	6
DDS Architecture.....	7
Technicolor DDS features.....	10
Overview.....	10
Design goals.....	11
Centralized Discovery Daemon.....	12
DDS Debug features.....	12
Data Types and Topics.....	13
Introduction.....	13
Technicolor DDS Type Support.....	14
Example Type usage.....	18
Using Sequences.....	23
Using Dynamic types.....	29
WaitSets and Listeners.....	36
Listener functions.....	36
WaitSets.....	37
Technicolor DDS memory handling.....	39
Memory pools.....	39
Configuring memory pools.....	40
Mobile device support.....	43
Controlling DDS activities.....	43
The Background Notification Service.....	44
Configuration variables.....	46
Configuration mechanisms.....	46
Configuration groups.....	47
Global configuration parameters.....	47
Pool configuration parameters.....	48
RTPS configuration parameters.....	49
IPv4 and IPv6 configuration parameters.....	50
IPv4 Multicast configuration parameters.....	52
IPv4 interface configuration parameters.....	53
IPv4 and IPv6 address scopes.....	55
IPv4 and IPv6 mode of operation parameters.....	55
UDP configuration parameters.....	56
TCP configuration parameters.....	57
Battery saving configuration parameters.....	58
Technicolor DDS components.....	59
Overview.....	59
Central Discovery Daemon.....	60
Shapes demo.....	63
Bandwidth program.....	66
Latency program.....	68
Chat program.....	69
Dump program.....	70
IDL to TSM converter.....	71

The DCPS functionality test program.....	72
The DCPS API test suite.....	73
Library build options.....	75
Using secure domains.....	75
DDS Debug shell.....	78
Introduction.....	78
Command Overview.....	78
Resource display commands.....	82
Timers.....	82
Strings.....	83
Pools.....	83
Sockets.....	87
Locators.....	90
Quality of Service.....	91
Domain data.....	92
Topics.....	96
Endpoints.....	99
History Cache.....	101
Proxy Endpoints.....	102
Security data.....	103
RTPS Send/Receive contexts.....	108
Memory dumps.....	109
Control commands.....	109
Tracing and Profiling commands.....	110

Overview

The purpose of this document is to make the user familiar with working with the Technicolor DDS Middleware implementation.

Technicolor DDS provides a Middleware approach that minimizes the memory footprint of DDS-applications, without sacrificing their speeds, in order to make the DDS Publish/Subscribe methodology usable on low-memory, low-performance embedded devices, all this both over LAN and WAN and fully secure.

The next chapters will give a small introduction in the Middleware concept and how the Technicolor DDS can make it easier for application writers to use DDS.

The Technicolor DDS Data Type system will be described in detail, since this is one of the most important changes in Technicolor DDS compared to other DDS implementations.

The document will also give some guidelines to application writers which kind of DDS data eventing approaches are available, and which is more appropriate in what situation, e.g. using either Listener functions, or using a WaitSet-based approach.

A number of environment variables are available in Technicolor DDS that can be used to influence the usage of Multicast and of Interface IP addresses. These environment variables will also be discussed.

A drastic per-node DDS memory reduction is possible using the Technicolor Central Discovery Daemon.

For a good look into how applications are actually using their DDS resources, the Technicolor DDS implementation can provide lots of ways to examine this in detail. Some of those ways are by default compiled out in order to minimize the code size, but it is quite easy to enable more exhaustive debugging features. The available compilation options will be explained in detail.

We will also explain how to optimise memory usage even more using the Technicolor DDS memory pool system.

A number of DDS example programs are provided which are quite useful on their own, such as the Shapes demo program, and the Bandwidth and Latency test programs. These will also be explained.

Middleware concepts and DDS

Need for a middleware

Software applications are becoming increasingly distributed nowadays. Whereas previously standalone application programs were sufficient for most purposes, this is no longer the case.

However, making a distributed application, even though it allows a much more powerful systems, has an associated cost (nothing ever comes for free).

Not only does the application writer need to take into account the difficulties of communication over LANs and WANs (with potential NAT issues), but communication needs to be secure as well, using modern cryptography mechanisms.

We can envisage a distributed application on several data distribution levels:

- Multiple threads within a process
- Multiple processes within a single node
- Multiple processes in distinct nodes.

For each of these approaches there are advantages and disadvantages.

The first level, i.e. multiple threads communicating on, typically, shared data is one of the easier approaches, when using a multi-threading library such as pthreads. There are, however a number of problems associated with this:

- A misbehaving thread can impact other well-behaving threads, making the resulting system a lot less robust than expected.
- It does not scale well, and leads to huge system components.
- Makes it very hard to pull apart the different threads into different processes in a later stage, since the possibility of data sharing is typically used in unexpected situations by the application writers.

For the second level of data distribution, e.g. multiple processes within a single node, there are also some mechanisms, out of the box available on many systems, such as Dbus. This mechanism typically works very well within a single node, but it doesn't scale when an application outgrows a single node, and when it becomes necessary to use multiple heterogenous nodes communicating together. Take the example of a Technicolor Gateway, a Technicolor Settop box and some management software running on a management PC. In this case, you automatically end up with multiple heterogenous nodes, all needing to communicate with each other.

The third level of distribution is the most general but leads to a lot more complex communication mechanisms. The communication overhead then becomes a heavy burden on application writers, since they need to be constantly aware of it while developing their software.

The idea of the middleware now, is to shield the details of the communication and the communication overhead from the application writers. Which data distribution level is used, then becomes a last-minute build decision, instead of having this visible on application level. Application writers simply use a standardized way of sharing their data with others, not knowing where the other application components are located, be it in the same process, another process on the same node, or on an entirely different node altogether. Not having to deal with the communication then allows application writers to focus much more on the usage of their data, which is a lot more important to them.

Publish-subscribe methodology

Various types of distributed middlewares exist, based on different approaches into how end-to-end communication occurs:

- **Client-server.** This works well when all data is centralized, such as in databases, transaction processing systems and in file servers. In this approach, a central server handles client requests and sends responses when ready. Typically this doesn't scale well, since the server is a central point of failure and quickly becomes a bottleneck in a system. Also, when multiple nodes are able to generate data, they all have to send this data to the server before it can be redistributed.
- **Message-Passing.** The message-passing approach extends the client-server system to a more distributed topology. It uses queues of messages to individual processes, which then process a single message at the time and responding to them if appropriate. Message passing is a lot more efficient than client-server, since it allows peer-to-peer communication. However, there are also a number of problems associated with this:
 1. The message passing approach is not data-centric, and requires the application to know the details as to 'where' the data is available, and then asking the appropriate data handling process for the data it is interested in.
 2. If providing processes become unavailable, applications need to know this and must recover properly, which is often quite complex to do.
- **Publish-subscribe** The Publish-subscribe approach makes communication a lot easier, since it adds a data-centric layer to the message passing approach. In this way, it is no longer required for the application to know where the data is located. It just subscribes to a specific data topic and will be notified when this data is available, is updated and is removed. Data publishers also don't need to listen to requests for data. They simply publish the data and the middleware will take care of proper distribution between the publishers and the subscribers.

The publish-subscribe approach is clearly the most powerful and most application-friendly approach to sharing data. Since it is fully peer-to-peer, it is also very fast and scales very well to a reasonable amount of distinct components (a hundred components seems feasible with modest memory requirements). Its design also makes it a lot easier to recover on various error situations, such as:

- Components disappearing due to crashes, and restarting.
- Components starting up in random order.

A number of publish-subscribe middlewares exist currently. The most widely known are Java Messaging Service (JMS) and Data Distribution System (DDS). The main differences between the two are:

- JMS is a lot slower than DDS, but allows both publish-subscribe and message-passing.
- DDS was designed from the start for real-time systems, and is much better suited for this.
- DDS has no builtin message-passing mechanism, although it is not that difficult to do this on top. In fact, the Remote Message Invocation (RMI) extension to DDS is a recent development, which is doing exactly this.

DDS in a nutshell

So what exactly does DDS deliver to applications?

The DDS publish-subscribe model connects information producers (publishers) with information consumers (subscribers).

The overall distributed application is typically composed of processes called "participants," each running in a separate address space, possibly on different computers and in different execution environments. A DDS participant may simultaneously publish and subscribe to a number of typed data 'topics'.

Type safety is an important aspect in DDS and the data interface to and from the applications, and even between different participants is therefore strongly typed.

Topic data types are described either statically (at compile time) or dynamically while the application program is executing (X-Types).

The first is achieved using a data type definition language that is programming language independent. This can be either XML-based or a language that closely resembles the familiar C/C++ data definition notations (IDL file). Later on, an off-line tool allows generation of typecode data and type-specific functions appropriate for a specific destination execution environment.

This has a number of benefits:

1. Since DDS knows about the data types, it can not only take care of the differences between little- and big-endian machines, but it can overcome the different data representations between multiple execution environments, such as C, C++, Java, etc.
2. Static typed data interfaces are generated, functions as well as data structure definitions, making the DDS API for applications less error-prone.
3. Data types can be "keyed" by marking specific fields in the type as "key" fields, so that multiple instances of data, e.g. multiple data streams can be possible, identified by its "key".
4. All data types are communicated over the wire and verified to ensure that types in both reader and writer are compatible.

Dynamic types are more difficult to use, since every type detail needs to be fully described by the application. This is mainly used to dynamically create DDS types from type descriptions done in higher order languages (ex. Java introspection), or from foreign data definitions (ex. JSON).

Due to the powerful introspection possibilities of Dynamic types, they can even be used to discover and use as yet unknown types that are discovered in the DDS domain.

Both static types and dynamic types can work seamlessly together. It doesn't matter how types are defined. One application using a dynamically created type can communicate with another application using the static type definition.

The DDS standard defines a communications relationship between the various publishers and subscribers. The data communications are decoupled in space (nodes can be anywhere), time (delivery may be immediately after publication or later), and flow (delivery may be reliably or best-effort and made at controlled bandwidths).

To increase scalability, topics may contain multiple independent data channels identified by "keys". This allows nodes to subscribe to many, possibly thousands, of similar data streams with a single subscription. When the data arrives, the middleware can sort it by the key and deliver it for efficient processing.

There is no requirement for a one-to-one relationship between publishers and subscribers. On the contrary, an any-to-any relationship is the overall model.

The DDS standard is fundamentally designed to work over both unreliable transports, such as UDP or wireless networks as well as over other transport mechanisms such as plain or encrypted TCP. Efficient, direct, peer-to-peer communications, or multicasting, will be used to optimize data transfers, reliable (even over multicast) or best-effort.

The communication aspects are typically specified by the application when subscribing to data topics. This is done with optional so-called Quality of Service parameters.

Filtering of received data samples is possible, either on data contents (using SQL-like expressions) or by specifying the requested rate of received data. Whether this filtering occurs on the publisher side or on the subscriber side depends on the DDS implementation. Good DDS implementations are able to do both.

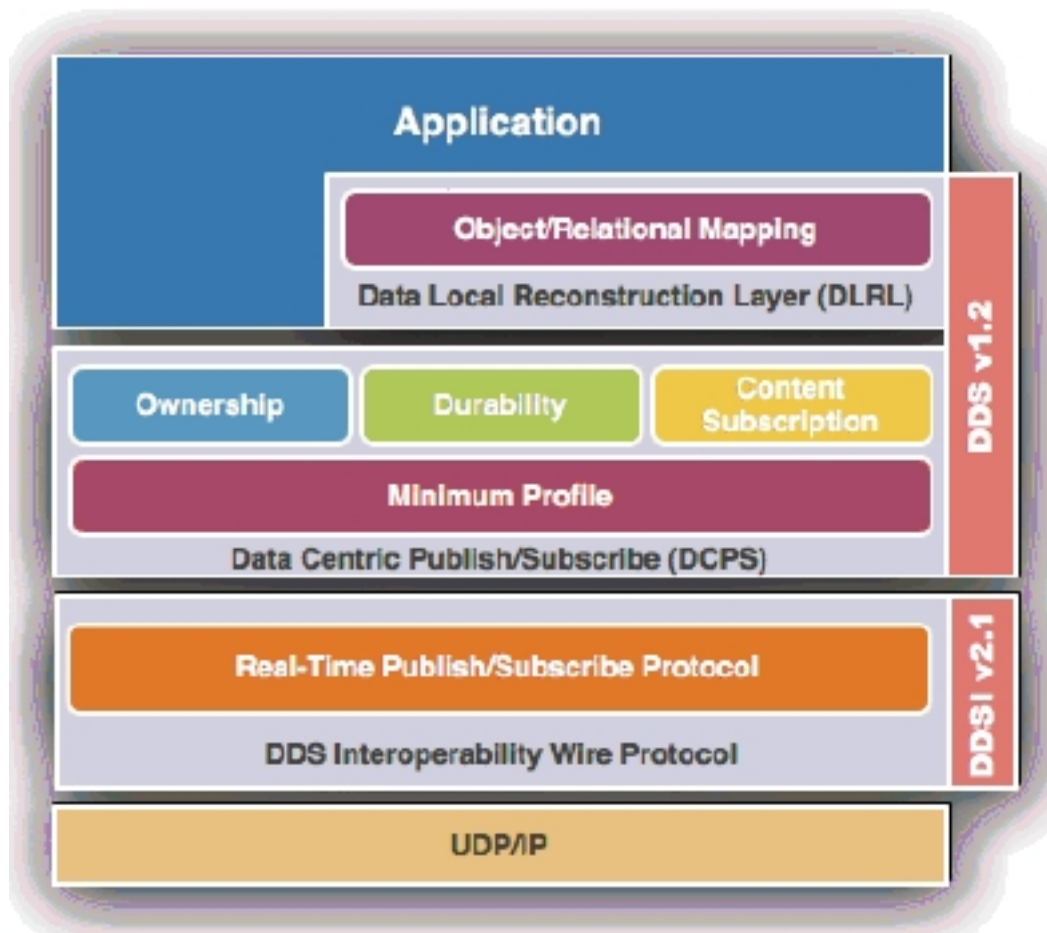
DDS data security is very important and can be done either on individual transport connections (DTLS/UDP or TLS/TCP) or very fine-grained using the DDS Security standard which provides the following additional features:

- Data can be plain, encrypted and signed (AES128/SHA1 or AES256/SHA256), or signed (SHA1 or SHA256).
- Encryption and signing can be done either on a full RTPS message (disallowing DDS routers/forwarders), on individual Reader/Writer Submessages, or only the Payload data is encrypted/signed.
- All this is selectable per Topic and encryption keys are different for different topics.

Standard DDS security requires Certificate-based authentication of participants, and uses signed permissions data.

DDS Architecture

As can be seen in the figure below, a DDS middleware is structured in a number of layers:



The DDS specification gives a choice of 2 distinct APIs to the user, although most DDS vendors implement only one of those APIs.

The most common API, which every DDS vendor currently supports, is called the Data Centric Publish Subscribe (DCPS) API, which is the most basic of the two.

The second API is the Data Local Reconstruction Layer (DLRL) API, an optional layer on top of DCPS, which is a lot more Object-Oriented. It allows applications to describe objects with methods as well as attributes, which are either local or shared. DLRL does not do RPCs however, all methods are local. Also, since only a few DDS vendors support DLRL, using this API would lead to interoperability issues with different DDS vendors.

For both APIs, the OMG standard gives a description in a language-independent format (IDL) in order to allow multiple execution environments and also to allow multiple programming languages. Due to this choice, the conversion from this format to an effective programming language API still leaves some design choices open. A disadvantage of this approach thus means that a specific programming language DCPS API may differ slightly from one DDS vendor to the next.

In order to avoid these minor differences, and also because for object-oriented languages it is possible to wrap the DCPS API into language-specific class-based APIs which are of course a lot easier to work with, OMG has standardized some specific programming language specifications for the DCPS API. Specifications are defined for both C++ and Java.

The DDS specification allows many different types of data durability, from volatile up to persistent durability. The first will remove data as fast as possible from the data caches. The persistent durability stores data automatically on a storage medium (such as a disk) without the users intervention.

For effective communication over a network, an additional software layer below DCPS is responsible. This layer, which is called the Real-Time Publish Subscribe (RTPS) layer is able to exchange data between DDS components in a standardized manner, providing necessary compatibility between different DDS vendors.

RTPS is responsible for coping with the differences between different nodes (big/little endianess) and between subtle different data layouts between different execution environments. This is achieved in practice by marshalling local data representations into a network-specific format, which is execution environment agnostic, before sending the data on the wire. On reception of messages, the data is unmarshalled to the local specific execution environment data representation.

Since RTPS is designed to operate over multiple networking protocols, there is no requirement for a specific networking protocol, and DDS thus allows both operation over local LANs as well as over Wide Area Networks (WAN), typically using different protocols.

Only the interface between RTPS and UDP/IP is defined in the DDS standards, but in practice, many vendors offer support for IPv6, secure transports, TCP, and even shared-memory transports within device nodes.

Technicolor DDS features

Overview

As described before, the publish-subscribe way of exchanging data between components is a very powerful idea.

Of course, an idea is one thing, but an efficient implementation is another thing altogether.

While examining DDS as a technology, it became apparent that current commercial DDS implementations had huge memory and performance requirements, making them unfit for use on embedded platforms. It also was not clear why exactly this seemed to be the case, since DDS does (from our point of view) allow very efficient and memory-constrained implementations.

The focus of DDS vendors at that time was (not yet) intended for small platforms, but was instead mainly driven by use cases like:

- High-speed financial transaction processing.
- Avionics industry
- Defense industry

Since most of these do not have tight memory requirements, although they do have high performance requirements, this was not seen as a big issue.

Also, the speed aspects of DDS were seen as prevalent to resource usage. Optimizing an implementation for speed typically sacrifices resources due to strategies such as heavy caching, and overly preallocating receive buffers.

The result of all this is that it became necessary for Technicolor to build their own DDS implementation in order for DDS to be optimally usable on the current Gateway and Settop box platforms.

The current version of Technicolor DDS runs on multiple Operating Systems, i.e. Unix, Linux and Windows and is able to run and is actively supported on x86 platforms in both 32-bit and 64-bit modes. It has also been tested on ARM-based Android platforms and Apple OSX and iOS devices, and is used on gateways, settop boxes and cable modems based on both big- and little-endian MIPS processors and Intel Atom processors.

Currently, Technicolor DDS is an integral part of the Qeo Component Development framework.

Except for the following, not often used features, the complete DDS specification is implemented:

- Because there was no real need for Persistent and Transient durability, and because it is a quite complex feature to implement, these durability modes are currently not present.
- No support for Multitopics. Multitopics are only implemented by a few DDS vendors, are quite complex, and there was no use case for them in Technicolor Qeo.
- No support for coherent topic updates. The Qeo API does not need this and a lot of DDS vendors don't implement this, so this was not a priority either.
- Because the DLRL layer requires an additional software layer on top of the DCPS API with additional memory and performance requirements which does not bring much added value, and because only a few DDS vendors support this, Technicolor DDS has chosen not to support this API.

The DCPS API for Technicolor DDS is available for the C, Java and Python programming languages.

Data sample sizes are supported up to 256 MBytes per sample. In Qeo, limits on the Android platform reduced this to 4MBytes per sample, which is enough to transfer complete pictures as single samples.

DDS Dynamic types (X-Types) are fully supported except for XML-based transport and reading of XML-types at runtime.

Technicolor DDS has support for RTPS over UDP/IPv4 by default. Support for RTPS over UDP/IPv6 is optional but it can easily be added at compile-time when the `-DDDS_IPV6` compile option is used. In the latter case, it can be controlled via either environment variables or other configuration mechanisms as to which IP variant (or both) is to be used. When both are used simultaneously, it is also selectable which protocol has precedence and which of the protocols will have discovery enabled.

In addition to UDP, Technicolor DDS supports the following transport mechanisms:

- DTLS/UDP (secure peer-to-peer transport)
- Plain TCP and TLS/TCP (secure remote transport)

The latter two mechanisms are primarily used for remote DDS access over NAT/Firewalls, where TCP connections are setup from within the NAT-ed LANs.

DDS Security is accomplished by means of:

- Transport security (DTLS/UDP and TLS/TCP), and/or
- Fine-grained security, based on the DDS Security specification, using either standard or Qeo-specific handshaking mechanisms.

In addition to simple peer-to-peer connectivity, Technicolor DDS supports a DDS forwarder (or relay) functionality that allows to extend a local DDS domain/cloud in order to include devices located remotely over Internet via TCP. This mechanism is invisible for application writers. DDS takes care of all the routing aspects.

Design goals

Due to Technicolor DDS being used on platforms having severe memory footprint restrictions, and also because it is an intrinsic part of the Technicolor core middleware, which comprises more than DDS alone, and instead provides a full component development framework, the design of Technicolor DDS is slightly different from other commercial DDS implementations.

Following design goals were envisaged:

- Memory usage needs to be kept as low as possible, both for core data structures where DDS metadata must be stored in the most optimal manner, as well as for real data buffers, storing topic data.
- Never create an entity until it is really needed to have it.
 1. RTPS Reader/Writers are only created when a remote entity is detected that has an interest in the topic data.
 2. Data caches can autonomously transfer data between a local Writer and a local Reader without the need for RTPS.
 3. Discovery data readers are only created when the user expresses an interest in it.
- When possible, reuse the same data (sometimes determined by hashes), using reference counts. This is done, for example, for:
 1. Stored user data, which is shared between DDS history caches.
 2. Locator structures.
 3. Discovery data.
 4. QoS parameter sets.
 5. Character strings are stored only once.
- Using memory pools instead of 'pure' `malloc()/free()` has several advantages, such as:
 1. Severely limiting allocation overhead and memory fragmentation.
 2. Allocation/free of data chunks is much faster.
 3. Allows to specify upfront what the pool limits may be.
 4. The application's memory usage can be monitored due to the pool's bookkeeping.
- Contrary to other DDS implementations, data type marshalling and unmarshalling is data-driven instead of function-driven (see TSM type specification).

- Data types are generated primarily from Qeo layer XML-based data. Qeo is an additional, simplified API on top of Technicolor DDS. This does not preclude an IDL data definition compiler, in fact, a small IDL compiler is included, which allows to create TSM definitions from IDL.
- Since the Qeo layer is a C-component, and there are already Qeo bindings for various execution environments (Java/JNI and Python), there would be no need for other execution environments on top of DDS. There are, however, DDS bindings available for other execution environments (Java/JNI and Python).
- Focus on easy debugging of application as well as Technicolor DDS operations.
- Use a node-centralized Discovery data manager, the Central Discovery Daemon (CDD), in order to prevent components of learning topics via the DDS Discovery mechanisms in which they are not interested.

Centralized Discovery Daemon

Ordinarily, DDS implementations learn all topics that are present in the DDS domains in which they are interested, even though they usually are only interested in a subset of the collective Topic data.

This is done, by design, because it is not known initially which topics will be referenced in the future by applications. It could be, for example, that Topics are discovered first, after which they might be referenced by local Readers or Writers. A strategy that simply throws away Discovered data for which there is no local Topic yet, will not work at all. The data that was thrown out will not come back automatically afterwards, just like that.

The solution to this problem is a node-central Discovery manager, called the CDD component. The Centralized Discovery Daemon is a stand-alone (daemonized) program that allows to optimize this per-node Discovery data so that discovered data does not need to be kept in components that are not interested in this data.

The CDD, when active, allows to communicate the state of Discovery data to other Technicolor DDS components. This way, those DDS components will be informed of matching discovery data the moment they subscribe to Topics and there are already 'early joiners' for this topic.

DDS Debug features

The Technicolor DDS implementation has the following features that are useful for debugging user- as well as DDS-specific code and data:

1. A debug command handler (shell) can easily be attached to application code.
2. An application debugger (gdb, ddd, etc.) can 'call' Technicolor DDS debug commands.
3. Any Technicolor DDS application can optionally be verified/monitored remotely using the telnet protocol to a built-in debug server giving access to the Debug shell. This debug server is either started automatically, or spawned on request of applications.
4. The Debug shell can give detailed information on memory usage, created as well as discovered DDS entities, and it allows detailed inspection of History Cache data.
5. RTPS message tracing can be enabled for any topic, even the builtin Discovery topics.

Data Types and Topics

Introduction

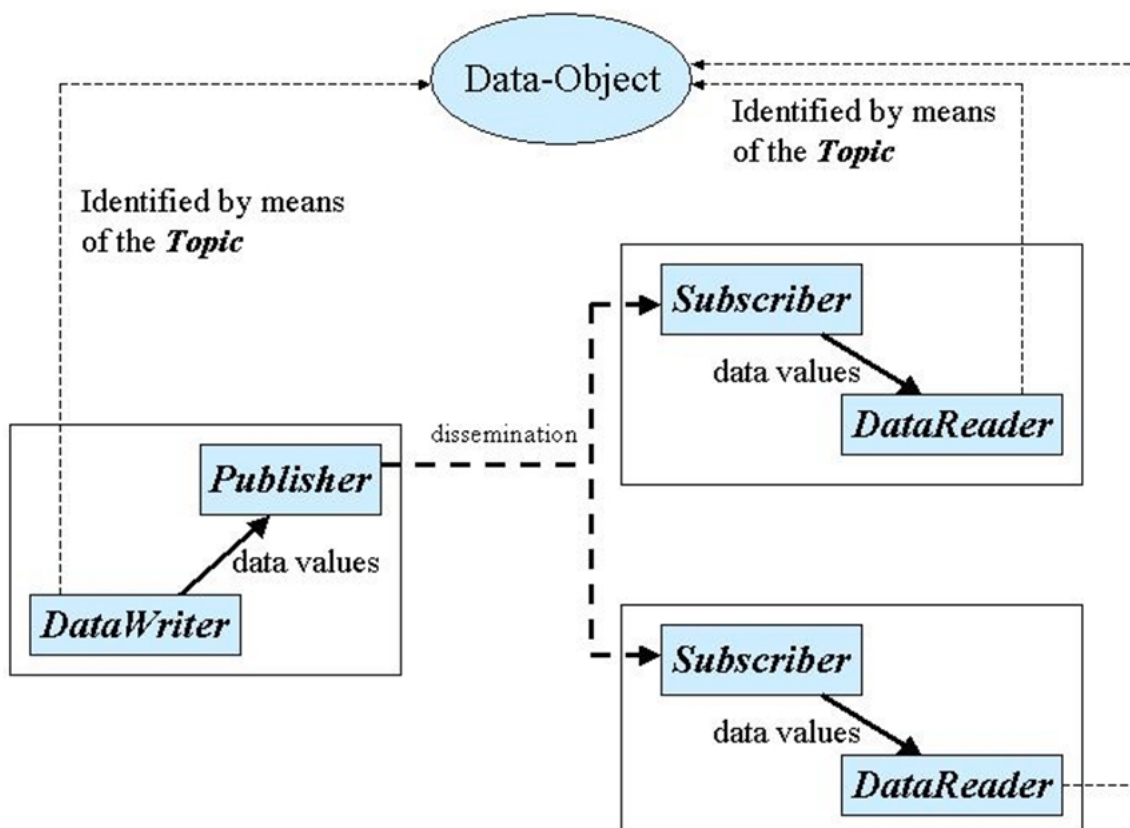
As explained before, data is transferred between Data Writers and Data Readers using properly typed writer and reader functions, based on an intimate knowledge of the exact data types that are used.

This is done using the concept of topics and types, where a DDS topic corresponds to a single data type. However, several topics may refer to the same data type.

Topics may range from a single instance to a whole collection of instances of the given type. Different instances must be distinguishable. This is achieved by means of the value of some data fields of the objects that form the key to that data set.

The exact type of a topic and the fields it contains can either be constructed dynamically (using the Dynamic Type system), or can be specified statically using external tools. In the latter case, a descriptive language such as IDL or XML is used to define the data contents. From this, the code/data that is needed to handle the Topic data and the on-the-wire encoding/decoding logic is generated.

See the figure below for the relationship between Writers, Readers, Topics and types:



Types that can be specified are:

- Integer types (16-bit, 32-bit, 64-bit).
- Unsigned types (16-bit, 32-bit, 64-bit).

- Floating point types (32-bit, 64-bit, 128-bit)¹
- Fixed point types (up to 31 digits).²
- Bytes
- Booleans
- Chars (8-bit and 32-bit) - Allows UTF-8, UTF-16 and UTF-32.
- Bitsets
- Enumerations
- Alias
- Collections (Map, Sequence, Array and String).
- Aggregations (Union and Struct)

Some metadata is supported as well:

- Annotations
- Verbatim text



Note: The definition of which types can be expressed has been extended recently with an additional DDS specification, called X-types, or the Extended Types specification, which has added support for the Map type, Annotations and Verbatim text. Originally these features were not available. The X-types specification has also added support for sparse data structures and evolution of data types, as well as providing additional marshalling methods, and an additional API for dynamically adding types and for dynamically accessing data. Introspection of types, locally generated or received from peers is also possible.



Note: Technicolor DDS properly supports the extensions that are described in the X-types specification, as will be demonstrated further.

Technicolor DDS Type Support

In Technicolor DDS, static types are specified in an intermediate type description format, which is basically an array of C-data structures.

This type description can be generated in different ways:

- Created from the Qeo component framework, which is a layer on top of Technicolor DDS, and which provides tools to generate this type description from an XML-based service description.
- A tool is available that can be used to generate the intermediate type from an IDL-based type specification.
- It can be hand-written, although the preferred way is of course to use one of the tools described above.

The intermediate type data structure is called the TSM format and has the following layout (as defined in dds/dds_types.h):

```
struct dds_typesupport_meta_st {
    CDR_TypeCode_t      tc;          /* Type code. */
    tsm_flags            flags;       /* Key field, Dynamic data, ... */
    const char           *name;       /* Name of the field (info). */
    size_t               size;        /* Size of the corresponding
                                     container or C string. */
    size_t               offset;      /* Offset of the field in its
                                     corresponding container. */
    unsigned int         nelem;        /* Number of elements in the
                                     container. */
    int                  label;        /* The label of a union member. */
    const DDS_TypeSupport_meta *tsm;   /* When TYPEREF is instantiated. */
};
```

¹ Technicolor DDS only supports 128-bit floating point numbers when compiled with a special preprocessor definition.

² Not currently supported in Technicolor DDS.

The actual data in this structure depends on the type code id, which can have the following possible values as described below:

```
typedef enum {
    CDR_TYPECODE_SHORT = 1,      /* 16-bit signed integer. */
    CDR_TYPECODE_USHORT,        /* 16-bit unsigned integer. */
    CDR_TYPECODE_LONG,          /* 32-bit signed integer. */
    CDR_TYPECODE_ULONG,         /* 32-bit unsigned integer. */
    CDR_TYPECODE_LONGLONG,      /* 64-bit signed integer. */
    CDR_TYPECODE_ULONGLONG,     /* 64-bit unsigned integer. */
    CDR_TYPECODE_FLOAT,         /* 32-bit floating point number. */
    CDR_TYPECODE_DOUBLE,        /* 64-bit floating point number. */
#ifdef LONGDOUBLE
    CDR_TYPECODE_LONGDOUBLE,    /* 128-bit floating point number. */
#endif
    CDR_TYPECODE_FIXED,         /* Fixed-point number. */
    CDR_TYPECODE_BOOLEAN,       /* Boolean flag. */
    CDR_TYPECODE_CHAR,          /* 8-bit character. */
    CDR_TYPECODE_WCHAR,         /* 32-bit wide character. */
    CDR_TYPECODE_OCTET,         /* 8-bit byte. */
    CDR_TYPECODE_CSTRING,       /* 8-bit character string. */
    CDR_TYPECODE_WSTRING,       /* 32-bit wide character string. */
    CDR_TYPECODE_STRUCT,        /* C-struct. */
    CDR_TYPECODE_UNION,         /* C-union. */
    CDR_TYPECODE_ENUM,          /* C enumeration type. */
    CDR_TYPECODE_SEQUENCE,      /* Variable length element array. */
    CDR_TYPECODE_ARRAY,         /* Fixed length element array. */
    CDR_TYPECODE_TYPEREF,       /* Nest another type inline. */
    CDR_TYPECODE_TYPE           /* Refers to another type. */
} CDR_TypeCode_t;
```

The flags field currently has the following valid bits:

```
#define TSMFLAG_KEY      1      /* Is a key field. */
#define TSMFLAG_DYNAMIC  2      /* Container structure has dynamic data. */
#define TSMFLAG_MUTABLE  4      /* Extra container fields can be added later.
 */
#define TSMFLAG_OPTIONAL 8      /* Field is optional, i.e. is a pointer field
 that
 can be skipped. */
#define TSMFLAG_SHARED   16     /* Field is a pointer field instead of a real
 data
 field. */
#define TSMFLAG_GENID    32     /* Generate the memberId when mutable. */
```

A constructed DDS type must always start with a *struct* TSM, which has additional TSMs directly following the first for each individual field of the structure.

The significance of the TSM fields is described in the following table, depending on the type code:

Table 1: TSM fields

tc	flags	name	size	offset	nelem	label	tsm
SHORT	KEY	char pointer to name ³ .	-	Container offset in bytes.	-	-	-
USHORT	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-

³ Although the type name is only really required for the top-level TSM-struct, the name field will also be used when SQL content filter specifications refer to these fields, in order to find the proper data offset in received data samples.

tc	flags	name	size	offset	nelem	label	tsm
LONG	KEY	<i>idem</i>	-	<i>idem</i>	-	Enum value (or not significant).	-
ULONG	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
LONGLONG	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
ULONGLONG	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
FLOAT	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
DOUBLE	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
LONGDOUBLE	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
FIXED	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
BOOLEAN	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
CHAR	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
WCHAR	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
OCTET	KEY	<i>idem</i>	-	<i>idem</i>	-	-	-
CSTRING	KEY, DYNAMIC	<i>idem</i>	-	<i>idem</i>	C-string size or 0 if dynamic (char *)	-	-
WSTRING	KEY, DYNAMIC	<i>idem</i>	-	<i>idem</i>	C-wstring size or 0 if dynamic (wchar *)	-	-
STRUCT	KEY, DYNAMIC, MUTABLE	<i>idem</i>	Total C-struct data size in bytes.	<i>idem</i>	Number of struct fields (that follow as TSMs).	-	-
UNION	KEY, DYNAMIC, MUTABLE	<i>idem</i>	Total C-union data size in bytes	<i>idem</i>	Number of union fields (that follow as TSMs).	-	-
ENUM	KEY, MUTABLE	<i>idem</i>	-	<i>idem</i>	Number of enum elements (which follows as LONG TSMs).	-	-
SEQUENCE	KEY, DYNAMIC, MUTABLE	<i>idem</i>	-	<i>idem</i>	Maximum number of sequence elements or 0. Sequence	-	-

tc	flags	name	size	offset	nelem	label	tsm
					type follows as TSM.		
ARRAY	KEY, DYNAMIC, MUTABLE	<i>idem</i>	Total C-array size in bytes.	<i>idem</i>	Number of array elements (type follows as additional TSM).	-	-
TYPEREFP	KEY	<i>idem</i>	-	<i>idem</i>	-	-	Points to included type.
TYPE	KEY	<i>idem</i>	-	<i>idem</i>	-	-	Points to referred type.

Notes on the flags field.

Fields with the TSMFLAG_OPTIONAL and TSMFLAG_SHARED are converted to C-pointer types to actual data. If the data is not present, a NULL pointer will be used.

A field with TSMFLAG_GENID set will have its MemberId calculated via a callback function. Qeo typically uses a CRC32 function on the member name in order to create MemberIds. This MemberId is used in the context of structures and unions that are declared as MUTABLE, specifying the index of the member. If this flag is not set, the user can specify it in the label field (non-zero label field), or let DDS generate the next MemberId automatically (label set to zero).

Typically C-structs are allowed to embed substructures. This can easily be expressed with the TSM mechanism by specifying fields of a CDR_TYPECODE_STRUCT type, which are then followed again by structure field definitions in a recursive manner.

The handling of the flags field in embedded structs as well as for CDR_TYPECODE_TYPEREFP and CDR_TYPECODE_TYPE referenced structs can become somewhat complex due to this. The following rules are therefore used in Technicolor DDS to ensure a consistent behavior:

- If a struct reference has the KEY attribute, but none of its fields has the KEY attribute, this KEY attribute will automatically be applied to all structure fields, as if each field had it set specifically.
- If fields have the KEY attribute, but the struct reference doesn't, no fields will be used as keys. The parent structure must thus always have the KEY attribute set in order for any of its member fields to have this also.
- The DYNAMIC attribute **must** be set for all structure fields that have pointer data, such as dynamic strings (having a 0-length, i.e. indicating *char **) and sequences.
- The DYNAMIC attribute must be properly set for parents. I.e. if any field is dynamic, the parent structure must have the DYNAMIC attribute as well, and this must propagate up to the highest level container struct.

Registering a type in DDS

Once a type is created as a list of TSMs, it can be used as the foundation of a new Technicolor DDS type. This can be done using the `DDS_DynamicType_register()` function as follows:

```
DDS_TypeSupport *DDS_DynamicType_register(const DDS_TypeSupport_meta *tc);
```

Where tc is a pointer to the first TSM in the list.

When this function is successful, i.e. it has managed to convert the TSM list to the type representation used internally in Technicolor DDS, a new type is returned that can be used for registering the type within the DDS domains. If not successful, the function returns a NULL- result.

For each domain where the type is to be used, the following function should to be called:

```
DDS_ReturnCode_t DDS_DomainParticipant_register_type(
    DDS_DomainParticipant self,
    const DDS_TypeSupport *ts,
    const char *type_name
);
```

Where self is the DomainParticipant which should have been created previously, ts is the TypeSupport pointer that was registered via *DDS_DynamicType_register()*, and type_name is the name of the type within the domain.

If something went wrong with registering the type, a standard DDS error code will be returned.

This function may be called multiple times, in case the type is needed in multiple domains.

If a type is no longer needed in a domain, the *DDS_DomainParticipant_unregister_type()* function can be used. This is not really needed, however, since this unregister function is used implicitly for each registered type while cleaning up DomainParticipants, when the *DDS_DomainParticipant_delete_contained_entities()* function is used.

If the type is no longer needed in any DDS domain, the *DDS_DynamicType_free()* function can be called to free all internal type representation data.

Example Type usage

Shapes type.

A very simple example from the shapes demo as an IDL type definition:

```
const long MAXC = 127;

struct Shape_t {
    string<MAXC> color;    //@Key
    int          x;
    int          y;
    int          shapesize;
};
```

As a C-language type definition:

```
#define MAXC 127

typedef struct {
    char color[MAXC+1];    /*@Key*/
    int  x;
    int  y;
    int  shapesize;
} Shape_t;
```

This would be specified as a static type definition in TSM format as:


```
#include <stddef.h>

static DDS_TypeSupport_meta shape_tsm [] = {
    { CDR_TYPECODE_STRUCT, TSMFLAG_KEY, "Shape_t",
      sizeof (Shape_t), 0, 4, 0, NULL },
    { CDR_TYPECODE_CSTRING, TSMFLAG_KEY, "color",
```

```

    MAXC+1, offsetof (Shape_t, color), 0, 0, NULL },
{ CDR_TYPECODE_LONG, 0, "x",
  0, offsetof (Shape_t, x), 0, 0, NULL },
{ CDR_TYPECODE_LONG, 0, "y",
  0, offsetof (Shape_t, y), 0, 0, NULL },
{ CDR_TYPECODE_LONG, 0, "shapessize",
  0, offsetof (Shape_t, shapessize), 0, 0, NULL }
};

```

 **Note:** This conversion can be done either manually, or by using the IDL data compiler.

In order to use this type in a DDS client program, the type still needs to be converted to a DDS_TypeSupport type:

```

DDS_TypeSupport *shape_ts;

shape_ts = DDS_DynamicType_register (shape_tsm);
if (!shape_ts)
    return (DDS_RETCODE_OUT_OF_RESOURCES);

```

This type representation can be registered in a DDS Domain with the following function:

```

error = DDS_DomainParticipant_register_type (part, shape_ts,
"ShapeType");
if (error)
    return (error);

```

Sending data of this type could then be done as follows:

```

Shape_t data;

strcpy (data.color, "RED");
data.x = 10;
data.y = 50;
data.shapessize = 25;
error = DDS_DataWriter_write (w, &data, 0);

```

When a sample of this type is received, the following code might be used:

```

error = DDS_DataReader_take (dr, &ssample, &sinfo, 1, ss, vs, is);
if (!error && DDS_SEQ_LENGTH (sinfo)) {
    info = DDS_SEQ_ITEM (sinfo, 0);
    if (info->valid_data) {
        sample = DDS_SEQ_ITEM (ssample, 0);
        printf ("%s x=%u, y=%u, size=%u, sample->color,
                sample->x, sample->y,
                sample->shapessize);
    }
    DDS_DataReader_return_loan (dr, &ssample, &sinfo);
}

```

Simple types and enums.

A slightly more complex example but using only simple types and an enumeration type:

First we show the IDL definition:

```

enum anenum2 {
    a,
    b,
    c
}

```

```
};

struct s1 {
    unsigned short ul6;
    short il6;
    unsigned long u32;
    long i32;
    unsigned long long u64;
    long long i64;
    float fl;
    double d;
    char ch;
    bool b;
    octet o;
    anenum2 anenumelement;
};
```

The corresponding C-type:

```
#include <stdint.h>

typedef enum {
    a,
    b,
    c
} anenum2;

struct s1 {
    uint16_t ul6;
    int16_t il6;
    uint32_t u32;
    int32_t i32;
    uint64_t u64;
    int64_t i64;
    float fl;
    double d;
    char ch;
    unsigned char b;
    unsigned char o;
    anenum2 anenumelement;
};
```

Resulting in the following TSM definitions⁴:

```
#include <stddef.h>

static DDS_TypeSupport_meta tsm1[] = {
    { .tc = CDR_TYPECODE_STRUCT, .name = "s1",
      .size = sizeof(struct struct1), .nelem = 12 },
    { .tc = CDR_TYPECODE_USHORT, .name = "ul6",
      .offset = offsetof(struct s1, ul6) },
    { .tc = CDR_TYPECODE_SHORT, .name = "il6",
      .offset = offsetof(struct s1, il6) },
    { .tc = CDR_TYPECODE_ULONG, .name = "u32",
      .offset = offsetof(struct s1, u32) },
    { .tc = CDR_TYPECODE_LONG, .name = "i32",
      .offset = offsetof(struct s1, i32) },
    { .tc = CDR_TYPECODE_ULONGLONG, .name = "u64",
      .offset = offsetof(struct s1, u64) },
    { .tc = CDR_TYPECODE_LONGLONG, .name = "i64",
```

⁴ This time we use the GCC extension for naming fields in constant declarations.

```

        .offset = offsetof(struct s1, i64) },
    { .tc = CDR_TYPECODE_FLOAT, .name = "fl",
      .offset = offsetof(struct s1, fl) },
    { .tc = CDR_TYPECODE_DOUBLE, .name = "d",
      .offset = offsetof(struct s1, d) },
    { .tc = CDR_TYPECODE_CHAR, .name = "ch",
      .offset = offsetof(struct s1, ch) },
    { .tc = CDR_TYPECODE_BOOLEAN, .name = "b",
      .offset = offsetof(struct s1, b) },
    { .tc = CDR_TYPECODE_OCTET, .name = "o",
      .offset = offsetof(struct s1, o) },
    { .tc = CDR_TYPECODE_ENUM, .name = "anenumelement",
      .offset = offsetof(struct s1, anenumelement), .nelem = 3 },
    { .tc = CDR_TYPECODE_LONG, .name = "a", .label = a },
    { .tc = CDR_TYPECODE_LONG, .name = "b", .label = b },
    { .tc = CDR_TYPECODE_LONG, .name = "c", .label = c }
};

```

Arrays, sequences and strings

Following example contains some fixed, as well as dynamic array types and some dynamic and static strings.

The example in IDL looks like this:

```

const long MSGLEN = 20;

struct s2 {
    string<MSGLEN>    message;
    string            dynmsg;
    double            darr [10][2];
    sequence<float>    floats;
};

```

This example uses a dynamic array, called a sequence. Although sequences are handled in more detail in the following chapter, a small introduction is given here in order to be able to understand the next code fragments.

Since there is no direct C-language representation of a dynamic array, a macro is used to define a sequence of float values.

```

#include "dds/dds_seq.h"

DDS_SEQUENCE (float, fseq);

```

This macro will expand to the following C-type definition:

```

typedef struct fseq_st {
    unsigned _maximum;
    unsigned _length;
    unsigned _esize;
    int      _own;
    float     *_buffer;
} fseq;

```



Note: Sequences should not be handled directly. There is a large set of sequence handling prototypes defined in `dds_seq.h` that should be sufficient for the user and which allows almost direct use of all the sequence fields. See the next chapter for details.

The example uses this defined sequence:

```

#define MSGLEN 20

struct s2 {

```

```

char message [MSGLEN];
char *dynmsg;
double darr [10][2];
fseq floats;
};

```

This will result in the following TSMs:

```

#include <stddef.h>

static DDS_TypeSupport_meta tsm2[] = {
    { .tc = CDR_TYPECODE_STRUCT, .flags = TSMFLAG_DYNAMIC, .name = "s2",
      .size = sizeof (struct s2), .nelem = 4 },
    { .tc = CDR_TYPECODE_CSTRING, .name = "message", .nelem = MSGLEN,
      .offset = offsetof (struct s2, message) },
    { .tc = CDR_TYPECODE_CSTRING, .flags = TSMFLAG_DYNAMIC,
      .name = "dynmsg", .nelem = 0, },
      .offset = offsetof (struct s2, dynmsg ) },
    { .tc = CDR_TYPECODE_ARRAY, .name = "darr",
      .size = sizeof (double) * 20, .nelem = 10,
      .offset = offsetof (struct s2, darr)},
    { .tc = CDR_TYPECODE_ARRAY, .size = sizeof (double) * 2, .nelem = 2 },
    { .tc = CDR_TYPECODE_DOUBLE },
    { .tc = CDR_TYPECODE_SEQUENCE, .flags = TSMFLAG_DYNAMIC,
      .name = "floats", .nelem = 0,
      .offset = offsetof (struct s2, fseq)},
    { .tc = CDR_TYPECODE_FLOAT}
};

```

If a dynamic type, such as this, is given to the user via a DDS Reader, it is up to the user to determine where the dynamic data will be located. How this works will be described when sequences are discussed in more detail.



Note: It is possible to constrain the size of embedded sequences in a TSM, if deemed necessary, by simply setting the *nelem* field to the maximum allowed number of elements in the sequence. If not constrained, the number of elements in the sequence may be unlimited, which is clearly impractical.

Union types

Union types in DDS are handled somewhat differently from normal C-unions, since DDS requires a special and unique discriminant value for each distinct union value. To make DDS union handling therefore somewhat easier, the `DDS_UNION` macro should be used:

```
DDS_UNION (union_type, discr_type, un);
```

Which expands to:

```

typedef struct un_st {
    discr_type discriminant; /* Extra discriminant field. */
    union_type u;           /* Union data fields. */
} un;

```

The following example union will demonstrate how to handle TSMs for unions:

First we'll show the example using a union in IDL:

```

enum discr {
    ANINT,
    ASTRING
};

const long STRLEN = 20;

```

```

union dds_union switch (discr) {
    case ANINT: int i;
    case ASTRING: string<STRLEN> str;
};

struct uls {
    dds_union aunion;
};

```

The same example, but resulting C-code:

```

#include "dds/dds_types.h"

typedef enum {
    ANINT,
    ASTRING
} discr;

#define STRLEN 20

typedef union {
    int i;
    char str [STRLEN+1];
} ul;

DDS_UNION (ul, discr, dds_union);

typedef struct uls_st {
    dds_union aunion;
} uls;

```

Which will then result in the following TSMs:

```

#include <stddef.h>

static DDS_TypeSupport_meta uls_tsm[] = {
    { .tc = CDR_TYPECODE_STRUCT, .name = "uls",
      .size = sizeof (uls), .nelem = 1 },
    { .tc = CDR_TYPECODE_UNION, .name = "aunion", .nelem = 2,
      .size = sizeof (dds_union), .offset = offsetof (uls, aunion) },
    { .tc = CDR_TYPECODE_LONG, .name = "along",
      .offset = offsetof (dds_union, u), .label = ANINT },
    { .tc = CDR_TYPECODE_CSTRING, .name = "acstring",
      .offset = offsetof (dds_union, u),
      .size = STRLEN+1, .label = ASTRING }
};

```

Using Sequences

Introduction

Support for dynamic arrays, called sequences in DDS is supported directly via a set of sequence handling functions and macros. There is really no need for direct access to the sequence type itself, since the support functions should be able to provide all required accesses.

Sequences are strictly typed. Any previously defined type can be used as the element of a sequence, even sequences of sequences are allowed.

Whereas a number of programming languages have direct support for sequences, not all of them do. In C++, sequences are handled using templates and overloading of standard operators. In C, there is no direct support in the programming language, and a macro is used to define the sequence and a set of access functions and macros is used for working with it.

Creating a new sequence type

Creating a new sequence type is done via the following macro:

```
#include "dds/dds_seq.h"

DDS_SEQUENCE(<element_type>, <sequence_type>);
```

Where the `<element_type>` can be whatever previously defined type.

The result will be the following type definition:

```
typedef struct name <sequence_type>_st {
    unsigned    _maximum;
    unsigned    _length;
    unsigned    _esize;
    int         _own;
    <element_type> *_buffer;
} <sequence_type>;
```



Note: Although we will describe each sequence field in detail, it is not the intention that the user accesses these fields directly. This is just for understanding how sequences work. Actual handling of sequences should be via the access functions and macros, since this is an internal representation which might still be changed in future revisions of the software. Using the access functions and macros will thus ensure portability of user code to new Technicolor DDS versions.

The `_maximum` field specifies the maximum number of sequence elements that are currently available in `_buffer`. If the length of `_buffer` changes due to more elements being added to the sequence, the `_maximum` field will be adapted appropriately.

The `_length` field specifies the actual number of elements in the sequence. The `_length` field might vary between 0 and `_maximum`, but it may never be larger than `_maximum`. Resetting a sequence can be done simply by setting `_length` to 0. No memory needs to be allocated for adding elements unless the sequence contains (`_length == _maximum`) elements. In that case the `_buffer` field as well as `_maximum` will need to be updated before the element can be added.

The `_esize` field specifies the size of a sequence element. It is usually set when the instance is created, and initialized with `sizeof(<element_type>)`.



Note: Although one might expect that the `_esize` type, which indicates the size of a particular sequence element, would be defined as a `size_t` type, this is not the case. This was a design choice. On 64-bit machines the compiler typically generates a 32-bit type for unsigned types, and a 64-bit type for `size_t` types, which we feel is clearly impractical in the context of DDS, where memory savings are more important, even on 64-bit machines. Besides, a 32-bit type is already capable of representing up to 4 Gigabytes of memory, and transferring chunks of more than 4 Gigabytes in a single message is clearly out of the question.

The `_own` field specifies ownership of the sequence element data. If `_own` is set, it means that the user of the sequence may change any element within the sequence. If not set, the sequence was passed to the called function by reference only, and the function is not supposed to change the sequence contents. See further for some specific use-cases.

The `_buffer` field is a pointer to an array of sequence elements. This array will initially be set to NULL, but will be allocated and reallocated appropriately when elements are subsequently appended to the sequence.

Creating and deleting sequences

Sequences can be either statically initialized when declared, or initialized explicitly. Static initialization is done using the `DDS_SEQ_INITIALIZER()` macro:

```
#include "dds/dds_seq.h"

DDS_SEQUENCE(float, FloatSeq);

static FloatSeq fs = DDS_SEQ_INITIALIZER(float);
```

The alternative explicit initialization uses the `DDS_SEQ_INIT()` macro:

```
#include "dds/dds_seq.h"

DDS_SEQUENCE(double, DoubleSeq);

int test(void)
{
    DoubleSeq dseq;

    DDS_SEQ_INIT(dseq);
    ...
}
```

Once a sequence has been used, i.e. elements were appended to it, it can be reset for reuse with the following generic function:

```
void dds_seq_reset(void *seq);
```

This function doesn't really cleanup all data, it just resets the sequence so that it becomes empty.

If a sequence is no longer needed, all its contained data can be disposed completely with the following generic function:

```
void dds_seq_cleanup(void *seq);
```



Note: Both functions will only update the sequence descriptor. If the element type is a pointer type, the user should take care that all pointed to data is also disposed, possibly by walking over the sequence first, freeing the data that is pointed to, before calling the functions.



Note: Don't use the `dds_seq_reset/cleanup()` functions on sequences that are created by Technicolor DDS. They are really meant for user-created sequences only. For DDS-originated sequences, although these functions might or might not work, depending on the element type, there are specific cleanup functions predefined that know exactly how to handle a correct cleanup. Therefore it is prudent to use only the DDS-specific cleanup functions in those situations, as described further.

Handling sequence elements

Once a sequence is initialized, data can be added to it, element by element in the following manner:

```
DDS_ReturnCode_t dds_seq_append(void *seq, void *value);
DDS_ReturnCode_t dds_seq_prepend(void *seq, void *value);
DDS_ReturnCode_t dds_seq_insert(void *seq, unsigned pos, void *value);
```

The first function adds an element at the end of the sequence, the second prepends an element, before all the existing elements, by shifting all these elements one position further in the sequence. The last function inserts a sequence element at a specified position in the sequence, hereby shifting all existing elements from that position onwards one position further.

A sequence can be completely initialized from an array of elements, sequence data can be stored in an array, using the following functions:

```
DDS_ReturnCode_t dds_seq_from_array(void *seq, void *array, unsigned len);
unsigned dds_seq_to_array(void *seq, void *array, unsigned len);
```

The first populates an empty sequence with array data. Note that the sequence must be empty for this to work properly. The second copies existing sequence data to an array, without changing the sequence contents.

Any existing element in a sequence can be replaced using the *dds_seq_replace()* function:

```
DDS_ReturnCode_t dds_seq_replace(void *seq, unsigned pos, void *value);
```

Removing elements from a sequence can be done with the following functions:

```
DDS_ReturnCode_t dds_seq_remove_first(void *seq, void *value);
DDS_ReturnCode_t dds_seq_remove_last(void *seq, void *value);
DDS_ReturnCode_t dds_seq_remove(void *seq, unsigned pos, void *value);
```

These functions remove elements from the sequence from the first, the last or from an arbitrary position in the sequence respectively.

Processing sequence elements often requires walking over the sequence, element by element. There are a number of ways to do this:

```
DDS_SEQ_FOREACH(seq,i) { ... }
DDS_SEQ_FOREACH_ENTRY(seq,i,ptr) { ... }
```

The first macro sets *i* initially at 0 and keeps incrementing it until the sequence end is reached. If the sequence is empty, nothing is done, otherwise the statement body will be called for each element in succession.

The difference between the two macros is that the latter sets *ptr* to the location of each element while walking over the sequence.

A more structured way of walking is with the following function:

```
void dds_seq_every(void *seq,
                  void (*func)(void *ptr, void *data),
                  void *data);
```

The *dds_seq_every()* function will call the subfunction *func* for each element, with the element pointer as the first argument and the passed data pointer as the second.

Although these functions are very generic, sometimes it is necessary to handle sequence elements in a more direct manner. Following macros give access to a sequence element at a specified position:

```
DDS_SEQ_ITEM(seq,i)
DDS_SEQ_ITEM_SET(seq,i,v)
DDS_SEQ_ITEM_PTR(seq,i)
```

The first simply returns the element of a sequence at position *i*. The second replaces the element at the given position with *v*.



Note: Neither *DDS_SEQ_ITEM()* nor *DDS_SEQ_ITEM_SET()* validates the given element position. This must have been checked by the caller upfront.

The last macro returns a pointer to a sequence element location or NULL if the element is not in the sequence, so this does check the position.

Direct access to the sequence descriptor fields is possible using a number of macros, although one should take extreme care when using these.

```
DDS_SEQ_LENGTH(seq)
DDS_SEQ_MAXIMUM(seq)
DDS_SEQ_OWNED(seq)
DDS_SEQ_OWNED_SET(seq, o)
DDS_SEQ_ELEM_SIZE(seq)
```

The *DDS_SEQ_LENGTH()* returns the current sequence length, so this could be used to implement an alternative iterator. The *DDS_SEQ_MAXIMUM()* returns the current element buffer size. The *DDS_SEQ_OWNED()* and *DDS_SEQ_OWNED_SET()* are getter and setter functions for the sequence ownership field. *DDS_SEQ_ELEM_SIZE()* returns the size in bytes of a sequence element.

DDS-specific sequences

The former sequence operations are of course general manipulations of any sequence, user-defined or Technicolor DDS derived.

In practice, however, the user will also be confronted with the sequences that are defined in DDS itself. Technicolor DDS uses and has already defined a number of DDS sequences and sequence operations for specific DDS purposes.

Many of these sequences have a set of associated type-specific functions defined to make it easier for the user to create and dispose of them.

These functions come in a number of flavors and typically have the following prototypes:

```
<seq_type> *<seq_type>__alloc(void);
void <seq_type>__free(<seq_type> *arg);
void <seq_type>__init(<seq_type> *arg);
void <seq_type>__clear(<seq_type> *arg);
```

Here, the *<seq_type>__alloc()* function will attempt to allocate from the heap and reset appropriately a new sequence of the given type. If there is no more memory, a NULL value will be returned.

The *<seq_type>__free()* function is the reverse of the previous operation and will cleanly dispose the full sequence contents before releasing the sequence back to the heap.

The *<seq_type>__init()* function can be used to initialize a variable of the specific sequence type for the first time. It should not be called on an already initialized sequence since this will destroy (and leak) the contained data.

The last function, i.e. *<seq_type>__clear()* can be used to properly dispose of existing sequence data. It results in a clean empty sequence.

These functions are defined in Technicolor DDS for the following sequence types as well as for a number of container types that contain sequences:

Table 2: DDS Sequences and container types

Sequence or Container type	Element Type	Description
DDS_OctetSeq	unsigned char	Byte sequence, used as QoS parameter for User, Topic and Group data.
DDS_StringSeq	char *	A sequence of strings.
DDS_DataSeq	void *	An abstract data sequence.
DDS_SampleInfoSeq	DDS_SampleInfo *	Sequence of sample information as used in <i>_read/_take()</i> and <i>_return_loan()</i> functions.
DDS_InstanceHandleSeq	DDS_InstanceHandle_t	Sequence of instance handles.

Sequence or Container type	Element Type	Description
DDS_DataReaderSeq	DDS_DataReader	Sequence of DataReaders.
DDS_GroupDataQosPolicy	-	Group data QoS contains a DDS_OctetSeq.
DDS_PartitionQosPolicy	-	Partition QoS contains a DDS_StringSeq.
DDS_TopicDataQosPolicy	-	Topic data QoS contains a DDS_OctetSeq.
DDS_UserDataQosPolicy	-	User data QoS contains a DDS_OctetSeq.
DDS_ParticipantBuiltinTopicData	-	Contains a DDS_UserDataQosPolicy element.
DDS_TopicBuiltinTopicData	-	Contains various dynamic data elements.
DDS_PublicationBuiltinTopicData	-	Contains various dynamic data elements.
DDS_SubscriptionBuiltinTopicData	-	Contains various dynamic data elements.
DDS_DomainParticipantQos	-	Contains a UserDataQosPolicy element.
DDS_TopicQos	-	Contains a TopicDataQosPolicy element.
DDS_SubscriberQos	-	Contains various dynamic data elements.
DDS_PublisherQos	-	Contains various dynamic data elements.
DDS_DataReaderQos	-	Contains a DDS_UserDataQosPolicy element.
DDS_DataWriterQos	-	Contains a DDS_UserDataQosPolicy element.
DDS_ConditionSeq	DDS_Condition	Sequence of Conditions.

Using the `_own` field in DDS DataReader functions

In practice, the `_own` field is only used in the *DDS_DataReader_read/take()* functions to indicate to which layer the data buffers in the sequence belong. The following three mechanisms can be distinguished:

1. The user specifies own sample data buffers (`_own = 1` in *DDS_DataReader_read/take()* functions), and the dynamic contained data areas will be already allocated by the user, by setting the individual pointer fields to a non-NULL value. For the above example, the *dynmsg* and *dptr* fields and the *floats._buffer* would already be populated with pointers to valid and large enough data areas. The sequence maximum length (`_maximum`) must then also be sufficiently large to contain the maximum size of received elements. No *DDS_DataReader_return_loan()* is necessary after processing of the data samples and the user is completely responsible for managing its own data.
2. The user specifies, as before, sample data buffers (`_own = 1` in *DDS_DataReader_read/take()* functions), but the dynamic data fields, i.e. *dynmsg* and *dptr* fields and the *floats._buffer* are initialized to NULL. In this case, Technicolor DDS will allocate the dynamic memory using the builtin C-memory allocator (typically *malloc()*⁵). When done processing the samples, the user is responsible for cleaning up the data using the free memory function of the C-memory allocator (typically *free()*). As in the first case, no *DDS_DataReader_return_loan()* is necessary after processing of the data samples.
3. The user doesn't give sample data buffers (`_own = 0` or both sample data and sample info sequences initialized to empty) in *DDS_DataReader_read/take()* functions,. In this case Technicolor DDS is responsible for allocating all data, typically using another, more efficient but not disclosed, allocation strategy. Note that the use of this mechanism is a lot easier from the point of view of the caller, but it does mandate the use of *DDS_DataReader_return_loan()* when processing of data samples is done. It is also the more efficient method of processing received data since no copying is necessary - the user merely 'borrows' the buffers from DDS.

⁵ In the future, the memory allocator function will be settable by the user, but this is not the case yet.

Using Dynamic types

Dynamic Shapes type.

The same Shapes type as defined before, can be constructed using the Extended types API ⁶ as demonstrated in the following code samples⁷:

```
#include "dds/dds_xtypes.h"

DDS_ReturnCode_t create_dyn_shape (void)
{
    DDS_TypeDescriptor      *desc;
    DDS_MemberDescriptor    *md;
    DDS_DynamicTypeMember   *dtm;
    DDS_AnnotationDescriptor *key_ad;
    DDS_DynamicTypeBuilder  sb, ssb;
    DDS_DynamicType         s, ss;

    desc = DDS_TypeDescriptor__alloc ();
    if (!desc)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    md = DDS_MemberDescriptor__alloc ();
    if (!md)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    key_ad = DDS_AnnotationDescriptor__alloc ();
    if (!key_ad)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    key_ad->type = DDS_DynamicTypeBuilderFactory_get_builtin_annotation
("Key");
    if (!key_ad->type)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    desc->kind = DDS_STRUCTURE_TYPE;
    desc->name = "ShapeType";
    sb = DDS_DynamicTypeBuilderFactory_create_type (desc);
    if (!sb)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    ssb = DDS_DynamicTypeBuilderFactory_create_string_type (128);
    if (!ssb)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    ss = DDS_DynamicTypeBuilder_build (ssb);
    if (!ss)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    md->name = "color";
    md->index = md->id = 0;
    md->type = ss;
    md->default_value = NULL;
    if (!md->type)
        return (DDS_RETCODE_OUT_OF_RESOURCES);
```

⁶ The Extended types API is available when Technicolor DDS has been built with the `-DXTYPES_USED` build flag.

⁷ In order not to clutter the code with error handling, the following examples don't properly cleanup all the resources in case errors should occur. Actual error handling when using dynamic types in production code should really be done somewhat more rigidly.

```

error = DDS_DynamicTypeBuilder_add_member (sb, md);
if (error)
    return (error);

dtm = DDS_DynamicTypeMember__alloc ();
if (!dtm)
    return (DDS_RETCODE_OUT_OF_RESOURCES);

error = DDS_DynamicTypeBuilder_get_member (sb, dtm, 0);
if (error)
    return (error);

error = DDS_DynamicTypeMember_apply_annotation (dtm, key_ad);
if (error)
    return (error);

free (dtm->descriptor.name);
dtm->descriptor.name = NULL;
DDS_DynamicTypeBuilderFactory_delete_type (dtm->descriptor.type);
dtm->descriptor.type = NULL;
md->name = "x";
md->index = md->id = 1;
md->type = DDS_DynamicTypeBuilderFactory_get_primitive_type
                                                (DDS_INT_32_TYPE);

if (!md->type)
    return (DDS_RETCODE_OUT_OF_RESOURCES);

error = DDS_DynamicTypeBuilder_add_member (sb, md);
if (error)
    return (error);

md->name = "y";
md->index = md->id = 2;
md->type = DDS_DynamicTypeBuilderFactory_get_primitive_type
                                                (DDS_INT_32_TYPE);

if (!md->type)
    return (DDS_RETCODE_OUT_OF_RESOURCES);

error = DDS_DynamicTypeBuilder_add_member (sb, md);
if (error)
    return (error);

md->name = "shapesize";
md->index = md->id = 3;
md->type = DDS_DynamicTypeBuilderFactory_get_primitive_type
                                                (DDS_INT_32_TYPE);

if (!md->type)
    return (DDS_RETCODE_OUT_OF_RESOURCES);

error = DDS_DynamicTypeBuilder_add_member (sb, md);
if (error)
    return (error);

s = DDS_DynamicTypeBuilder_build (sb);
if (!s)
    return (DDS_RETCODE_OUT_OF_RESOURCES);

DDS_DynamicTypeBuilderFactory_delete_type (ssb);
DDS_DynamicTypeBuilderFactory_delete_type (sb);
dtype = s;
shape_ts = DDS_DynamicTypeSupport_create_type_support (s);
DDS_DynamicTypeBuilderFactory_delete_type (ss);
DDS_TypeDescriptor__free (desc);
DDS_MemberDescriptor__free (md);

```

```

DDS_AnnotationDescriptor__free (key_ad);
DDS_DynamicTypeMember__free (dtm);
return (DDS_RETCODE_OK);
}

```

A Shapes dynamic data sample can then be constructed and written as shown in the following code fragment⁸:

```

DDS_DynamicData dd;

dd = DDS_DynamicDataFactory_create_data (dtype);
if (!dd)
    fatal_printf ("Can't create dynamic data!");

id = DDS_DynamicData_get_member_id_by_name (dd, "color");
rc = DDS_DynamicData_set_string_value (dd, id, "RED");
if (rc)
    fatal_printf ("Can't add data member(color)!");

id = DDS_DynamicData_get_member_id_by_name (dd, "x");
rc = DDS_DynamicData_set_int32_value (dd, id, 10);
if (rc)
    fatal_printf ("Can't add data member(x)!");

id = DDS_DynamicData_get_member_id_by_name (dd, "y");
rc = DDS_DynamicData_set_int32_value (dd, id, 60);
if (rc)
    fatal_printf ("Can't add data member(y)!");

id = DDS_DynamicData_get_member_id_by_name (dd, "shapessize");
rc = DDS_DynamicData_set_int32_value (dd, id, 25);
if (rc)
    fatal_printf ("Can't add data member(shapessize)!");

rc = DDS_DynamicDataWriter_write (w, dd, handle);

```

When a dynamic data is received, its member fields can be retrieved as shown in the following code:

```

error = DDS_DynamicDataReader_take (dr, &sample, &info, 1, ss, vs, is);
if (error) {
    if (error != DDS_RETCODE_NO_DATA)
        printf ("Unable to read: error = %s!\r\n", DDS_error (error));
    return;
}
if (DDS_SEQ_LENGTH (info)) {
    inf = DDS_SEQ_ITEM (info, 0);
    if (inf->valid_data && dyn_data) {
        dd = DDS_SEQ_ITEM (sample, 0);
        if (!dd)
            fatal_printf ("Empty dynamic sample!");

        id = DDS_DynamicData_get_member_id_by_name (dd, "color");
        error = DDS_DynamicData_get_string_value (dd, color, id);
        if (error)
            fatal_printf ("Can't get data member(color)!");

        id = DDS_DynamicData_get_member_id_by_name (dd, "x");
        error = DDS_DynamicData_get_int32_value (dd, &x, id);
        if (error)
            fatal_printf ("Can't get data member(x)!");
    }
}

```

⁸ The DDS_DynamicData_get_member_id_by_name() function calls are not really required since the id value is fixed when the type is created and is typically known in the program.

```

        id = DDS_DynamicData_get_member_id_by_name (dd, "y");
        error = DDS_DynamicData_get_int32_value (dd, &y, id);
        if (error)
            fatal_printf ("Can't get data member(y)!");

        id = DDS_DynamicData_get_member_id_by_name (dd, "shapessize");
        error = DDS_DynamicData_get_int32_value (dd, &shapessize, id);
        if (error)
            fatal_printf ("Can't get data member(shapessize)!");

        /* ... do something with the sample (color/x/y/shapessize) ... */
    }
}

```

Dynamic sparse data type:

The sparse type with the following IDL specification:

```

//@Extensibility(MUTABLE_EXTENSIBILITY)
struct struct2m {
    uint16_t u16;    //@ID(10) //@Key
    int16_t i16;     //@ID(20)
    uint32_t u32;    //@Key
    int32_t i32;     //@ID(50)
    uint64_t u64;
    int64_t i64;
    float fl;
    double d;
    char ch;         //@ID(5) //@Key
};

```

can be constructed using the given C-code:

```

#include "dds/dds_xtypes.h"

#define ADD_FIELD(s,md,n,idx,i,t) md->name=n; md->index=idx; md->id=i;\
    md->type=DDS_DynamicTypeBuilderFactory_get_primitive_type(t);\
    fail_unless(md->type != NULL);\
    rc = DDS_DynamicTypeBuilder_add_member (s,md);\
    if (rc) return (rc)

DDS_ReturnCode_t set_key_annotation (DDS_DynamicTypeBuilder b,
                                     const char *name)
{
    DDS_DynamicTypeMember dtm;
    DDS_ReturnCode_t ret;
    static DDS_AnnotationDescriptor ad = { NULL, };

    if (!b && ad.type) {
        DDS_AnnotationDescriptor__clear (&ad);
        return (DDS_RETCODE_OK);
    }
    DDS_DynamicTypeMember__init (&dtm);
    if (!ad.type) {
        ad.type = DDS_DynamicTypeBuilderFactory_get_builtin_annotation
("Key");
        if (ad.type == NULL)
            return (DDS_RETCODE_OUT_OF_RESOURCES);
    }
    ret = DDS_DynamicTypeBuilder_get_member_by_name (b, &dtm, name);
    if (ret)

```



```

        return (ret);

    ret = DDS_DynamicTypeMember_apply_annotation (&dtm, &ad);
    if (ret)
        return (ret);

    free (dtm.descriptor.name);
    dtm.descriptor.name = NULL;
    DDS_DynamicTypeBuilderFactory_delete_type (dtm.descriptor.type);
    dtm.descriptor.type = NULL;
    return (DDS_RETCODE_OK);
}

DDS_ReturnCode_t set_id_annotation (DDS_DynamicTypeBuilder b,
                                   const char          *name,
                                   DDS_MemberId         id)
{
    DDS_DynamicTypeMember      dtm;
    DDS_ReturnCode_t           ret;
    unsigned                   n;
    char                       buf [12];
    static DDS_AnnotationDescriptor ad = { NULL, };

    if (!b && ad.type) {
        DDS_AnnotationDescriptor__clear (&ad);
        return (DDS_RETCODE_OK);
    }
    DDS_DynamicTypeMember__init (&dtm);
    if (!ad.type) {
        ad.type = DDS_DynamicTypeBuilderFactory_get_builtin_annotation
("ID");
        if (!ad.type)
            return (DDS_RETCODE_OUT_OF_RESOURCES);
    }
    n = snprintf (buf, sizeof (buf), "%u", id);
    if (!n)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    ret = DDS_AnnotationDescriptor_set_value (&ad, "value", buf);
    if (ret)
        return (ret);

    ret = DDS_DynamicTypeBuilder_get_member_by_name (b, &dtm, name);
    if (ret)
        return (ret);

    ret = DDS_DynamicTypeMember_apply_annotation (&dtm, &ad);
    if (ret)
        return (ret);

    free (dtm.descriptor.name);
    dtm.descriptor.name = NULL;
    DDS_DynamicTypeBuilderFactory_delete_type (dtm.descriptor.type);
    dtm.descriptor.type = NULL;
    return (DDS_RETCODE_OK);
}

DDS_ReturnCode_t set_ext_annotation (DDS_DynamicTypeBuilder b,
                                    const char          *ext)
{
    DDS_ReturnCode_t           ret;
    static DDS_AnnotationDescriptor ad = { NULL, };

    if (!b && ad.type) {

```

```

        DDS_AnnotationDescriptor__clear (&ad);
        return (DDS_RETCODE_OK);
    }
    if (!ad.type) {
        ad.type = DDS_DynamicTypeBuilderFactory_get_builtin_annotation
("Extensibility");
        if (!ad.type)
            return (DDS_RETCODE_OUT_OF_RESOURCES);
    }
    ret = DDS_AnnotationDescriptor_set_value (&ad, "value", ext);
    if (ret)
        return (ret);

    ret = DDS_DynamicTypeBuilder_apply_annotation (b, &ad);
    return (ret);
}

DDS_ReturnCode_t create_dyn_struct2m (void)
{
    DDS_TypeSupport          *ts;
    DDS_DynamicTypeBuilder sb2;
    DDS_DynamicType          s2;
    DDS_TypeDescriptor       *desc;
    DDS_MemberDescriptor     *md;
    DDS_ReturnCode_t         rc;

    /* 1. Create the type. */
    desc = DDS_TypeDescriptor__alloc ();
    if (!desc)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    desc->kind = DDS_STRUCTURE_TYPE;
    desc->name = "dstruct2m";

    sb2 = DDS_DynamicTypeBuilderFactory_create_type (desc);
    if (!sb2)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    md = DDS_MemberDescriptor__alloc ();
    if (!md)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    ADD_FIELD (sb2, md, "i16", 0, 99, DDS_INT_16_TYPE);
    rc = set_id_annotation (sb2, "i16", 20);
    if (rc) return (rc);

    ADD_FIELD (sb2, md, "u32", 1, 21, DDS_UINT_32_TYPE);
    rc = set_key_annotation (sb2, "u32");
    if (rc) return (rc);

    ADD_FIELD (sb2, md, "i32", 2, 100, DDS_INT_32_TYPE);
    rc = set_id_annotation (sb2, "i32", 50);
    if (rc) return (rc);

    ADD_FIELD (sb2, md, "u16", 0, DDS_MEMBER_ID_INVALID, DDS_UINT_16_TYPE);
    rc = set_id_annotation (sb2, "u16", 10);
    if (rc) return (rc);

    rc = set_key_annotation (sb2, "u16");
    if (rc) return (rc);

    ADD_FIELD (sb2, md, "u64", 5, 51, DDS_UINT_64_TYPE);
    ADD_FIELD (sb2, md, "i64", 5, DDS_MEMBER_ID_INVALID, DDS_INT_64_TYPE);

```

```

ADD_FIELD (sb2, md, "fl", 6, 53, DDS_FLOAT_32_TYPE);
ADD_FIELD (sb2, md, "d", 7, 54, DDS_FLOAT_64_TYPE);
ADD_FIELD (sb2, md, "ch", 8, 55, DDS_CHAR_8_TYPE);
rc = set_id_annotation (sb2, "ch", 5);
if (rc) return (rc);

rc = set_ext_annotation (sb2, "MUTABLE_EXTENSIBILITY");
if (rc) return (rc);

s2 = DDS_DynamicTypeBuilder_build (sb2);
if (!s2) return (DDS_RETCODE_OUT_OF_RESOURCES);

DDS_DynamicTypeBuilderFactory_delete_type (sb2);

ts = DDS_DynamicTypeSupport_create_type_support (s2);
if (!ts) return (DDS_RETCODE_OUT_OF_RESOURCES);

DDS_TypeDescriptor__free (desc);
DDS_MemberDescriptor__free (md);
return (DDS_RETCODE_OK);
}

```

A data sample of this type can then be constructed as follows:

```

#define SET_FIELD(dd,id,type,v) rc=DDS_DynamicData_set_##type##_value\
(dd,id,v);\
    if (rc) return (rc)

DDS_ReturnCode_t write_sample(DDS_DynamicType s2,
                              DDS_DataWriter w,
                              DDS_HANDLE h)
{
    DDS_DynamicData dd;
    DDS_ReturnCode_t rc;

    dd = DDS_DynamicDataFactory_create_data (s2);
    if (!dd)
        return (DDS_RETCODE_OUT_OF_RESOURCES);

    SET_FIELD (dd, 50, int32, -5);
    SET_FIELD (dd, 10, uint16, 0xDEAD);
    SET_FIELD (dd, 20, int16, INT16_MIN);
    SET_FIELD (dd, 21, uint32, UINT32_MAX);
    SET_FIELD (dd, 52, int64, 100);
    SET_FIELD (dd, 53, float32, 0.5f);
    SET_FIELD (dd, 54, float64, 100e-5);
    SET_FIELD (dd, 5, char8, 'd');
    SET_FIELD (dd, 51, uint64, 5010000);

    rc = DDS_DynamicDataWriter_write (w, dd, h);
    return (rc);
}

```

WaitSets and Listeners

DDS allows two distinct methods in order to be notified of events, and both can be used simultaneously, i.e. Listener functions and WaitSets.

There are advantages and disadvantages to each of these two mechanisms, and care must be taken in order to properly handle additional DDS function calls in either case.

Following discussion will explain each mechanism in detail.

Listener functions

The first and foremost mechanism is the Listener function. This mimics a bit the Unix *signal* mechanism, in that it results in an asynchronous callback function in a thread context that is not the thread of the user's application, so it is very important that some precautions are taken to ensure proper use of this mechanism.

Listener callbacks can be specified on different entities as well as on encompassing entities on top of the former. For a listener to be active, both the function callback must be set and the Entity event mask indicating the specific listener must be set. The primary entities and their optional listeners are listed below:

Table 3: Listener functions

Entity	Specific Listener functions
DataReader	<i>requested_deadline_missed()</i> , <i>requested_incompatible_qos()</i> , <i>sample_rejected()</i> , <i>liveliness_changed()</i> , <i>data_available()</i> , <i>subscription_match()</i> , <i>sample_lost()</i>
Subscriber	<i>data_on_readers()</i>
DataWriter	<i>offered_deadline_missed()</i> , <i>offered_incompatible_qos()</i> , <i>liveliness_lost()</i> , <i>publication_match()</i>
Topic	<i>inconsistent_topic()</i>

The entities encompassing these primary entities inherit the listener functions of the contained entities. For example, the Subscriber entity has both its own specific listener (*data_on_readers()*), but also inherits every DataReader listener functions. The same can be set of the Publisher entity, since it also contains the listener functions of the DataWriter.

The DomainParticipant then inherits all listener functions, since it is a parent of all the Topics, DataReaders, DataWriters as well as of all Publishers and Subscribers in the domain.

When multiple Listeners are available on all levels, the general rule is that the most specific Listener will always be called. The listeners on the higher levels are only called if there's no listener on a more specific entity. There is however a single exception of this rule for the Subscriber.

The *data_on_readers()* listener is somewhat special, because even though there might be more specific *data_available()* listeners on the contained DataReaders, if this specific Subscriber listener is enabled, it will be called first. The reason for this is that it is still possible to call the specific *data_available()* listeners from within the *data_on_readers()* callback function, using the *DDS_Subscriber_notify_datareaders()* function. Refer to the DCPS specification for more details on this.

As stated before, care must be taken as to which functions are used within a listener callback. Since these callbacks occur from within a Technicolor DDS internal thread, it is very important that a callback function tries to do its work as fast as possible and doesn't go to sleep for a long time. If there is a lot of work to do, or if the workload might block for a long time it is necessary to dispatch this work to another thread, or just use the alternative WaitSet mechanism which was really meant for these use-cases.

Technicolor DDS ensures that the state of DDS is valid for a number of function calls. However, function calls that would block will return immediately with an error code, and function calls that delete entities may not be called.

Within these restrictions, it is possible, and this is effectively one of the more commonly used use-cases for listeners, to read data directly from a DataReader cache from within the listener callback, and even to send data on another DataWriter. For this, it is sufficient to call any of the *DDS_DataReader_read/take()* function calls, followed by an optional *DDS_DataReader_return_loan()* when finished with it.

WaitSets

WaitSets are a lot easier to work with and are a lot safer than listeners. In fact, a WaitSet can be seen as a kind of 'overloaded' *select()* or *poll()* function call. Both time-outs as well as waiting for data and waiting for other events, including notifications from other threads, can be handled properly using this mechanism.

WaitSets don't have any of the restrictions that are specific to listener callbacks. Everything can be done, i.e. every DDS function may be called, since everything occurs in the thread context of the application thread.

Handling WaitSets requires a bit more preparation work, however. Before a WaitSet can be used, it needs to be created properly, and this requires a number of function calls.

The typical workflow is as follows:

- Call *DDS_WaitSet__alloc()* to create a new WaitSet.
- Call *DDS_ConditionSeq__alloc()* or use *DDS_ConditionSeq__init()* on an already existing variable to create/initialize an empty Condition sequence. This is needed in order to add conditions that we are interested in to the WaitSet conditions.
- Add one or more conditions using the *DDS_WaitSet_attach_condition()* function. Conditions come in many flavors, such as:
 1. ReadConditions. These are DataReader conditions, where we specify a DataReader and the set of View, Instance and Sample states that we are interested in.
 2. QueryConditions. A superset of a ReadCondition, where in addition to the states, it is possible to specify an SQL-based expression on the data samples that we're interested in. DDS will only signal that the condition is ready when both the expression matches and the requested states are available.
 3. StatusConditions. Any state can be signalled as a condition on any entity, similar to the various Listener callbacks:

Table 4: Status bits

Status	Description	Entity
INCONSISTENT_TOPIC_STATUS	Topic is not compatible with an existing topic, or with a discovered topic with the same type.	Topic
OFFERED_DEADLINE_MISSED_STATUS	When Deadline QoS is enabled, Writer has missed a deadline for writing instance data.	Writer
REQUESTED_DEADLINE_MISSED_STATUS	idem, where Reader missed a deadline for reading instance data.	Reader
OFFERED_INCOMPATIBLE_QOS_STATUS	Offered Writer QoS settings are incompatible with Reader's QoS.	Writer
REQUESTED_INCOMPATIBLE_QOS_STATUS	Reader QoS settings are incompatible with Writer QoS.	Reader
SAMPLE_LOST_STATUS	Sample was lost due to a slow/late reader.	Reader

Status	Description	Entity
SAMPLE_REJECTED_STATUS	Sample was not accepted due to resource limits.	Reader
DATA_ON_READERS_STATUS	One of the DataReaders has data available.	Subscriber
DATA_AVAILABLE_STATUS	Data is available for reading.	Reader
LIVELINESS_LOST_STATUS	When Liveliness QoS is enabled, Writer was too slow with sending data.	Writer
LIVELINESS_CHANGED_STATUS	idem, where Reader is not responsive anymore.	Reader
PUBLICATION_MATCHED_STATUS	Writer has found a matching Reader.	Writer
SUBSCRIPTION_MATCHED_STATUS	Reader has found a matching Writer.	Reader

4. GuardConditions. A guardcondition can be used to wakeup a thread that is waiting for a number of Conditions in the WaitSet. An example use would be waiting for both one or more sockets being ready as well as on a number of DDS Conditions. The main DDS waiting thread would be blocked on the WaitSet, with an additional I/O monitoring thread waiting on *select()*, *read()* or *fread()* functions. When data is available via the I/O thread, this data/event could be appended to an input queue, followed by a *set_trigger_value()* function call to unblock the main DDS waiting thread.
- Call *DDS_WaitSet_wait()* with a suitable time-out. This function will return only if either of the given Conditions became active, or a time-out occurred. It is easy to see which Conditions were ready, since the list of ready Conditions is returned in a sequence of Conditions which can be checked individually, based on the type of condition. When processing is done, the *DDS_WaitSet_wait()* can be called again, for as many times as needed, until the program finishes.

Technicolor DDS memory handling

The way in which Technicolor DDS uses its data internally can be configured extensively in order optimize it for specific application components.

The following description will attempt to explain the mechanism that is used and will describe how to adapt DDS for the applications.

Memory pools

In order to speed up regular allocations and deallocations of memory blocks, a generic pool mechanism was introduced in Technicolor DDS.

For each type of often used memory, a pool is defined that is able to group a collection of preallocated memory chunks, chained together as a linked list of identical sized blocks.

Whenever a memory block is requested from a pool, the first available block is returned and the linked list contains one less element.

If a memory block is requested, but the pool has been drained completely, it is configurable as to what will happen then.

- If the pool limit was set to exactly the number of preallocated memory blocks for that pool, a NULL pointer is returned and an internal out-of-memory statistics counter is incremented.
- If the limit is higher than the number of preallocated memory blocks, a second list (the grow list) is checked for spare memory blocks. If there are memory blocks in that list, the new block will be taken from the grow list.
- If the grow list is empty and the memory constraints allow it, a new block will be allocated via a standard memory allocation function and that one will be returned.

When a memory block is released back to a pool from which it was allocated, it is either:

- Put back in front of the preallocated block linked list, if it was effectively a preallocated block, or
- it was allocated and the grow list is not full yet, it is added to the grow list, or
- it is released via a standard memory release function.

It is fully configurable by the user as to:


- How many memory blocks are preallocated in a pool.
- What the absolute limit of blocks is that may be in use for a pool.
- The maximum size of the grow list in relation to the preallocated list size (as a percentage, i.e. 0..infinity).


A number of statistics are kept per pool, as shown in the following table:

Table 5: Memory Pool statistics

Name	Description
MPUse	Maximum number of blocks that were ever allocated from the preallocated block list.
CPUse	Number of preallocated blocks that are currently in use (i.e. allocated).
MXUse	Maximum number of extra blocks that ever needed to be allocated in addition to the preallocated blocks.
CXUse	Number of extra blocks that are currently in use.
Alloc	Total number of extra block allocations that were done.

Name	Description
NoMem	Total number of times that a block was requested, but the pool was empty.

 **Note:** The names in the table correspond with the names that are displayed in the output of the 'spool' or 'spoola' debug commands.

 **Note:** If the -DFORCE_MALLOC compile option is given when Technicolor DDS is built, no preallocated blocks will be allocated and all allocations/frees will use the standard heap memory allocator, although it will still do accounting on the allocated pool blocks.

Configuring memory pools

The memory pools can be configured from application components using the following Technicolor DDS functions (located in `<dds/dds_aux.h>`):

```
DDS_ReturnCode_t DDS_get_default_pool_constraints (DDS_PoolConstraints *pars,
                                                    unsigned          max_factor,
                                                    unsigned          grow_factor);
```


This function will initialize the pars parameter, which is a structure containing all pool parameters, to suitable values in accordance with the maximum and grow factors.

The max_factor specifies the relation between the total number of allocatable blocks and the number of preallocated pool blocks as an extra percentage factor. I.e. if max_factor is set to 0, the maximum number of blocks will be equal to the number of preallocated blocks. If max_factor is set to 100, the maximum number of blocks will be twice the number of preallocated blocks. If set to ~0, the maximum number of blocks will not be limited and all available system memory may be used.

The grow_factor specifies the maximum length of the grow list for each of the memory pools as a percentage of the number of extra items that must not be returned.

```
DDS_ReturnCode_t DDS_set_pool_constraints (const DDS_PoolConstraints *pars);
```

This function should be used to apply the pool constraint parameters.

 **Note:** These functions should be used before the first DDS DomainParticipant is created, or they will not have any effect.

The following table summarizes the pool parameters that can be modified:

Table 6: Memory Pool Parameters

Name	Minimum	Maximum	Description
Domain	<na>	max_domains	Local DDS DomainParticipant
Subscriber	min_subscribers	max_subscribers	Local DDS Subscriber
Publisher	min_publishers	max_publishers	Local DDS Publisher
LocalReader	min_local_readers	max_local_readers	Local DDS DataReader
LocalWriter	min_local_writers	max_local_writers	Local DDS DataWriter
Topic	min_topics	max_topics	DDS Topic
FilteredTopic	min_filtered_topics	max_filtered_topics	DDS Filtered Topic
TopicType	min_topic_types	max_topic_types	DDS Topic Type.

Name	Minimum	Maximum	Description
ReaderProxy	min_reader_proxies	max_reader_proxies	Remote DDS DataReader Proxy
WriterProxy	min_writer_proxies	max_writer_proxies	Remote DDS DataWriter Proxy
RemoteParticipant	min_remote_participants	max_remote_participants	Discovered DDS DomainParticipant
RemoteReader	min_remote_readers	max_remote_readers	Discovered DDS DataReader
RemoteWriter	min_remote_writers	max_remote_writers	Discovered DDS DataWriter
PoolData	min_pool_data	max_pool_data	Data Buffers for sample storage
RxBuffer	<na>	max_rx_buffers	Receive buffers for RTPS
Change	min_changes	max_changes	Cache Change descriptor.
Instance	min_instances	max_instances	Cache Instance descriptor.
ApplicationSample	min_application_samples	max_application_samples	Application sample references.
LocalMatch	min_local_match	max_local_match	Local matching DataReader for a DataWriter.
CacheWaiter	min_cache_waiters	max_cache_waiters	Maximum threads waiting for a Cache.
CacheTransfer	min_cache_transfers	max_cache_transfers	Samples being queued between two local Data History Caches.
TimeFilter	min_time_filters	max_time_filters	Active Time Filter.
TimeInstance	min_time_instances	max_time_instances	Time Filter Instance.
String	min_strings	max_strings	Unique string descriptor
StringData	min_string_data	max_string_data	Unique string data area
Locator	min_locators	max_locators	Locator descriptor (see RTPS).
QoS	min_qos	max_qos	Unique DDS Quality of Service data.
List	min_lists	max_lists	Dynamic list
ListNode	min_list_nodes	max_list_nodes	Dynamic list node.
Timer	min_timers	max_timers	Timer
IPSocket	<na>	max_ip_sockets	IP socket
IPv4Address	<na>	max_ipv4_addresses	IPv4 address
IPv6Address	<na>	max_ipv6_addresses	IPv6 address
WaitSet	min_waitsets	max_waitsets	DDS WaitSet
StatusCondition	min_statusconditions	max_statusconditions	DDS StatusCondition
ReadCondition	min_readconditions	max_readconditions	DDS ReadCondition
QueryCondition	min_queryconditions	max_queryconditions	DDS QueryCondition
GuardCondition	min_guardconditions	max_guardconditions	DDS GuardCondition
Notification	min_notifications	max_notifications	Outstanding Listener notification.
TopicWait	min_topicwaits	max_topicwaits	Threads waiting for a Topic

Name	Minimum	Maximum	Description
Guard	min_guards	max_guards	Used if any of the Liveliness, Deadline, Lifespan and AutoPurge QoS parameters are active.
DynamicType	min_dyn_types	max_dyn_types	A DDS Dynamic type (see X-Types)
DynamicData	min_dyn_samples	max_dyn_samples	A DDS Dynamic sample (see X-Types)
DataHolder	min_data_holders	max_data_holders	A DDS DataHolder (Security: handshake/tokens)
Property	min_properties	max_properties	A DDS Property (Security)
BinaryProperty	min_binary_properties	max_binary_properties	A DDS Binary Property (Security)
HolderSequence	min_holder_sequences	max_holder_sequences	A DDS DataHolder contained sequence (Security)



Note: To make it a bit more easy for the user to configure these pool parameters, the Technicolor DDS debug shell can be used when an application has been running in steady-state for some time. The resulting `spool` (or `spoola`) debug command output can then be used to figure out how much memory was actually used in DDS itself by looking at the `MPUse` and `MXUse` parameters (see before). The total memory ever used will then be the sum of these two parameters.

Mobile device support

In order to use DDS effectively on low-power devices such as smartphones and tablets, it is mandatory to make sure that the CPU of these devices can be put in a deep sleep as often as possible.

Since these low-power requirements are actually in contradiction with the DDS peer-to-peer, always-on way of working, some extra mechanisms are needed: a suspend/resume mechanism, as well as a remote wakeup feature, handled by a dedicated background notification server.

Suspend/resume makes sure that when an application is about to enter a low-power state, DDS can detach itself from all external connections and stops all internal activity, effectively freezing itself until coming out of low-power. In this state, all entities such as DomainParticipants, Topics, Types, DataReaders, DataWriters (as well as their caches) and all discovered Participants, DataReaders and DataWriters are preserved but inactive, until activated again.

For the remote wakeup feature, an external server is required, which will be explained later.

Controlling DDS activities

The suspend/resume functionality is based on fine-grained control of internal DDS activities.

Following activities are defined (see `/dds/dds_aux.h` for the exact data definitions):

Table 7: DDS activities

Activity	Description
DDS_TIMER_ACTIVITY	Controls whether internal timers are running.
DDS_UDP_ACTIVITY	Whether the UDP/IP and UDP/IPv6 transport subsystems are active (sockets are present).
DDS_TCP_ACTIVITY	Whether the TCP/IP and TCP/IPv6 transport subsystems are active (sockets are present).
DDS_DEBUG_ACTIVITY	Whether the internal DDS debug shell stays active.



Note: An additional definition exists to specify all activities at once: `DDS_ALL_ACTIVITY`

Following functions are present to control these activities from within application programs:

```
void DDS_Activities_suspend (DDS_Activity activities);
void DDS_Activities_resume (DDS_Activity activities);
```

When an application needs to go to a deep sleep, it should call `DDS_Activities_suspend (DDS_ALL_ACTIVITY)` to force DDS into a frozen operating state. In this state, all DDS entities still exist, even the discovered participants are still present, but the low-level protocol sockets are closed, all timers are frozen and the debug shell is no longer working. The net result is that all datacaches stay populated with data and DDS can resume at any time.

When the application is ready to resume operation after waking up, it should call `DDS_Activities_resume (DDS_ALL_ACTIVITY)`.

This will recreate all underlying sockets, and continue all active timers, adjusting them for the amount of time that has actually passed. As a result of this, a number of possible things might happen:

- Either the participant hasn't timed out yet, and operation will seamlessly resume, or
- The participant has timed out and all associations with it will be removed completely.

Either way, operation of DDS resumes and other participants can now be discovered again.

The net result of DDS being frozen leads to this participant becoming timed out in the participants it was talking to. Of course, once it resumes operation, all peers will rapidly become available again.

In practice, no UDP or TCP sockets stay open while frozen, so the application cannot be awoken by other participants. Unless the Background Notification service is used, as discussed in the next topic.

The Background Notification Service

Sometimes it is important to be awoken automatically when there is urgent data available that the application *must* respond to.

This can be done by connecting to a remote Background Notification Server, which then takes over the reading of important notifiable topics (by creating its own DataReaders), and signalling, i.e. waking up the sleeping application.

In practice, the application needs to connect to this Background Notification Server and communicate to this server with a proprietary suspend/resume protocol, which can be configured in 2 ways:

- Either by using a dedicated TCP connection to this server, or
- Piggy-backing the suspend/resume protocol on the existing TCP control channel which is used for controlling DDS/TCP/IP data connections.

Which domain is used and whether it is a separate or in-band TCP connection, is configured via DDS parameters, as will be explained later.

Client side API functions

A number of functions are available on the client side to control the wakeup procedure.

In order to let the server know which topics are important to be notified, following functions are available:

```
DDS_Activities_notify (DomainId id, const char *topic, const char *type);
DDS_Activities_unnotify (DomainId id, const char *topic, const char *type);
```

The first function registers a Reader Topic at the remote notification server, and the second function reverses this (for cleanup purposes, for example).

Note that these functions optionally take wildcards in their topic name and type specification, making it much simpler for the app developers to specify which topics/types are important. Simple Unix wildcards are permitted, so '*' and '?' characters can be used in name specifications.

The normal behaviour when a client is awoken is to do an automatic, i.e. implicit DDS_Activities_resume().

If this is not useful, i.e. sometimes the client application wants to be notified but is not really waking up yet, i.e. it just wants to send a message to the user, who will then decide whether to wake up or to continue sleeping.

For this reason, the following function is defined:

```
typedef void (*DDS_Activities_on_wakeup) (
    const char *topic_name,
    const char *type_name,
    unsigned char id [12]
);

typedef void (*DDS_Activities_on_connected) (
    int fd,
    int connected
);

DDS_EXPORT void DDS_Activities_register (
```

```

DDS_Activities_on_wakeup wakeup_fct,
DDS_Activities_on_connected connect_fct
);

```

If the function is installed properly, i.e. with valid callback functions, then the `wakeup_fct` will be called when the Notification Service has received data on one of the proxy readers, and the `topic_name/type_name` and `id` will be filled with the Topic Name, Type Name and Participant identity of the writer.

Note that this `id` is sufficient to identify the peer participant, since it is actually the GUIDPrefix, which is known as the `DDS_BuiltinTopicKey_t` key in the `DDS_ParticipantBuiltinTopicData` data structure.

It becomes the responsibility of the application developer to unfreeze the DDS component when this `wakeup_fct` is called. If not called eventually, the DDS software will no longer function properly.

The `connect_fct` will be called whenever the TCP connection to the Background Notification Service is established, or when it is closed. The idea is to tell the Operating System in the callback that this TCP connection is allowed to stay up when sleeping. Some Operating Systems (like iOS) require this, or the connection will be broken by the OS itself when it forces the application in a deep sleep.

Server side API functions

It is possible to be notified on the Background Notification Server itself that a client is connected or disconnected and its state changes can be observed.

The following function allows this:

```

typedef enum {
    DDS_ACTIVITIES_CLIENT_ACTIVE,
    DDS_ACTIVITIES_CLIENT_SLEEPING,
    DDS_ACTIVITIES_CLIENT_DIED
} DDS_ActivitiesClientState;

typedef void (*DDS_Activities_on_client_change) (
    DDS_DomainId_t domain_id,
    DDS_BuiltinTopicKey_t *client_key,
    DDS_ActivitiesClientState state
);

DDS_EXPORT void DDS_Activities_client_info (
    DDS_Activities_on_client_change client_fct
);

```

The `client_fct` function, when installed is able to signal the following client states:

Table 8: Notification client states

State	Description
DDS_ACTIVITIES_CLIENT_ACTIVE	Either a new client is detected, or a transition from suspended to active occurred.
DDS_ACTIVITIES_CLIENT_SLEEPING	A connected client started sleeping, i.e. is in suspended state.
DDS_ACTIVITIES_CLIENT_DIED	A connected client became disconnected from the server. Either its battery is dead, or it is associated to a completely different network.

Configuration variables

Technicolor DDS can be configured extensibly. This can be done with system environment variables, configuration files and specific API functions.

Configuration mechanisms

A number of DDS parameters are present. Either DDS global parameters or specific group-related parameters can be configured.

A parameter name is typically formatted in configuration files and API function arguments as:

[<group>_]<name>

Where **<group>** is the name of the group and **<name>** is the parameter name.

When referred to a parameter as an environment variable, it gets an additional **TDDS_** prefix.

Example: **export TDDS_IP_MCAST_TTL=20** sets an environment variable specifying the **MCAST_TTL** parameter in the **IP** group.

DDS configuration files can be used to automatically configure Technicolor DDS on startup and have the typical .ini file syntax with group names between brackets and parameters in a group following the group name.

Example configuration file contents:

```
MAX_SAMPLE=100000      # Max. size of data samples

[TCP]
PORT=7400              # Set TCP server port.

[IPV6]
MODE=disabled          # No IPv6
```

These files can be located in different places:

- Either the location is specified by an application with a DDS API function as one of the first actions on startup, before calling any other DDS function:

```
void DDS_configuration_set (const char *filename);
```

- or it is specified with the **TDDS_CONFIG=<location>** environment variable.
- or one of the following standard filesystem locations are searched in the following order:
 - > The current directory: *tdds.conf*
 - > The user's home directory: *~/.tdds.conf*
 - > The system configuration directory: */etc/tdds.conf*

Configuration parameters can also be changed dynamically by applications with the following API functions:

```
DDS_ReturnCode_t DDS_parameter_set (
    const char *name,
    const char *value
);

DDS_ReturnCode_t DDS_parameter_unset (
    const char *name
);
```

```
const char *DDS_parameter_get (
    const char *name,
    char buffer [],
    size_t size
);
```

Configuration groups

The following configuration parameter groups are currently defined:

Table 9: Parameter groups

Name	Description
	DDS global parameters (no group prefix).
POOL	Memory pools configuration parameters.
RTPS	RTPS behaviour parameters.
IP	Internet Protocol parameters.
IPv6	IPv6 parameters
UDP	User Datagram Protocol (UDP) parameters.
TCP	Transmission Control Protocol (TCP) parameters.

Global configuration parameters

A number of global DDS parameters can be configured:

Table 10: Global parameters

Parameter name	Description	Type
NAME	Specifies the DDS name that will be used during discovery.	String
ENVIRONMENT	Identifies the DDS environment type. By default this is supposed to be the 'C'-environment, but can be changed for other programming languages. DDS will try to optimize data transfers between participants with the same environment and endianness, and will try to avoid marshalling in that case. See the environment definitions in <i>dds/dds_aux.h</i> for the numbers to use.	Number
PURGE_DELAY	Minimum time to delay to cleanup data when handling <i>*_delete_contained_entities()</i> . The default time is 50 us.	Number
SAMPLE_SIZE	Maximum size of data samples (in bytes). Default this is set to 256 Mbytes per sample.	Number
FORWARD	This is only meaningful when an application is compiled with -DDDS_FORWARD, to enable RTPS forwarding. It should be set to the actual forwarding mode (typically, either 0 to disable forwarding, which is the default, or 15 to enable full forwarding).	Number

Examples:

EXPORT TDDS_FORWARD=15

Enables forwarding.

EXPORT TDDS_NAME="TDDS Simple Application"

Sets the DDS application name to "TDDS Simple Application".

Pool configuration parameters

The memory pools can be configured, both with the `DDS_set_pool_constraints()` function (see: Configuring memory pools) as well as with Pool configuration parameters.

Following parameters are defined:

Table 11: Pool configuration parameters

Parameter name	Description	Type
POOL_DOMAINS	Local DDS DomainParticipant	Number
POOL_SUBSCRIBERS	DDS Subscriber	Range
POOL_PUBLISHERS	DDS Publisher	Range
POOL_READERS	Local DDS Reader	Range
POOL_WRITERS	Local DDS Writer	Range
POOL_TOPICS	DDS Topics	Range
POOL_FILTERED	DDS Filtered Topics	Range
POOL_TYPES	DDS Types	Range
POOL_RPROXIES	RTPS RemoteReader	Range
POOL_WPROXIES	RTPS RemoteWriter	Range
POOL_RPARTICIPANTS	Discovered DDS Participant	Range
POOL_RREADERS	Discovered DDS Reader	Range
POOL_RWRITERS	Discovered DDS Writer	Range
POOL_POOL_DATA	Sample data buffers	Range
POOL_RX_BUFFERS	Receive data buffers for RTPS	Number
POOL_CHANGES	Cache samples	Range
POOL_INSTANCES	Cache instances	Range
POOL_SAMPLES	Application sample references	Range
POOL_LOCAL_MATCH	Locally matching readers/writers	Range
POOL_CACHE_WAIT	Maximum threads waiting for a Cache.	Range
POOL_CACHE_XFER	Samples being queued between two local Data History Caches.	Range
POOL_TIME_FILTERS	Active Time Filter.	Range
POOL_TIME_INSTS	Time filter instance.	Range
POOL_STRINGS	Unique string descriptor	Range
POOL_STRING_DATA	Unique string data area size	Range

Parameter name	Description	Type
POOL_LOCATORS	Locator descriptor (see RTPS).	Range
POOL_QOS	Unique DDS Quality of Service data set.	Range
POOL_LISTS	Dynamic list	Range
POOL_LIST_NODES	Dynamic list node	Range
POOL_TIMERS	Timer	Range
POOL_WAITSETS	DDS WaitSet	Range
POOL_STATUSCONDS	DDS StatusCondition	Range
POOL_READCONDS	DDS ReadCondition	Range
POOL_QUERYCONDS	DDS QueryCondition	Range
POOL_GUARDCONDS	DDS GuardCondition	Range
POOL_NOTIFICATIONS	Outstanding Listener notification.	Range
POOL_TOPIC_WAITING	Threads waiting for a Topic	Range
POOL_GUARDS	Used if any of the Liveliness, Deadline, Lifespan and AutoPurge QoS parameters are active.	Range
POOL_DYN_TYPES	A DDS Dynamic type (see X-Types)	Range
POOL_DYN_SAMPLES	A DDS Dynamic sample (see X-Types)	Range
POOL_GROWTH	Default growth factor	Number

Example:

```
export TDDS_POOL_GROWTH=200
```

Sets the default pool growth parameter to 200%.

RTPS configuration parameters

The RTPS layer of DDS can be configured indirectly via a number of configuration parameters.



Note: Although the DDS RTPS specification combines both RTPS and RTPS transport mechanisms (UDP/IP), the configuration parameters for the transport layers are defined within the IP, UDP and TCP groups.

Following parameters are defined:

Table 12: RTPS parameters

Parameter name	Description	Type	Default
RTPS_MODE	Mode of operation.	Mode	Enabled
RTPS_SL_RETRIES	Number of stateless fast (200ms) retransmissions of SPDP messages. Used both when a participant starts up for the first time and when it sees another participant for the first time, in order to minimize the chance of packet loss that would necessitate having to wait for RTPS_RESEND_TIME before a new attempt to communicate is done.	Number	2

Parameter name	Description	Type	Default
RTPS_RESEND_TIME	Normal delay between successive SPDP messages (in seconds).	Number	10s
RTPS_HB_TIME	Heartbeat retransmission time (ms).	Number	100ms
RTPS_NACK_RESP_TIME	Nack response time (ms). The delay before Data/Gaps are sent upon receiving an AckNack requesting data.	Number	0
RTPS_NACK_SUPP_TIME	Nack suppression time (ms). The delay within AcnNack messages are silently discarded by a Writer.	Number	0
RTPS_LEASE_TIME	Default lease time before remote participants are forgotten.	Number	30s
RTPS_HB_RESP_TIME	Heartbeat response time (ms). The delay before an AckNack is set in response to a received Heartbeat message.	Number	0
RTPS_HB_SUPP_TIME	Heartbeat suppression time (ms). The delay within Heartbeat messages are silently discarded by a Reader when received too fast after a previous Heartbeat ⁹ .	Number	0
RTPS_MSG_SIZE	The maximum RTPS message length that determines whether RTPS submessages will be appended to a message.	Number	1452
RTPS_FRAG_SIZE	The default Fragment size for fragmented messages.	Number	8192
RTPS_FRAG_BURST	The maximum number of fragments that can be sent in a single burst.	Number	512
RTPS_FRAG_DELAY	The delay between bursts of fragments.	Number	0
RTPS_DTRACE	Default tracing mode of new created endpoints.	Number	0



Note: The RTPS_HB_TIME has a randomized exponential backoff mechanism behind it in order not to flood the network in case of unresponsive readers. This means that subsequent retransmission times can double up to 256 times the initial value (which is 25.6 seconds with the default settings).

Examples:

export TDDS_RTPS_FRAG_SIZE=1024

Sets the default fragment size to 1024 bytes.

IPv4 and IPv6 configuration parameters

A lot of parameters can be configured for the Internet transport mechanisms.

The following tables summarizes them:

Table 13: IPv4 and IPv6 configuration parameters

Parameter name	Description	Type	Default value
IP_SOCKETS	Defines the maximum number of IP and IPv6 sockets to be used simultaneously by	Number	1024

⁹ This parameter is currently not yet implemented in Technicolor DDS

Parameter name	Description	Type	Default value
	DDS. Once this maximum number is used, no more sockets will be opened.		
IP_MODE	Mode of operation of IPv4	Mode	Preferred
IP_SCOPE	Scope of own IPv4 addresses that may be used.	Range	Link .. Global
IP_INTF	The IP interface names that may be used for IPv4.	String (name{'.' name})	
IP_ADDRESS	The IPv4 addresses that may be used.	String (IPv4 address filter)	
IP_NETWORK	The IPv4 networks that may be used.	String (IPv4 network filter)	
IP_NO_MCAST	The IPv4 interfaces with the given addresses may not be used for multicast transmits. May be set to 'any' to disable IPv4 multicast altogether.	String (IPv4 address filter 'any')	
IP_MCAST_TTL	IPv4 multicast TTL value	Number	1
IP_MCAST_DEST	IPv4 multicast destination interface addresses	String (IPv4 address)	
IP_MCAST_SRC	IPv4 multicast source interface addresses	String (IPv4 address filter)	
IP_GROUP	IPv4 multicast destination address for DDS	String (IPv4 multicast address)	239.255.0.1
IPV6_MODE	Mode of operation of IPv6	Mode	Enabled
IPV6_SCOPE	Scope of own IPv6 addresses that may be used.	Range	Site-local
IPV6_INTF	The IP interface names that may be used for IPv6	String (name{'.' name})	
IPV6_MCAST_HOPS	Maximum number of IPv6 multicast hops.	Number	
IPV6_MCAST_INTF	Default IPv6 interface for outgoing multicast traffic.	String (interface name)	
IPV6_GROUP	IPv6 multicast destination address for DDS.	String (IPv6 multicast address)	ff03::80

As can be seen in the table, both IP and IPV6 groups can be configured.

IPv4 Multicast configuration parameters

Table 14: Multicast configuration parameters

DDS configuration parameter	Description
IP_MCAST_TTL	Can be used to change the default Time-To-Live value for multicast datagrams. If not specified, this TTL value is set to 1. Possible values: 1..255.
IP_MCAST_SRC	Specifies the IP addresses of the interfaces that may be used as a source interface for Multicast IP packets.
IP_MCAST_DEST	Specifies the IP address of the interface that must be used in order to route Multicast IP packets. This variable must be set in cases when there is no default gateway defined and there is no valid Multicast route present in the system. If set, it strictly overrides any automatic Multicast forwarding mechanisms that are present in the system, and can thus be used to restrict DDS access to other interfaces. It is allowed to specify the loopback address (e.g. 127.0.0.1) as a crude way of preventing Technicolor DDS from reaching other nodes outside the own node.
IP_NO_MCAST	This environment variable can be set to indicate which DDS peers do not accept Multicast traffic or on which networks it is not wise to use Multicast traffic (such as wireless networks). The argument can be a combination of IP address and IP network specifications similar to the specification that is used in TDDS_IP_ADDRESS and TDDS_IP_NETWORK (see further).
IPV6_MCAST_HOPS	Specifies the maximum number of hops for outgoing IPv6 multicast packets. It should be set to a value between 0 and 255.
IPV6_MCAST_INTERFACE	This option can be used to set the default outgoing interface for IPv6 multicast packets. Contrary to the TDDS_IP_MCAST_DEST environment variable, the argument should be a valid interface name (such as lo or eth0).
IPV6_GROUP	This variable should be used in case the default IPv6 Discovery multicast address (FF03::80) is not available. It should be set as a normal IPv6 address specification.

Examples:

export TDDS_IP_MCAST_TTL=2

Sets the Time-To-Live value for IPv4 Multicast datagrams to 2, so Multicast IP packets will be able to be routed 1 node further from the own node, but not further.

export TDDS_IP_MCAST_DEST=10.0.0.64

Set the Multicast route to the IP interface on which the address 10.0.0.64 is present, disregarding default IP Multicast routes. Note that this interface must be present for Multicast routing to be possible.

export TDDS_IP_NO_MCAST=10.0.0.32+31;144.11.0.12

Technicolor DDS will not send Multicast datagrams to DDS peers in the range 10.0.0.32..10.0.0.63, nor to a DDS peer located on 144.11.0.12.

The mechanism that Technicolor DDS uses to send data to multiple peers depends on this environment variable. If a datagram must be sent to a group of DDS peers with at least one peer in the given range, Technicolor DDS will choose to use specific Unicast packets to each individual destination within the group, even to those that do allow Multicast. If none of the destinations is mentioned in the variable, multicast will still be used.

IPv4 interface configuration parameters

The general behavior of Technicolor DDS as to the use of IP interface addresses is very simple. On startup, all active IP interfaces are queried, and if a valid IP address is found, this will be used. If no valid IP address is found, Technicolor DDS will wait until a valid IP address is assigned.¹⁰¹¹

In practice, this means that all IP addresses on a device will be used, even some that are really not meant to be used (such as *vmWare* addresses when a *vmWare* player is operating in bridged mode). This makes Technicolor DDS extremely open to the IP-network for all DDS domains, which might not always be a good idea from a security point of view.

To alleviate these issues somewhat, there are two ways to cope with this.

Firstly, it is possible to configure Technicolor DDS at build time to always filter away *vmWare* addresses. Since this is only useful in environments where we know that the *vmWare* addresses are effectively an issue, and is even wrong when *vmWare* is used in other modes (routing with or without NAT), this is rather crude and somewhat of a hack.

Another, second, mechanism was therefore introduced which is a lot more powerful and which allows the user to specify allowed addresses with a high degree of freedom.

This advanced mechanism is based on two parameters: **IP_ADDRESS** and **IP_NETWORK**, which together allow the user to create address and network blacklists as well as whitelists.

Table 15: IP addresses parameter

Parameter name	Description
IP_ADDRESS	This variable can be used to specify IP addresses which may (whitelist) or may not (blacklist) be used. Multiple addresses can be specified, and addresses can even differ per DDS domain.



Note: This is a filter specification which is applied to the found IP interface addresses in the system. It can not be used to override existing network addresses.

The EBNF syntax of the **IP_ADDRESS** parameter is as follows:

```
<tdds_addr_par> = "IP_ADDRESS=" <address_spec> { ';' <address_spec> }
<address_spec> = [<not>] [<domain_id> ':' ] <ipv4addr>
<not> = '~' | '^'
<domain_id> = <number> -- 0..230
<ipv4addr> = <ipv4octet> { '.' <ipv4octet> } -- 2..4 IPv4 octets
<ipv4octet> = <number> -- 0..255
```



Note: If less than 4 octets are specified, 0's will be inserted between the last octet and the octet before that one in order to get to a length of 4 octets. Example: 10.5 is converted to 10.0.0.5.

Examples:

```
export TDDS_IP_ADDRESS=10.100;120.10.6.90
```

¹⁰ If addresses change (due to a DHCP lease being lost or due to some explicit configuration changes), connectivity might be temporarily lost. However, once the IP interfaces are stable again, data transfers will resume. Technicolor DDS monitors IP address changes and automatically copes with changes when they occur.

¹¹ Only IPv4 addresses can be specified currently with these environment variables. Since public IPv6 addresses are not assigned by default to an IPv6 connection, and since Link-Local IPv6 addresses (which are generated automatically at startup) are not picked up by Technicolor DDS (see IPV6_SCOPE), this is not seen as a problem. When there is a clearly perceived need for IPv6 address selection, support for this might still be added by introducing new environment variables.

Only addresses 10.0.0.100 and 120.10.6.90 may be used by Technicolor DDS.

export TDDS_IP_ADDRESS=100:10.1.0.200

For domain 100, only address 10.1.0.200 may be used. Other domains can use any address.

export TDDS_IP_ADDRESS=^10:10.5;10.5

For domain 10, all addresses except 10.0.0.5 may be used. Other domains may only use 10.0.0.5.

Table 16: IP networks parameter


Parameter name	Description
IP_NETWORK	This variable can be used to define IP networks which may or may not be used. Similar to IP_ADDRESS, multiple specifications may be given, and the specification may differ per DDS domain.

The EBNF syntax of the **IP_NETWORK** parameter is as follows:

```

<tdds_net_par> = IP_NETWORK=<network_spec> { ';' <network_spec> }
<network_spec> = [<not>] [<domain_id> ':' ] <ipv4net>
<ipv4net> = <ipv4addr> |
              <ipv4addr> '/' <width> |      -- A single IPv4 address
              <ipv4addr> '+' <count>         -- A masked IPv4 subnet.
              -- Range of IPv4 addresses.
<width> = <number>                          -- 4..31 mask bits.
<count> = <number>                          -- 1..64K-1 extra addresses.

```

 **Note:** Contrary to IPv4 address specifications, specifying less than 4 octets in a network address will result in 0-byte padding after the last octet to get a length of 4 octets. Example: 10.5 is converted to 10.5.0.0.

Examples:

export TDDS_IP_NETWORK=10/8

Only addresses in the range 10.0.0.0 to 10.255.255.255 may be used.

export TDDS_IP_NETWORK=10/8;~10.1/16;50:100.10/14;10.1.1.20+24

All addresses in the range 10.0.0.0 to 10.255.255.255 that are not in the range 10.1.0.0 to 10.1.255.255, with the exception of 10.1.1.20 to 10.1.1.44 may be used for all domains. In addition, domain 50 may use addresses in the range 100.10.0.0 to 100.13.255.255.

Both **IP_ADDRESS** and **IP_NETWORK** can be used simultaneously, giving the opportunity to the user to specify both addresses and network ranges.

If none of these environment variables is specified, all valid interface addresses that are detected will be used.

The existence of address and/or network specification environment variables implies that filtering is used on all the detected interface IP addresses. This filtering uses a matching algorithm that operates in the following manner:

- If no match is found with any address or network specifications, the address will not be used.
- If a single match occurs, the address or network specification will either accept the address or reject it, depending on the presence of the <not> operator. If <not> was given in the specification, the address is rejected, otherwise it is accepted.
- If multiple matches occur, the first valid match with the longest prefix mask will be used.

Both blacklisting (using the <not> operator) and whitelisting (listing usable addresses/networks) can be used simultaneously as described above.

IPv4 and IPv6 address scopes

Table 17: Address Scope parameters

Parameter name	Description
IP_SCOPE	May be used to limit the scope of IPv4 addresses that will be used as source addresses and that will be announced to peer participants. If not specified, all source addresses with Link, Site, Organisation and Global scope will be used (see further). A valid scope expression should be specified as explained below.
IPV6_SCOPE	Specifies the scope of IPv6 addresses that will be used as source addresses and that will be announced to peer participants. If not specified all source addresses with Site, Organisation and Global scope will be used. A valid scope expression should be used, as explained below.

A scope expression has the form of: *<ScopeNumber>-<ScopeNumber>*, where *<ScopeNumber>* is a decimal number from 0 to 4, with the following significance of each number:

Table 18: Address scope numbers

Number	Scope	Description
0	Node	Node-local addresses are addresses that should never be used outside of a node as a source address (the loopback address).
1	Link	Link-local addresses are typically generated by a system. They should not be used outside the link on which they were defined. On IPv4: auto-ip addresses. On IPv6: the automatically generated link addresses.
2	Site	Site-local addresses should not be used outside a site. For IPv4, this includes the 192.168/16 and the 10/24 address ranges. For IPv6, the address scope is defined with the address.
3	Organisation	Organisation-local addresses should not be used outside a specific organisation. This is typically not known for IPv4 addresses, but can be specify with IPv6 addresses.
4	Global	Global or Public addresses can always be reached from anywhere over the Internet.

The scope expression thus makes it easy to restrict the scope of IPv4 or IPv6 addresses that may be used by Technicolor DDS.



Note: The expression only influences which local IPv4 (or IPv6) addresses will be announced to DDS peers and does not mean that addresses learned from other peers that are outside this scope will be ignored when received from peer DDS participants. It is not clear currently whether a mechanism for ignoring these remote addresses would be useful.

IPv4 and IPv6 mode of operation parameters.

Table 19: IPv4 and IPv6 mode parameters

Parameter name	Description
IP_MODE	Specifies the default mode of operation of IPv4. By default, IPv4 operation is enabled and preferred over IPv6. The argument should be set to one of the valid IP modes as explained below.

Parameter name	Description
IPV6_MODE	Specifies the default mode of operation of IPv6. By default, if IPv6 is enabled at compile time, it is enabled, but not preferred. The argument should be set to one of the valid IP modes.

Valid arguments for IP_MODE and IPV6_MODE are:

Table 20: IPv4 and IPv6 mode arguments

Mode	Description
DISABLE[D]	The specific IP transport is disabled and remains passive. IP addresses of the transport type will not be used and discovery of this transport will be inactive.
ENABLE[D]	The specific IP transport is enabled, but is not prioritized. This means that the transport will function and will announce its source addresses, but it will only be used if the other transport has no means of reaching a certain destination.
PREFER[RED]	The specific IP transport is enabled and preferred.



Note: If both IPv4 and IPv6 have priority, then IPv4 will still be used.



Note: Although described in uppercase in the mode table, alternatively all lower-case modes can be used as well.

The combination of mode settings will, in the current implementation, enable only one discovery protocol. If IPv6 is enabled in cooperation with IPv4, discovery will be done over IPv4, even though IPv6 might be preferred. Only if IPv4 is disabled will Discovery be performed over IPv6. The rationale here is that Discovery over IPv4 is clearly defined in the DDS specification as the default Discovery protocol.

UDP configuration parameters

The following parameters can be configured for UDP/IP:

Table 21: UDP configuration parameters

Parameter name	Description	Type	Default
UDP_MODE	Specifies whether UDP is enabled	Mode	enabled
UDP_PB	RTPS Port Base	Number	7400
UDP_DG	RTPS DomainId gain	Number	250
UDP_PG	RTPS ParticipantId gain	Number	2
UDP_D0	RTPS D0 parameter	Number	0
UDP_D1	RTPS D1 parameter	Number	10
UDP_D2	RTPS D2 parameter	Number	1
UDP_D3	RTPS D3 parameter	Number	11


The RTPS parameters are used to construct UDP port numbers, as described in the specification.


TCP configuration parameters

The following parameters can be configured for RTPS over TCP:

Table 22: TCP configuration parameters

Parameter name	Description	Type	Default
TCP_MODE	Specifies whether UDP is enabled	Mode	enabled
TCP_[SEC_]PORT[6]	TCP Server mode: server port.	Number	
TCP_[SEC_]SERVER[6]	TCP Client mode: server URL:port	String	
TCP_PUBLIC	Public TCP/IP address behind NAT-Firewall	String (IPv4 address[:'port'])	If not defined, port is assumed to be 7400
TCP_PRIVATE	Allow to connect to private IP addresses for TCP connectivity.	Mode	Disabled
TCP_PB	RTPS Port base.	Number	7400
TCP_DG	RTPS DomainId Gain.	Number	250
TCP_PG	RTPS ParticipantId Gain.	Number	2
TCP_D0	RTPS D0 parameter	Number	0
TCP_D1	RTPS D1 parameter	Number	10
TCP_D2	RTPS D2 parameter	Number	1
TCP_D3	RTPS D3 parameter	Number	11

 **Note:** The optional SEC_ prefix may be used to specify a secure TCP connection, i.e. by using DDS/TLS/TCP i.o. plain DDS/TCP.

 **Note:** The optional 6 suffix may be used to specify an IPv6 connection instead of the default IPv4.

To enable plain TCP server mode on a DDS application, it is sufficient to set the server port on which the TCP server will listen:

export TDDS_TCP_PORT=7400

On the other DDS application that needs to communicate in TCP client mode to the first DDS application, you only need to set the server URL and port:

export TDDS_TCP_SERVER=my-home.dynip.com:7400

It might be necessary to configure the home gateway for port forwarding through the NAT-firewall, in order to enable reachability to the in-LAN TCP-server DDS application.

An alternative scenario is connectivity to a public DDS TCP server from within a local LAN behind a NAT-firewall.

In this case it would be sufficient to enable the server on the public DDS host, and just configure the local LAN DDS with the URL of the public DDS TCP server. No port forwarding needs to be enabled in this case, since all TCP connections are setup from the client to the server.

In the Qeo case, there is already support for a forwarder in a LAN, using the TCP mechanism to configure a TCP server on the forwarder, automatically punching a hole in the local NAT-firewall, and publishing the TCP URL on a location server somewhere in the cloud.

This server can be a small always-on device, like a Raspberry Pi. Qeo supports this.

Variables for the Background Notification Service

As explained in the previous chapter, this Service requires some parameters in order to be configured properly.

Following parameters can be distinguished:

Table 23:

Environment variable	Description
BGNS_DOMAIN	Specifies the DDS domain that is used for the Notification service. Both client and server must set this correctly or the service will use id 0.
BGNS_[SEC]PORT[6]	Used by the server to specify the TCP port number on which the server is listening for incoming connections. If set for the first time, a server for the notification service is started. Either a valid port number should be specified for a dedicated TCP server, or @ may be specified, which indicates a connection that is in-band, i.e. multiplexed on the connection of the DDS/TCP control channel server.
BGNS_[SEC]SERVER[6]	Used by the client to specify the TCP server URL of the remote Notification service. If set to @ the client connection will be in-band, i.e. multiplexed on the client connection of the DDS/TCP control channel. Once this variable is set, the client will attempt to connect to the remote server.



Note: The optional SEC prefix indicates that a secure TCP connection is needed for the server.



Note: The optional 6 suffix indicates that the TCP connection is over IPv6 instead of over IP.

Examples

export TDDS_BGNS_DOMAIN=7

Set the DDS domain to 7 for the Background Notification service.

export TDDS_BGNS_PORT=5500

Start a dedicated Background Notification server on port 5500.

export TDDS_BGNS_PORT=@

Start a Background Notification server multiplexed with the DDS TCP control channel.

export TDDS_BGNS_SERVER=www.technicolor.com/bgns:5500

Start a dedicated Background Notification service to the server at the given URL on port 5500.

export TDDS_BGNS_SERVER=@

Start a Background Notification client multiplexed with the DDS TCP control channel.

Technicolor DDS components

Technicolor DDS is delivered as a single build component, consisting of a dynamically loadable shared library containing all core DDS functionality, with a number of prebuilt utility and demo programs on the supported platforms.

There is usually no need for changing the component. However, there are still some compile/link-time options available for 'tweaking' the behavior somewhat. It is even possible to link statically if there would be a need for it.

Overview

Shared library

The `libdds.so.0` library is a dynamically loadable shared library that can be built on Unix/Linux platforms to provide the core Technicolor DDS functionality. Applications link to the library via a `-ldds` linker options.

Central Discovery Daemon

The Central Discovery Daemon (`cdd`) is responsible for keeping per-node Discovery data central.

This prevents local components to all keep their own copy of Discovery data which would multiply the amount of memory needed for storing Discovery data with the number of components. The amount of memory used can be very high in cases where a large number of distinct topics are used.

Shapes demo

The shapes program (`shapes`) demonstrates DDS interoperability between different DDS vendors.

It is a technology demo, showing in real-time various shapes (Circle, Square and Triangle) that are moving around. In the program, the user can create new shapes, as well as subscribe to shapes.

A bandwidth test program

The bandwidth test program (`bw`) can be used to test DDS performance.

The program can be started either as a data producer or as a data consumer, or even both, so it can test performance between multiple DDS nodes and processes. It attempts to stress the system by sending bursts of data at controlled speeds, while verifying that data is successfully received at the remote end.

Latency test program

This latency test program (`latency`) is able to accurately measure the latency between two DDS components.

It sends a number of data samples that are looped by a remote entity and registers the time it took for arrival of the response. When the test is completed, a report is generated on the measured latency.

Chat program

The chat program (`chat`) is able to chat to any number of other chat clients directly over DDS.

Dump program

The dump program (`dump`) can display all topics it detects in a DDS domain and can even connect to any topic and display received topic data in text format.

It can be used to monitor the data that is being published in the domain.

IDL to TSM conversion utility

An IDL parser (`idl`) is available that is very useful for generation of TSM data. Manual conversion of IDL specifications are typically very error-prone. A tool that does this automatically is therefore extremely useful.

The DCPS functionality test program

This program, called `dcps`, is used internally to test out various functionalities of DDS between different application components. At compile time, lots of options are available for creating Readers and Writers with various kinds of QoS settings. As many instances of the program as necessary can be run simultaneously. It can be interacted with in a dynamic manner in order to start and stop data transfers from individual DataWriters and to monitor received data from individual DataReaders.

DCPS API test suite program

The API test program (`api`) tests out the full DDS DCPS API in order to verify that the Technicolor DDS implementation works correctly. It is a limited conformance test suite, but still manages to test all public DCPS API functions at least once. It is limited since it doesn't have full coverage of all DDS code, but runs very quickly and is used as one of the tests in the battery of tests that are performed when a new DDS build is done.

Central Discovery Daemon

DDS Discovery mechanism

DDS uses the Simple Discovery protocol (*SPDP/SEDP*) to propagate Participant, Topic and Endpoint (i.e. Reader/Writer) information to all other discovered DDS peers. This is needed, since a DDS-component can only 'match' with peer endpoints if this information is locally available, i.e. discovered.

Although this is perfectly ok if there are only a limited amount of DDS components in a device or if there is plenty of memory in a device when there are a lot of DDS components, this can become a real problem on low-memory footprint devices when a lot of topics need to be discovered.

The amount of memory that is used on each device for discovery data (e.g. meta-data) would in practice be: the number of components times the total number of topics and endpoints in the DDS domain times the memory needed to store each of the endpoints.

In Technicolor DDS, this memory usage is not very high, i.e. less than 80 bytes per discovered endpoint, which is a lot less than what other DDS'es require. However, even this small amount can quickly add up! Assuming 30 components and 2000 endpoints in a domain on a single device, for example, this would end up occupying: $30 \times 2000 \times 80 = 4.8\text{MB}$ for discovery data alone.

In practice, a component has typically no interest in the majority of the discovered data. Only the data that matches its own endpoints is really needed for correct DDS operation. For small components that have only a few matching endpoints, having to store all the unused discovery data can be a real burden.

A naive approach to filter discovered information might be to just discard everything that is of no interest yet, but this simply doesn't work. The discovered endpoint/topic info might be needed afterwards. This can happen if the order of endpoint creation is different between the component and its peers and the component is a bit slower (intentionally or not) than its peers. A match would then not happen and communication would not be possible.

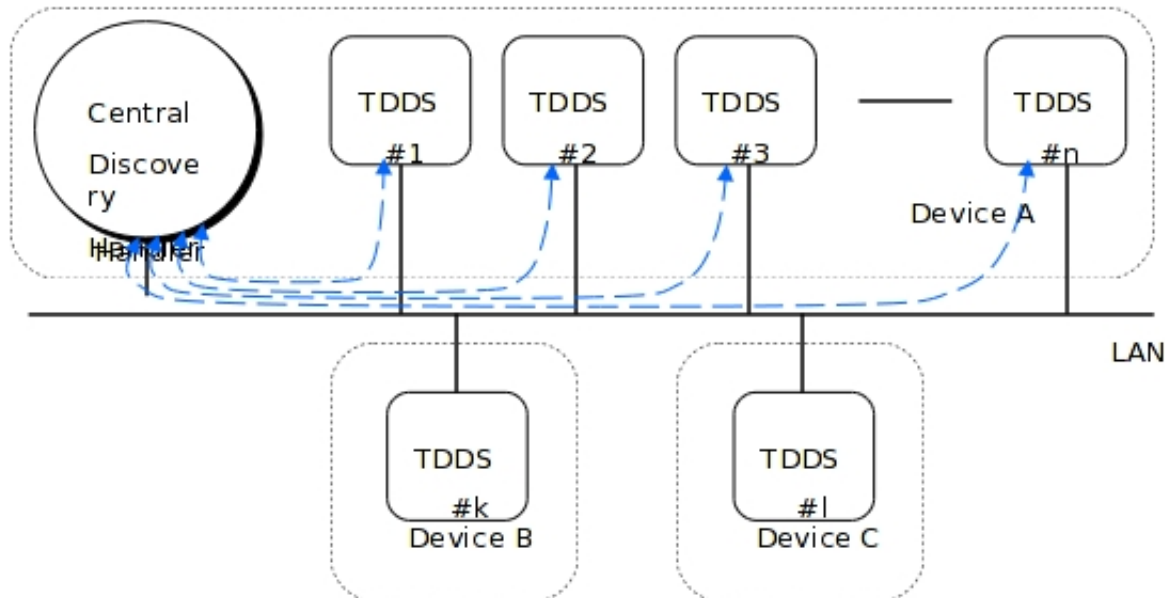
In fact, most embedded DDS's just disable the Simple Discovery protocol and use static (manually configured) discovery data. Although this can work with a limited amount of Topics/Endpoints, overall it is very cumbersome and also very error-prone, and it doesn't scale at all.

As shall be seen further, the idea of the Central Discovery handler solves the memory usage problems, while still allowing a fully dynamic Discovery mechanism.

Centralizing Discovery data

By centralizing the DDS Discovery meta-data on a device, a lot of memory can be saved. Using the numbers presented in the previous chapter and assuming that all Discovery data is needed only twice (not an unrealistic assumption), once in the Central Discovery handler and once distributed over the components of a device, the following numbers come up: $2 \times 2000 \times 80 = 320$ Kbytes, i.e. we can save almost 4.5 Mbytes of memory!

See below for an overview of the Central Discovery Daemon in action:



Of course, it is not all positive news since there are also a few disadvantages associated with a 'central' Discovery data store:

- It introduces a central point of failure in the device, which goes against the DDS 'philosophy'. If something happens with the central discovery daemon (CDD), crashes or otherwise, no more new matches will happen, although existing matches will continue to work.
- It slows down the time before discovery happens in the components, since there is more communication overhead and the matching decisions for late joining components are taken by the CDD instead of by the component itself.
- Some security issues might pop up if remote participants filter on participant USER_DATA to determine whether communication is allowed. The CDD is visible remotely as a separate participant and might be filtered out. Also, it might send matching information to local components, even though this was not allowed. Although the design of the daemon partially alleviates some of these issues, some problems might still occur. Some further study might be necessary to look at the various security issues.

CDD design principles

The following design principles are used:

- All communication between device components and the CDD will be over extra DDS builtin Central Discovery Endpoints, similar to the SPDP and SEDP endpoints.
- The Central Discovery endpoints will be invisible to external devices/components, since like builtin Discovery Endpoints they are not published.
- If a component starts up without a CDD being present (the CDD registers itself in the Technicolor DDS shared memory), or if specifically instructed not to use an existing CDD (via either an environment variable, or via a command-line flag), it will not use Centralized Discovery and just behave as a normal DDS component, storing all Discovery data.
- The CDD itself *will* be visible to external devices/components, but it will never announce Subscriptions or Publications over the normal SPDP/SEDP builtin endpoints.

- Since it is visible to the external world as an extra participant, the CDD captures all discovered information coming from components on the same device as well as from components on other devices, which it stores into its Discovery Data store.
- If the CDD detects that other participants on the *same* device publish new, i.e. not yet received Publications/Subscriptions, an additional action will happen; i.e. the Data store will be checked for new matches with existing endpoints in other components (local or remote). If a match is then found, this will be notified on one of the builtin Central Discovery endpoints.
- If the CDD detects that other participants on a *different* device publish new information, the CDD will not send anything to local components, since these still have their normal Discovery builtin endpoints and will react autonomously on the received information.

Starting `cdd`

The program, when used on Unix/Linux operating systems, daemonizes itself by default, and detaches itself from the standard I/O channels. It also then logs its operation to the Syslog service, and can be stopped by sending a Hangup (HUP) signal to it. This daemonizing behavior can be prevented with the `-D` option.

The program accepts a number of options that can be shown by typing `cdd -h`.

The following options can be given:

Table 24: `cdd` program options

Option	Description
<code>-i <domain></code>	Specifies the DDS Domain Identifier for which the Central Discovery function must be enabled. If not specified, Domain Identifier 0 is used.
<code>-f</code>	Continue if the Daemon is already running. This option is present to recover in cases where a previous instance of <code>cdd</code> stopped unexpectedly without proper cleanup. In case this happens, the shared memory region would wrongly indicate an active <code>cdd</code> , preventing the daemon from starting again.
<code>-r</code>	Reset the shared memory region before starting up. This can be used when DDS application components didn't complete properly, thereby messing up the count of active in-node DDS participants.
<code>-d</code>	Display/dump all Discovery data whenever changes are detected. Since this can lead to a lot of data being dumped, it is not advised to do this when daemonized (see <code>-D</code> option).
<code>-D</code>	Don't daemonize on startup, but behave as a normal user program. This option could be used while debugging, since <code>cdd</code> then displays all output on the normal terminal output channel, and its debug shell (when enabled) can be used directly from the command line.
<code>-p <name></code>	Specifies the name of a pipe on which <code>cdd</code> will send its <i>pid</i> when fully activated. This can be used by an init process to know the <i>pid</i> in order to be able to close <code>cdd</code> , but also to block until <code>cdd</code> is completely initialized before starting other DDS-based application processes.
<code>-v</code>	Verbose operation flag. This causes the program to log the main actions while it's active.
<code>-vv</code>	Very verbose operation flag. If enabled, <code>cdd</code> will give a lot of information while running, probably a lot more than desired. Don't use this while daemonized.
<code>-h</code>	Displays some information on program usage.

Shapes demo

Introduction

The shapes demo works by creating per basic shape form a DDS DataWriter, having a different instance per shape color (color = key), and periodically moving the shape in a specific direction (by publishing new shape data with different X and Y coordinates) until a wall is hit, whereafter the shape reverses direction.

Since only a limited number of shapes are defined, there are only up to 3 shape writers:

- square
- circle
- triangle

Subscribing to a shape form is done by creating a DDS DataReader, which then monitors the shape movements and draws the shape on the screen. There are also only three shape readers due to the limited number of shape forms.

The user is able to create new shapes dynamically via a builtin text GUI with context-specific menus.

It is also possible to create shape readers and shape writers from the command line at program startup.

A builtin debug shell can be enabled to monitor the DDS dynamic data usage, such as DataWriter and DataReader cache contents as well as other data while the program is active.

The shapes program can be started in as many terminals and on as many nodes as required. All shape writers will be able to communicate with every shape reader, whether locally, or over the network.

Starting shapes.

The shapes program can be run either using a full-screen text GUI interface, or in batch mode, where it just displays updates as messages on standard output.

The program accepts a number of options as shown by typing `shapes -h`.

The following options can be given:

Table 25: shapes program options

Option	Description
<code>-p <shape <color></code>	Publish a shape (square, circle or triangle) with the given color (red, green, yellow, blue, magenta, cyan).
<code>-s <shape></code>	Subscribe to a shape (square, circle or triangle).
<code>-t</code>	Trace messages, i.e. start in batch mode without a GUI.
<code>-b</code>	Black-on-white display instead of the default white-on-black.
<code>-x <strength></code>	Use exclusive ownership QoS with the given strength value.
<code>-d</code>	Start in debug mode.
<code>-i <domainid></code>	Specify the domain identifier (default = 0).
<code>-y</code>	Use a dynamic type definition instead of a static one.
<code>-v</code>	Verbose operation flag: logs overall functionality.
<code>-vv</code>	Very verbose flag: logs detailed functionality.
<code>-h</code>	Displays some information on program usage.

GUI mode

In GUI mode, the program starts up with the following screen:

```
Shapes demo (c) 2011, Technicolor  Publish  Subscribe  Command  Quit
#####
Subscriptions:                    #
#                                #
#                                #
#                                #
Publications:                     #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#                                #
#####
```

A limited number of commands are accepted here, as single characters:

When 'p' is typed, the Publish menu is displayed which gives the option to create a new Shapes instance writer. All local shape instance writers are listed under the Publications header of the GUI screen.

If 's' is typed, the Subscribe menu is displayed which allows to create a new Shapes reader, with the option of filtering received shapes based on shape coordinates. All local shape readers are listed under the Subscriptions header of the GUI screen.

If 'c' is typed, the shapes program enters the DDS Debug Shell, and displays the shell prompt. The user can then display DDS Debug commands and examine the debug responses. Published shapes are still updated, and received shapes are still monitored, so this allows an in depth examination of a running DDS application.

If 'q' is typed, the program exits immediately.

Publish menu

The Publish menu gives the user the opportunity to add new shape publications. It will query the shape type (Square, Circle or Triangle) and the shape color, and allows the user to select both parameters with a single keystroke from a menu of possible options.

It is entered via the 'p' command from the main GUI screen.

Following example creates a red square:

```
Shapes demo (c) 2011, Technicolor

  Publish
Shape type:
  Square
  Circle
  Triangle
or <esc> to return
Choice?
```

After 's' is typed, the following menu is presented:

```
Shapes demo (c) 2011, Technicolor

  Publish
```



```

Shape type = Square
Shape color:

  Red
  Green
  Yellow
  Blue
  Magenta
  Cyan

  or <esc> to return

  Choice?

```

After 'r' is typed a red square is added to the list of local publications, together with the current X and Y coordinates, which will constantly be updated while the program is running.

No square shape will be visible on the map if there is no shapes reader for square shapes. This needs to be created using the Subscribe menu.

Subscribe menu

The Subscribe menu allows the user to create new Shape readers, optionally specifying a coordinate-based shape data filter.

It is entered via the 's' command in the main GUI screen.

Following example creates a Square shapes reader:

Shapes demo (c) 2011, Technicolor

```

Subscribe

Shape type:

  Square
  Circle
  Triangle

  or <esc> to return

  Choice?

```

After 's' is typed, the following is shown:

Shapes demo (c) 2011, Technicolor

```

Subscribe

Shape type = Square
Filter:

  Yes
  No

  or <esc> to return

  Choice?

```

If the user types 'n', a non-filtered Shapes reader is created which is shown on the top GUI screen.

If 'y' is typed, the coordinates of the filter are requested. First the X-coordinate of the left corner of the visible square, next the X-coordinate of the right corner of the visible square, then the Y-coordinate of the top corner, followed by the Y-coordinate of the bottom corner:

Shapes demo (c) 2011, Technicolor

```

Subscribe

Shape type = Square
Filter = Yes

```

```
X1 (5..238) : 10
X2 (2..240) : 120
Y1 (5..254) : 50
Y2 (4..256) : 200
```

After the final coordinate is specified (end each coordinate input with <enter>), the filtered reader is created and the top GUI screen shows the visual shape rectangle:

```

Shapes demo (c) 2011, Technicolor
Publish Subscribe Command Quit
#####
Subscriptions:
#
#
Square
#
# # # # # # # #
Publications:
# #
#
Square RED 115 202
# #
#
# #
# #
# #
# #
# #
# #
#
# # # # # # # S#
#
#
#
#
#####

```

As can be seen from the screen display, whenever a reader is created and there are active shape writers, whether they are locally defined or generated by other DDS shapes writers, these shapes will be displayed on the top GUI screen in the map area.

Shapes are displayed in the map as a colored character, denoting the form of the shape ('S', 'C' or 'T'). The color of the square will correspond with the defined square color.

Command mode

The shapes command mode shows the builtin Debug shell as shown below:

```
Shapes command shell -- type 'h' or '?' for help.
>
```

Here, both the 'quit' command, as well as an additional 'menu' command is available to go back to the main GUI screen.

Refer to the Debug shell chapter for more details on how to use the command shell.

Bandwith program

Introduction

The DDS bandwidth test program (bw) is used to measure the bandwidth between two or more DDS components, where one is a sender of data and the others are receivers of the data.

It can be used to send timed data bursts at the Writer side in order to optimize the bandwidth to reach the highest throughput for a predetermined time period.

A large number of test parameters can be tweaked:

- Data is sent either reliably or in a best-effort manner.
- The size of the data chunks can be selected up to 32KBytes per data item.
- The number of data chunks in each burst is selectable.
- The delay between successive data bursts is selectable.
- The total delay of a test can be specified.

Test methodology

The same bandwidth program can be used either as a data producer or as a data consumer, or both, depending on the startup options given when the program is started.

It doesn't matter which of the two is started first, since the producer waits till all consumers are present before starting with the data transfers.

Separate command and status topics are used to:

1. The data producer commands the data consumers that a test is initiated or stopped.
2. The producer retrieves status information from the data consumers as to whether samples are missed.

Once all consumers are ready, the producer sends data of the specified size in bursts of a given length. When a burst is finished, the producer delays for some configured time, and while the test is not done yet, it repeats the burst.

When the test is finished, test results are displayed.

Starting the bw program

The bw program only runs from the command line and accepts a number of program options that can be shown by typing `bw -h`:

Table 26: bw program options

Option	Description
-o	Act as a data consumer.
-p	Act as a data producer.
-op po	Act as both, i.e. producer and consumer are in the same process, and exchange information locally. This is the default mode of operation if no options are given.
-n <nrdrs>	Number of data consumers listening for data. Default = 1.
-t <time>	Total test time in seconds. Default test time = 5 seconds.
-s <size>	Size of each data chunk in bytes. Default size = 1 byte. Sizes can be specified from 1 up to 32768 bytes.
-b <count>	Maximum number of data samples in each burst. Default burst size = 16. Burst size may be set from 1 up to 1000 data samples.
-d <usec>	Maximum delay in #seconds between each successive data burst. Default delay = 100#s.
-r	If given, data transfers will be reliable, i.e. data will be retransmitted when lost.
-q <history>	Number of samples to keep in the history. Default is all samples.
-g	Display a histogram of received samples in time.
-f	Repeat the tests forever.
-v	Verbose operation - logs overall functionality.
-vv	Very verbose - logs detailed functionality.
-h	Display some information on program usage.

Latency program

Introduction

The DDS latency test program (`latency`) is used to measure the latency, i.e. the end-to-end time delay that is consumed before samples are actually delivered from a data producer to a data consumer over DDS.

Test methodology

Testing latency is done using three distinct test entities:

1. A *Controller*, which controls the next two entities.
2. A *Looper*, which just loops back data to whoever sends data to it.
3. A *Generator*, which sends the samples to the looper entity.

The controller entity waits until all entities are present before starting a new test by sending a *START* command to the Generator.

The generator, on receiving this command, starts sending data samples periodically to the looper. When the looper receives a sample, it simply returns it.

When all samples are sent and looped, the generator stops and a latency report is generated.

All test entities can be combined in one program run by specifying appropriate options. Typical run examples are:

- `latency -ogl -s 1024` would run all functionalities in one program instance locally and test latency for a 1K sample size.
- `latency -og -n 100` on one node and `latency -l` on another node would test latency between two nodes, with the controller and generator combined in the first program, sending 100 samples.

Starting the latency program

The latency program only runs from the command line and accepts a number of program options that can be shown by typing `latency -h`:

Table 27: latency program options

Option	Description
<code>-o</code>	Controller mode (controls the other entities).
<code>-g</code>	Generator mode (sends data samples periodically and measures the latency from the returned, i.e. looped data sample).
<code>-l</code>	Looper mode (default setting if no mode is specified). In this mode, the program simply loops samples back to the generator.
<code>-n <nsamples></code>	Number of samples to sent for latency calculation. If not specified, 50 samples are sent.
<code>-s <size></code>	Data sample size in bytes. Default size = 1. Data sizes can range from 1 up to and including 32768.
<code>-d <msec></code>	Delay between each successive sample in milliseconds (default = 100ms).
<code>-r</code>	Use reliable transfers instead of best-effort operation.
<code>-q <history></code>	Number of samples to keep in the history. Default history is to keep all samples.
<code>-i</code>	Dump latency information for each sample. If not specified, only summary data containing minimum, maximum, average and median latency information is given.
<code>-v</code>	Verbose operation. Logs overall functionality.

Option	Description
-vv	Very verbose operation, logging detailed functionality.
-h	Display some information on program usage.

Chat program

Introduction

The chat program is an example of a typical publish-subscribe scenario. It allows multiple users to send information on multiple chatrooms, where everyone entering a chatroom can see all the messages that the other users are sending.

After starting the program, the user can just type messages. Everything typed, that is followed by an <enter> will be sent immediately. Messages that are received are displayed immediately.

In-chat commands

While chatting, the user can give in-chat commands. These all start with a '!' character, followed by a specific command:

Table 28: In-chat commands

Command	Description
!room <room_name>	Change the chatroom to the room with the given name.
!user <user_name>	Change the user name. By default, the user name is called .pid.%u where %u is the process id of the chat process.
!list	List all the attendees that are seen.
!info	Display the current chatroom and user name.
!busy	Set the user state as busy. The state will change back to normal once a new message is sent.
!away	Set the user state as not present. The state will change back to normal once a new message is sent.
!help	Display a summary of all the chat commands.
!q[uit]	Exit the chat program.
!!<command>	Execute a DDS Debug shell command.
!\$<command>	Execute an Operating System shell command.

Starting the chat program

Table 29: Chat program options

Option	Description
-r <chatroom>	Sets the chatroom name.
-u <name>	Sets the user name.
-i <domain_id>	Specify the domain identifier.
-v	Verbose operation. Logs overall functionality.

Option	Description
-vv	Very verbose operation, logging detailed functionality.
-h	Show the program options.

Dump program

Introduction

The dump program is a simple example on how to monitor a DDS domain and display any of its discovered topics.

It does this by listening to the builtin discovery readers and presenting a list of all the discovered topics to the user.

When the user decides to monitor any of these topics, the X-types functions can be used to create a new type support based on the discovered topic type. A subsequent DDS DataReader creation, referencing this type allows to read dynamic data samples that are published by the remote DDS DataWriters.

Since Dynamic data samples can be introspected using the X-types functions, they can be properly decoded and examined.

In-program commands

While the program is active, simple commands may be given to determine what should be monitored. These commands are an extension of the DDS Debug shell and provide for dump-specific extensions.

Following commands are available:

Table 30: Dump program commands

Command	Description
list	Lists all discovered domain topics.
auto <name_spec>	Automatically monitor the topics matching the <name_spec> argument, even if they haven't been discovered yet. Common wildcards like '*' and '?' are accepted.
monitor <name_spec>	Monitor the discovered topics matching the given argument.
ignore <name_spec>	Stop monitoring the discovered topics matching the argument.
trace	Toggle tracing of discovery traffic in real-time.

Starting the dump program

Following program arguments can be given:

Table 31: Dump program options

Option	Description
-p	Startup in paused state.
-d	Trace discovered data.
-i <num>	Domain id to use.
-a <name_spec>	Topics to be monitored automatically. This has the same effect as the in-program auto command.
-v	Verbose operation. Logs overall functionality.

Option	Description
-vv	Very verbose operation, logging detailed functionality.
-h	Show the program options.

IDL to TSM converter

Introduction

The TSM type that is used in Technicolor DDS was not really intended to be hand-written by users, although it is certainly possible to do so. It should really be seen as an intermediate type representation that is typically generated using special-purpose tools.

The complete core middleware development framework allows users to specify application components from an XML-definition file. This file contains sufficient information on the data types that will be used by applications. TSMs are then typically generated from this XML file using special core middleware tools.

In order to allow proper testing of the Technicolor DDS typecode framework, a need for a standalone TSM generation tool was identified.

The choice of a source type definition language fell on IDL, since it is a standardized data type mechanism that is source language independent, and using this would allow Technicolor DDS to take advantage of already existing IDL-based type definitions. An additional pleasant side-effect, of course, is that it is now possible to use Technicolor DDS as a standalone DDS middleware as well.

Using the `idlparser` program

The IDL parser program expects an IDL input, which is specified with the filename argument. If not specified, the parser will read from standard input.

If no output file is specified (using `-o`), the parser assumes `out.c` as the default output file, which will be automatically created or overwritten if it already exists.

The parser accepts a number of program options that can be seen by typing `idlparser -h`:

Table 32: `idlparser` program options

Option	Description
<code>-o <outfile></code>	Specifies the output file. If not specified, <code>out.c</code> is assumed.
<code>-t</code>	Generate test code for testing the generated type.
<code>-h</code>	Display program usage and options.

Example usage

Suppose the following input is used in `shapes.idl`:

```
struct ShapeType {
    string<128> color; //@Key
    long x;
    long y;
    long shapesize;
};
```

Then giving the command: `idlparser -o shapes.c shapes.idl`

Would result in the following *shapes.c* file contents:

```
#include <stdlib.h>
#include <stdio.h>
#include <dds/dds_dcps.h>

typedef struct _ShapeType_st {
    char color[128];
    int x;
    int y;
    int shapesize;
} ShapeType;

static DDS_TypeSupport_meta ShapeType_tsm [] = {
{CDR_TYPECODE_STRUCT    , TSMFLAG_KEY, "ShapeType", sizeof (struct
_ShapeType_st), 0, 4, 0, NULL},
{CDR_TYPECODE_CSTRING   , TSMFLAG_KEY, "color", 128, offsetof (struct
_ShapeType_st, color), 0, 0, NULL},
{CDR_TYPECODE_LONG      , 0, "x", 0, offsetof (struct _ShapeType_st, x), 0,
0, NULL},
{CDR_TYPECODE_LONG      , 0, "y", 0, offsetof (struct _ShapeType_st, y), 0,
0, NULL},
{CDR_TYPECODE_LONG      , 0, "shapesize", 0, offsetof (struct _ShapeType_st,
shapesize), 0, 0, NULL}
};
```

The DCPS functionality test program

Introduction

The dcps test program (dds) is used for testing of the Technicolor DDS middleware.

It allows to configure various different QoS combinations and allows the use of configurable data transfers between multiple DataReaders and DataWriters, locally, as well as on multiple hosts.

Program usage

By default, the program will behave as a single DataReader, waiting for data to arrive.

If started with the `-w` option, the program will create both a DataReader and a DataWriter, and unless the `-p` option is given, the program will start sending data immediately at a rate of 1 data sample per second.

Depending upon the `-v` option, sent and received data samples will be displayed on standard output.

Since the DDS Debug Shell is started by default, the user can examine DDS entities in detail while the test is running.

Some extra shell commands are present to allow some program-specific interactions:

- The 'pause' command allows to stop sending data samples on a DataWriter.
- The 'resume [*<nsamples>*]' command allows to continue sending data samples on the DataWriter, either continuously or for a given number of samples.
- The 'delay *<n>*' command updates the delay between the generation of successive data samples. Whereas this is by default set to 1 second, it can be set to values of 1 up to *<maxint>* milliseconds.
- The 'asp *<d>*' command can be used to invoke a `DDS_DomainParticipant_assert_liveliness()` function on the given domain.
- The 'ase *<endpoint>*' command can be used to invoke a `DDS_DataWriter_assert_liveliness()` function on the given endpoint handle.
- The 'quit' command can be used to exit the program, while properly cleaning up all resources.

Command line options

The program accepts a number of options as can be shown by typing `dds -h`.

The following options can be specified:

Table 33: dds program options

Option	Description
<code>-w</code>	Act as DataWriter, sending data samples.
<code>-r</code>	Act as a DataReader. This is the default, so doesn't need to be specified.
<code>-0 . . 3</code>	Number of local DataReaders in writer mode. Default is 1.
<code>-s <size></code>	Size of data samples sent in writer mode. By default, a 4-byte counter is sent, followed by string.
<code>-m <size></code>	Maximum data size of data samples. If set, the data sample sizes will increase from the <code>-s</code> argument and this one, and restart again.
<code>-a <offset></code>	Offset to add for each size increment.
<code>-n <count></code>	Maximum number of data samples to generate/receive. When the given amount is reached, the program stops sending/receiving or quits, depending on the <code>-q</code> option.
<code>-b <count></code>	Number of messages per burst.
<code>-q</code>	Quit when all packets are sent (writer side) or received (reader side).
<code>-f</code>	Flood mode when given. In flood mode, the writer doesn't wait, but sends as fast as possible.
<code>-d <msec></code>	Maximum delay to wait between sending samples in milliseconds. Default delay is 1 second (1000.ms). Minimum delay is 10ms.
<code>-p</code>	For a DataWriter: start in paused state (i.e. wait until a 'resume' command is entered).
<code>-t</code>	Trace transmitted/received messages.
<code>-i <domain></code>	Domain Id to use.
<code>-u <count></code>	Number of runs before exiting. Closes and reopens the domain for each subsequent run.
<code>-x</code>	Create an additional domain participant.
<code>-l</code>	Disable lower-layer RTPS functionality, so it can be used on the build farm as one of the tests in the automatic test suite.
<code>-g [<port>]</code>	Start a debug server.
<code>-y <filename></code>	Write received data to a file.
<code>-o <delay></code>	Delay before it starts publishing in writer mode.
<code>-v</code>	Verbose operation, logging overall functionality.
<code>-vv</code>	Very verbose operation, i.e. detailed functionality will be logged.
<code>-h</code>	Displays program usage.

The DCPS API test suite

Introduction

The DCPS API test suite (api) attempts to test all public DDS DCPS API functions with a number of different parameters.

It is intended to be run as a batch test, testing the total DCPS functionality, topic by topic. However, it is also possible to indicate specific topics in order to test only those specific functionalities.

Test methodology

Simple tests, if possible, just call functions within the to be tested topic.

More complex tests typically requires prologue and epilogue code to setup the edge conditions to be able to call a specific API function. Sometimes this requires quite an elaborate setup consisting of DomainParticipants, Topics, Publishers, Subscribers, DataReaders and DataWriters.

For the more complex tests, such as the discovery and data tests, up to 3 different domain participants within the same domain need to be created.

The advantage of this method is that everything runs in a single test program, while still enabling fundamental DDS functions such as RTPS messaging and state machines as well as active DDS Discovery mechanisms.

Starting the test suite

As mentioned above, the API test suite is divided in a number of individual test topics. These test topics can be shown with the `api -l` command, which then lists all the available topics:

Table 34: API test suite topics.

Topic	Description
extra	Auxiliary functions, e.g. Technicolor DDS-specific extensions to the DDS DCPS API.
seq	Tests for the Technicolor DDS-specific sequence functions.
type	Type registration tests.
factory	DomainParticipantFactory tests.
participant	DomainParticipant tests.
topic	Topic tests.
filter	Content-filtered topic tests.
multitopic	Multitopic tests ¹² .
pub	Publisher tests.
sub	Subscriber tests.
writer	DataWriter tests.
reader	DataReader tests.
data	Complex Data exchange tests.

If the program is started with as arguments a number of these topics, only the given topics will be tested.

The test suite accepts a number of startup options, that can be displayed with `api -h`.

Following options can be specified:

Table 35: API test suite options.

Option	Description
-l	Lists all test topics.

¹² Since Technicolor DDS doesn't support Multitopics yet, this just tests that the correct error codes are returned.

Option	Description
-d <msec>	Maximum delay to wait for responses (10..10000 milliseconds).
-n <num>	Total number of times that the tests need to run.
-r	Disable the lower-layer RTPS functionality.
-s	Enable the DDS Debug Shell.
-v	Verbose operation, logging overall functionality.
-vv	Very verbose, i.e. tries to log detailed test functionality.
-h	Display program usage.



Note: Since this program enforces severe restrictions on the number of discovered participants and discovered topics (it is a test program after all), it should only be used in clean environments, i.e. where there are no other participants active.

Library build options

The application and test programs described before each have a Makefile in their own directory and can be build as standalone programs.

The DDS library (i.e. *libdds.so.1*) should be build using the `Makefile_component` makefile in the `dds` root directory:

A large number of build options are available directly from the build system. These take the form of `dds_<option>=1`:

Table 36: DDS Library build options

DDS build option	Description
<code>dds_DEBUG=1</code>	This option automatically enables the debug shell using a builtin debug server that listens to port 3400 + <participant_id> for incoming Telnet connections. This is the default for debug builds but is normally disabled for release builds.
<code>dds_WERROR=1</code>	Compile everything with <code>-werror</code> .
<code>dds_SECURITY=1</code>	Enable DDS security software to be compiled in.
<code>dds_FORCE_MALLOC=1</code>	If not specified, the DDS library will use a builtin pool system (as described before), but it can be forced to automatically call <code>malloc/free</code> functions on pool allocations/releases for valgrind verification purposes. Using this option forces the use of <code>malloc/free</code> .
<code>dds_LOG_FILE=1</code>	The DDS log file will by default not be used in non-debug mode. Set this to keep using the DDS log file system.
<code>dds_QEO_SECURITY=1</code>	Enable the Qeo security plugin instead of the standard DDS security plugin.



Note: On Android and iOS devices, the <participant_id> number will always be 0, so port 3400 can always be used. On native systems (Linux, OSX, etc.) the <participant_id> will differ when multiple processes are using DDS. A command like `ss -t nmap` might indicate the correct debug port.

Using secure domains

All of the standalone programs can be run easily in any non-secure DDS domain. Some, however, have support for using secure domains as well, using extra command line options to specify the security parameters and using a special XML-based security policy file.

Most of the programs described in the previous chapters are security enabled:

- shapes
- bw
- latency
- chat
- fwd
- dcps
- cdd
- dump

To enable operation in a secure domain, the following options can be used:

Table 37: Security options for standalone programs

Option	Description
-e <name>	Security engine name (deprecated)
-c <path>	Path of the certificate to use
-k <path>	Path of the private key to use
-z <realm>	Optional Qeo realm name
-j <filename>	Filename of an XML-based security policy file.

If the -j option is not given, the *security.xml* file is searched in the current directory.

An example *security.xml* file that can be used for the dcps test program and which allows secure operation in domain 7 for topic *HelloWorld*, using DDS payload encryption:

```
<?xml version="1.0" encoding="utf-8"?>
<dds>
  <security>
    <domain id="7">
      <access>secret</access>
      <membership>exclusive</membership>
      <controlled>yes</controlled>
      <encryption>none</encryption>
      <transport protocol="UDP">
        <local>DDSSEC</local>
        <remote>DDSSEC</remote>
      </transport>
      <transport protocol="TCP">
        <local>TLS</local>
        <remote>TLS</remote>
      </transport>
      <topic>
        <name>*</name>
        <mode>restrict</mode>
        <controlled>no</controlled>
        <discovery>secure</discovery>
        <encryption>payload</encryption>
        <method>HMAC_SHA256</method>
      </topic>
      <partition>
        <name>*</name>
        <mode>open</mode>
      </partition>
    </domain>
  </security>
</dds>
```

```

<access>unclassified</access>
<membership>open</membership>
<transport>
  <local>open</local>
  <remote>open</remote>
</transport>
<topic>
  <name>*</name>
  <mode>open</mode>
</topic>
<partition>
  <name>*</name>
  <mode>open</mode>
</partition>
</domain>
<participant name="DCPS Test program">
  <access>secret</access>
  <domain id="7">
    <topic>
      <name>HelloWorld</name>
      <create>yes</create>
      <read>read_write</read>
      <controlled>yes</controlled>
      <discovery>secure</discovery>
      <encryption>payload</encryption>
      <method>AES128_HMAC_SHA1</method>
    </topic>
  </domain>
</participant>
<participant name="7fef45661ffa779f 9283 21a6">
  <access>secret</access>
  <domain id="7">
    <topic>
      <name>HelloWorld</name>
      <create>yes</create>
      <read>read_write</read>
    </topic>
  </domain>
</participant>
<participant name="7fef45661ffa779f 9050 21a5">
  <access>secret</access>
  <domain id="7">
    <topic>
      <name>HelloWorld</name>
      <create>yes</create>
      <read>read_write</read>
    </topic>
  </domain>
</participant>
</security>
</dds>

```

Note that there is still a need for both a valid certificate and a corresponding valid private key.

To start up the dcps test program with security enabled, the following command line could then be used:

```
./dds -k private_key.pem -c certificate.pem -i 7 -vwp
```

DDS Debug shell

Introduction

The DDS debug shell can be activated in a number of ways and allows users to examine/verify the internal workings of the Technicolor DDS middleware component.

It can be started, either directly from a specific application program linked to the Technicolor DDS middleware, or remotely via the DDS Debug server over a Telnet session. Both ways to connect to the shell result in approximately the same command interface with only minor differences.

Command Overview

The Debug shell, once activated will show a prompt and will wait for debug commands.

For example, when remotely connected over telnet, the following will be displayed:

```
Welcome to the TDDS Debug shell.
Type 'help' for more information on available commands.
>
```


Typing the 'help' command will show the list of all the DDS Debug commands that might be supported in an active shell:

```
TDDS Debug shell -- (c) Technicolor, 2012
Following commands are available:
  ssys          Display system-specific data.
  stimer        Display the timers.
  sstr          Display the string cache.
  spool         Display the pools.
  spoola        Display the pools (extended).
  scx [<cx>]    Display connections.
  scxa [<cx>]   Display connections (extended).
  scxq          Display queued connections.
  sloc          Display locators.
  sconfig       Display configuration data.
  sdomain <d> <lf> <rf> Display domain (d) info.
                  <lf> and <rf> are bitmaps for local/remote info.
                  1=Locator, 2=Builtin, 4=Endp, 8=Type, 10=Topic.
  sdisc         Display discovery info.
  sdisca        Display all discovery info (sdisc + endpoints)
  stypes [<name>] Display Type information.
  sqos          Display QoS parameters.
  stopic <d> [<name>] Display Topic information.
  sendpoints    Display the DCPS/RTPS Readers/Writers.
  scache <ep>   Display an RTPS Endpoint Cache.
  sdcache <ep>  Display a DCPS Endpoint Cache.
  qcach <ep> <query> Query cache data of the specified endpoint:
                    where: <ep>: endpoint, <query>: SQL Query string.
  sproxy <ep>   Display the Proxy contexts of an entry.
  rproxy [<ep>] Restart Proxy context.
  seqos <ep>    Display endpoint QoS parameters.
  scrypto <ep>  Display entity crypto parameters.
  sscache       Display security cache.
  rehs         Request a rehandshake.
  srx           Display the RTPS Receiver context.
  stx           Display the RTPS Transmitter context.
  sfd           Display the status of the file descriptors.
  pause        Pause traffic.
  resume [<n>]  Resume traffic [for <n> samples].
  delay <n>    Set sleep time in ms.
  asp <d>      Assert participant.
  ase <ep>     Assert writer endpoint.
  sdip         Display the dynamic IP state.
```

```

dtls          Display DTLS connection related info.
spdb          Display the policy database.
spv           Display the policy version numbers.
sfwd          Display the forwarder state.
ftrace <n>    Start forwarder tracing for <n> events.
dtrace [<t>]   Set/toggle default DDS tracing (t=0 or 1..15).
trace [<ep> [<t>]] Set/toggle tracing on endpoints.
               <t>: 1=Frame, 2=Signal, 4=State, 8=Timer.
ditrace [<t>] Set/toggle default IP tracing (0 or 1).
itrace [<h> [<t>]] Set/toggle tracing on IP connections (t:0/1).
cton          Enable the cyclic trace buffer.
ctoff         Disable the cyclic trace buffer.
ctclr         Clear the cyclic trace buffer.
ctmode <m>    Set tracing mode ('C'=cyclic/'S'=stop).
ctinfo        Display the state of the cyclic trace buffer.
ctaadd <pos> <name> <dlen> [<data>] <action> {<action>}
               Add an action list to a tracepoint.
ctaclr        Clear all action lists.
ctalist       Display all actions lists.
ctainfo [<id>] Display all id or all id:index values.
ctcounters    Display all counters.
ctclr         Clear all counters.
ctdump        Dump the cyclic trace buffer.
ctsave <filename> Save cyclic trace data to file.
profs         Display profile results.
cprofs <delay> <dur> Clear and restart profiling.
slstat        Display the status of all the locks.
d [<p> [<n>]]  Dump memory.
da [<p> [<n>]] Dump memory in ASCII.
db [<p> [<n>]] Dump memory in hex bytes.
ds [<p> [<n>]] Dump memory in hex 16-bit values.
dl [<p> [<n>]] Dump memory in hex 32-bit values.
dm [<p> [<n>]] Dump memory in mixed hex/ASCII.
indent <tab> <n> Set indent type (if <tab>=1: use TABs).
taflags <flags> Set type attribute display flags.
               <flags>: 1=header, 2=size, 4=elsize, 8=ofs.
server [<port>] Start debug server.
env           Display configuration data (same as sconf).
set <var> <value> Set the configuration variable to given value.
unset <var>     Unset the configuration variable.
suspend <value> Suspend with given mode.
activate <value> Activate with given mode.
help          Display general help.
quit          Quit main DDS program.
exit          Close remote connection.

```

 **Note:** This list includes all the commands that might be present in a DDS Debug enabled program. However, a lot of these commands depend on various compile options so only a subset of these is usually present. Only the commands that are available will be displayed when the help command is given.

Not all command characters need to be typed. Following table shows which command characters are significant and how to use them:

Table 38: Debug shell commands.

Command	Description	Compile option
#	Lines starting with '#' or empty lines are seen as comments and are simply ignored.	
<u>ssys</u>	Display system-specific data.	
<u>stimer</u>	Displays all active timers in the Technicolor DDS middleware.	
<u>sstr</u>	Displays the string cache.	
(<u>spool</u> <u>spoola</u>)	Display memory pool usage.	
(<u>scx</u> <u>scxa</u> <u>scxq</u>) [<cx>]	Display sockets that are in use.	RTPS_USED
<u>sloc</u>	Display the locators cache.	

Command	Description	Compile option
<u>sdomain</u> <d> <lf> <rf>	Display all domain data corresponding to the domain id and the local/remote flags.	
(<u>sdisc</u> <u>sdisca</u>)	Display the local and discovered entities.	
<u>stypes</u> [<name>]	Display registered types.	XTYPES_USED
<u>sqos</u>	Display the QoS cache.	
<u>stopic</u> <d> [<name>]	Display one or all Topics within the given domain.	
(<u>sendpoints</u> <u>sep</u>)	Display the Reader and Writer endpoints of all domains, local as well as discovered.	
(<u>scache</u> <u>sdcache</u>) <ep>	Display a local Reader or Writer cache.	
<u>qcache</u> <ep> [<query>]	Display local Reader or Writer cache data based on the given SQL-query expression.	
<u>sproxy</u> <ep>	Display the proxy Readers or proxy Writers corresponding with a local Writer or Reader.	RTPS_USED
<u>rproxy</u> <ep>	Restart a proxy reader or writer for test purposes.	RTPS_USED
<u>seqos</u> <ep>	Display the QoS parameters of any created entity.	
<u>scrypto</u> <ep>	Display entity crypto parameters.	DDS_NATIVE_SECURITY
<u>sscache</u>	Display security cache.	DDS_NATIVE_SECURITY
<u>rehs</u>	Request a rehandshake for test purposes.	
<u>srx</u>	Display the RTPS receive message context.	
<u>stx</u>	Display the RTPS transmitter context.	
<u>sfd</u>	Display the status of the file descriptors.	
<u>pause</u>	Pause sending/receiving data.	
<u>resume</u> [<n>]	Resume sending/receiving data for the given number of samples, or continuously if not specified.	
<u>delay</u> <n>	Sleep time between successive generated data samples (in milliseconds).	
<u>asp</u> <d>	Assert manual Participant Liveliness for the specified domain.	
<u>ase</u> <ep>	Assert Topic Liveliness for the specified endpoint.	
<u>sdip</u>	Display the dynamic IP address handler state.	DDS_DYN_IP
<u>dtls</u>	Display DTLS connection related info.	DDS_SECURITY
<u>spdb</u>	Display the policy database.	DDS_NATIVE_SECURITY
<u>spv</u>	Display the policy version numbers.	DDS_QEO_TYPES
<u>sfwd</u>	Display the forwarder state.	DDS_FORWARD
<u>ftrace</u> <n>	Start forwarder tracing for <n> events.	DDS_FORWARD
<u>dtrace</u> [<mode>]	Set the default tracing mode (default = no tracing) the the specified mode, or toggle the default tracing mode (none <-> all).	DDS_TRACE
<u>trace</u> [<ep> [<mode>]]	Set the tracing mode of an endpoint (default = no tracing) to the specified mode, or toggle the tracing mode of an	DDS_TRACE

Command	Description	Compile option
	endpoint (only <ep> specified) or toggle the tracing mode of all endpoints (if no arguments).	
<u>ditrace</u> [<t>]	Set/toggle default tracing of new IP connections.	MSG_TRACE
<u>itrace</u> [<h> [<t>]]	Set/toggle tracing on existing IP connections (t:0/1).	MSG_TRACE
<u>cton</u>	Start cyclic event tracing.	CTRACE_USED
<u>ctoff</u>	Stop cyclic event tracing.	CTRACE_USED
<u>ctclr</u>	Clear the cyclic event trace buffer.	CTRACE_USED
<u>ctmode</u> <ctmode>	Set cyclic trace mode to cyclic (<ctmode> = 'c') or one-shot (<ctmode> = 's').	CTRACE_USED
<u>ctinfo</u>	Display current cyclic trace settings and cyclic trace buffer filling level.	CTRACE_USED
<u>ctaadd</u> <pos> <name> <dlen> [<data>] <action> {<action>}	Add an action list to a tracepoint.	CTRACE_USED
<u>ctaclr</u>	Clear all action lists.	CTRACE_USED
<u>ctalist</u>	Display all actions lists.	CTRACE_USED
<u>ctainfo</u> [<id>]	Display all id or all id:index values.	CTRACE_USED
<u>ctcounters</u>	Display all counters.	CTRACE_USED
<u>ctcclr</u>	Clear all counters.	CTRACE_USED
<u>ctdump</u>	Dump the contents of the cyclic event trace buffer.	CTRACE_USED
<u>ctsav</u> e <filename>	Save the contents of the cyclic event trace buffer in a file.	CTRACE_USED
<u>profs</u>	Display profiling statistics information.	PROFILE
<u>cprof</u>	Clear profiling statistics.	PROFILE
<u>slstat</u>	Display the status of the DDS locks.	LOCK_TRACE
(<u>d</u> <u>da</u> <u>db</u> <u>ds</u> <u>dl</u> <u>dm</u>) [<p> [<n>]]	Dump memory commands to dump data in ASCII, hex (8/16/32 bits) or mixed hexadecimal/ASCII formats.	
<u>indent</u> <tab> <n>	Set indent type (if <tab>=1: use TABs).	
<u>taflags</u> <flags>	Set type attribute display flags.	
<u>server</u>	Start a DDS Debug server for remote access.	
<u>menu</u>	Start the program menu screen.	
(<u>sconfig</u> <u>env</u>)	Display configuration data.	
<u>set</u> <var> <value>	Set the configuration variable to given value.	
<u>unset</u> <var>	Unset the configuration variable.	
<u>suspend</u> <value>	Suspend with given mode.	
<u>activate</u> <value>	Activate with given mode.	
(<u>help</u> <u>H</u> <u>?</u>)	Display the debug commands.	
<u>quit</u>	Quit either the DDS Debug shell or the main program depending on the current context.	
<u>exit</u>	Exit from a remote debug shell connection.	

Resource display commands

The following DDS resources can be displayed:

- Timers. Use the `stimer` command to display the currently active timers.
- Strings. The `sstring` command can be used to display the string pool contents.
- File descriptors. The `sfd` command displays the currently used file descriptors.
- Memory pools. Use the `spool` or `spoola` command to display the current memory usage.
- IPv4 and IPv6 addresses. The `sdip` command can be used to display the currently assigned IP addresses.
- IPv4 and IPv6 connections. The `scx` command is useful to retrieve some socket statistics for a number of different IP connections (UDP, TCP, DTLS/UDP and TLS/TCP). For DTLS/UDP connections, the `dtls` command gives more detailed information.
- Locators. RTPS locators are used extensively, and the contents of the locators pool can be dumped with the `sloc` command.
- Domain entities. The domain entities can be displayed using the `sdomain`, `sdisc` and `sdisca` commands.
- Types. Registered types can be displayed with the `stypes` command.
- QoS pools. The QoS parameters are kept in a central QoS pool, which can be displayed with the `sqos` command.
- QoS parameters. Individual QoS parameters for any DDS entity that has it can be displayed with the `seqos` command.
- Topics. Topics can be displayed in detail using the `stopic` command.
- Endpoints. The created DataReader and DataWriter endpoints can be listed with the `sendpoints` and `sep` commands.
- History cache. The DDS history cache of both DataReader and DataWriter endpoints can be displayed with the `scache` and `sdcache` commands. The actual contents of a history cache can be queried with the `qcachec` command.
- RTPS Proxy contexts. RTPS Proxy Reader and Proxy Writer contexts can be shown with the `sproxy` command.
- RTPS Receiver context. The RTPS Receiver context as well as some statistics and a small error history can be dumped with the `srx` command.
- RTPS Transmitter context. Similar to `srx`, the `stx` command shows the RTPS Transmitter context.
- Participant and Endpoint crypto data. Crypto data can be displayed with the `scrypto` and `sscachec` commands.
- Security policies. The default security plugin allows to display its current configuration with the `spdb` and `spv` commands.

Some example resource display commands (taken from various running shapes and dds programs) will be given, together with an explanation of what is displayed in response.

Timers

Using the `stimer` command shows the following information:

```
>stimer
DDS running for 116.56s
# of timers started/stopped/active = 5387/2469/13
# of timeouts = 2905, busy-locks = 0
Now = 1351522, Timers:
0x1992b40: time = 1351529 (0.07s), user = 26815056, tcbf = 0x406cab, name = 'P:Square(Yellow)'
0x1992cb0: time = 1351530 (0.08s), user = 26815424, tcbf = 0x406cab, name = 'P:Circle(Blue)'
0x196ccd0: time = 1351530 (0.08s), user = 26814800, tcbf = 0x406cab, name = 'P:Square(Red)'
0x753320: time = 1351571 (0.49s), user = 0, tcbf = 0x40c4f5, name = 'SharedMemory'
0x1977310: time = 1351944 (4.22s), user = 26777384, tcbf = 0x4a27f7, name = 'RTPS-Resend'
0x7f83bc001150: time = 1352011 (4.89s), user = 140203771564752, tcbf = 0x40782c, name = 'S:Square(MAGENTA)'
0x7f83bc0010c0: time = 1352017 (4.95s), user = 140203771563984, tcbf = 0x40782c, name = 'S:Square(GREEN)'
0x7f83bc000d70: time = 1352017 (4.95s), user = 140203771563136, tcbf = 0x40782c, name = 'S:Square(YELLOW)'
0x7f83bc001400: time = 1352018 (4.96s), user = 140203771565504, tcbf = 0x40782c, name = 'S:Circle(CYAN)'
0x7f83bc000f30: time = 1352018 (4.96s), user = 140203771563584, tcbf = 0x40782c, name = 'S:Circle(BLUE)'
0x7f83bc000a70: time = 1352018 (4.96s), user = 140203771562880, tcbf = 0x40782c, name = 'S:Square(RED)'
0x1984110: time = 1355946 (44.24s), user = 26754896, tcbf = 0x47f08b, name = 'DiscParticipant'
0x1983e18: time = 1355946 (44.24s), user = 26754136, tcbf = 0x47f08b, name = 'DiscParticipant'
>
```

Strings

The `sstring` command displays the string pool:

```
1: 'string_128_character'*1 'PublicationBuiltinTopicData'*1
3: 'data'*1 'array_4_octet'*1
7: 05 00 00 00 43 59 41 4e - 00          ....CYAN.
*1
8: 'destination_participant_key'*1
9: 04 00 00 00 52 45 44 00          ....RED.
*2 'Int64'*1
11: 'Square'*1
13: 'SEDPbuiltinPublicationsReader'*1 'ParticipantStatelessMessage'*1 'ParticipantMessageData'*2
15: 'longlongs_value'*1 'class_id'*1
.
.
96: 06 00 00 00 47 52 45 45 - 4e 00          ....GREEN.
*1
```

The first number is the hash index, which is followed by the string as either ASCII characters or binary data which is displayed as mixed dump of hexadecimal numbers and ASCII characters.

As can be observed from this sample, not only ASCII strings, but binary strings as well are stored in the string pool.

The number of times a string is used in the system is denoted with the `*n` which follows each unique string. Here, the `n` parameter specifies the reference count of the string.

Pools

The `spool` command shows the following info:

```
>spool
Name      PSize  BSize  Rsrvd   Max  MPUse  CPUse  MXUse  CXUse  Alloc  NoMem
-----
List:
SLIST      1152   144    8       *    0     0     0     0     0     0
SL0        6144   32   192     *   58    58     0     0     0     0
SL1       2304   48    48     *   12    12     0     0     0     0
SL2        768   64    12     *    0     0     0     0     0     0
SL3        240   80     3     *    0     0     0     0     0     0
SL4         96   96     1     *    0     0     0     0     0     0
SL5        112  112     1     *    0     0     0     0     0     0
SL6        128  128     1     *    0     0     0     0     0     0
SL7        144  144     1     *    0     0     0     0     0     0
String:
STRING     1536   24    64     *   64    64    14    14    14     0
STR_REF    1184   16    74     *   74    74     4     4     4     0
Timer:
TIMER      1536   48    32     *    7     1     0     0     0     0
Buffers:
DATA_BUF(0) 8216  8216     1     *    0     0     0     0     0     0
DATA_BUF(1) 4120  4120     1     *    0     0     0     0     0     0
DATA_BUF(2) 2072  2072     1     *    0     0     0     0     0     0
DATA_BUF(3) 2096  1048     2     *    0     0     0     0     0     0
DATA_BUF(4) 2144   536     4     *    2     0     0     0     0     0
DATA_BUF(5) 2240   280     8     *    8     7     1     0    716     0
DATA_BUF(6) 2432   152    16     *    0     0     0     0     0     0
DATA_BUF(7) 2816    88    32     *    0     0     0     0     0     0
Locators:
LOCREF     1184   16    74     *   74    70     1     1     1     0
LOCATOR    1280   40    32     *    9     9     0     0     0     0
QoS:
QOS_REF     1344   16    84     *    3     3     0     0     0     0
QOS_DATA    1792  112    16     *    3     3     0     0     0     0
Cache:
HIST_CACHE  5200   200    26     *   12    12     0     0     0     0
CHANGE     4096   64    64     *   19    15     0     0     0     0
INSTANCE   4864  152    32     *   15    14     0     0     0     0
CCREF      4096   32   128     *   30    28     0     0     0     0
CREF        32    16     2     *    2     2     0     0     0     0
CWAIT      208   104     2     *    0     0     0     0     0     0
CXFER       144   72     2     *    0     0     0     0     0     0
XFLIST      112   56     2     *    0     0     0     0     0     0
```

```

FILTER      240    120    2    *    0    0    0    0    0    0
FINST       128     64    2    *    0    0    0    0    0    0
Domain:
DOMAIN      1688   1688    1    5    1    1    0    0    0    0
DPARTICIPANT 6080   760    8    *    2    2    0    0    0    0
TYPE        320    32    10   *    5    5    0    0    0    0
TOPIC       2304   144   16   *   10   10    0    0    0    0
FILTER_TOPIC 240    240    1    *    0    0    0    0    0    0
PUBLISHER   528    264    2    *    2    2    0    0    0    0
SUBSCRIBER   544    272    2    *    2    2    0    0    0    0
WRITER      1920   240    8    *    6    6    0    0    0    0
READER      4224   352   12   *    6    6    0    0    0    0
DWRITER     1920    80   24   *    9    9    0    0    0    0
DREADER     2304    96   24   *   10   10    0    0    0    0
GUARD       224    56    4    *    0    0    0    0    0    0
PREFIX      320    40    8    *    1    0    0    0    0    0
DDS:
NOTIFICATION 192    24    8    *    2    0    0    0    0    0
WS_DEFERED   24    24    1    *    0    0    0    0    0    0
CFG_UPDATE   24    24    1    *    0    0    0    0    0    0
DCPS:
SAMPLE_INFO  1792    56   32   *    2    2    0    0    0    0
WAITSET      240   120    2    *    0    0    0    0    0    0
STATUS_COND  112    56    2    *    0    0    0    0    0    0
READ_COND    128    64    2    *    0    0    0    0    0    0
QUERY_COND   304   152    2    *    0    0    0    0    0    0
GUARD_COND    96    48    2    *    0    0    0    0    0    0
TOPIC_WAIT   160    80    2    *    0    0    0    0    0    0
RTPS:
READER      1080    72   15   *    6    6    0    0    0    0
WRITER      968    88   11   *    6    6    0    0    0    0
REM_READER  2320   232   10   *   10   10    2    2    2    0
REM_WRITER  2080   208   10   *    9    9    0    0    0    0
CCREF       3072    48   64   *   10    2    0    0    0    0
MSG_BUF     7168   224   32   *    5    0    0    0    0    0
MSG_ELEM_BUF 10240   160   64   *    5    0    0    0    0    0
MSG_REF      256    16   16   *    0    0    0    0    0    0
IP:
IP_CX       246000   240  1025  1025    5    5    0    0    0    0
XTYPES:
DYNTYPE     192    24    8    *    0    0    0    0    0    0
DYNDREF      64    16    4    *    0    0    0    0    0    0
SECURITY:
DATA HOLDER  224    56    4    *    0    0    0    0    0    0
PROPERTY     64    16    4    *    0    0    0    0    0    0
BIN_PROPERTY 128    32    4    *    0    0    0    0    0    0
HOLDER_SEQ   96    24    4    *    0    0    0    0    0    0

Pool/max/xmax/used/xused memory = 378128/49968/1160/45400/880 bytes (737 mallocs) (13%/12%)
Dynamically allocated: 39244 bytes.
malloc statistics: 163 blocks, 53280 bytes, 0 failures.
realloc statistics: 4 blocks, 320 bytes, 0 failures.
free statistics: 94 blocks, 14356 bytes.
Dynamic pool block stats: <=64K - max/used/msize/size: 1/0/0/216
                        >64K - max/used           : 0/0
Total heap memory: 379288 bytes.

```

The display first shows some columns, followed by summary information.

The `spoola` command is similar, but gives two additional address columns, the pool block address and the first free block address.

The first column shows either the origin (ex: RTPS:) or the name of the memory pool:

Table 39: Types of memory blocks

Name	Description
SLIST	Skiplist list.
SL0..7	Skiplist node containing i*2 forward pointers. Ex.: SL3 contains 6 forward pointers.
STRING	String descriptor. Actual string data is stored in a separate string heap.
STR_REF	Reference to a string descriptor.

Name	Description
TIMER	Timer node.
DATA_BUF(0..7)	Data buffer pools having room for 64 up to 8K bytes as subsequent powers of 2, e.g. 64, 128, 256, etc. Buffers have a reference count and can be referenced from multiple readers and writers.
LOCREF	Reference to a locator node.
LOCATOR	A locator node contains locator storage and a reference count.
QOS_REF	Reference to a QoS node.
QOS_DATA	A QoS node contains Quality of Service parameters in a unified format, which is suitable for DataWriters, DataReaders and Topics either locally or remote.
HIST_CACHE	History cache descriptor as used in both DataWriters and DataReaders. It describes all data-related handling, and contains a list of samples, and possibly a list of instances.
CHANGE	A Change descriptor typically contains a reference to actual change data (handle, data stored in a DATA_BUF or allocated directly).
INSTANCE	Each instance, e.g. key-dependent data is a descriptor that contains a list of samples as well as various state information.
CCREF	The Change reference descriptor is used while queueing sample data. Keyed samples typically have two change references. Non-keyed is typically a single change reference.
CREF	The Reader Cache reference is used in order to chain matched local DataReaders for a local DataWriter.
CWAIT	The Cache Wait context is used when an application needs to wait until some condition needs to be fulfilled, such as room for writing a new sample, wait for acknowledgements, etc.
CXFER	Samples that need to be sent to local DataReader caches from a local DataWriter cache, but which can not be transferred yet, due to the destination cache having no room, are queued using Cache Transfer descriptors.
XFLIST	A list of Cache Transfer descriptors to be received when the local DataReader cache accepts new samples.
FILTER	The Time-Based filter descriptor is used for a local Reader cache to limit the receive rate from a specific Writer.
FINST	This a per-instance descriptor that contains a delayed sample for reception.
DOMAIN	A complete DomainParticipant descriptor.
DPARTICIPANT	A Discovered, i.e. remote DomainParticipant.
TYPE	Type descriptor.
TOPIC	Topic descriptor.
FILTER_TOPIC	A Filtered topic is derived from a normal topic.
PUBLISHER	Publisher descriptor.
SUBSCRIBER	Subscriber descriptor.
WRITER	A DDS DataWriter, suitable for local and remote use, referring to a history cache, optionally extended with an RTPS Writer context.
READER	A DDS DataReader, suitable for local and remote use, referring to a history cache, optionally extended with an RTPS Reader context.
DWRITER	A Discovered DataWriter.

Name	Description
DREADER	A Discovered DataReader.
GUARD	Guard contexts are used to implement various QoS parameters, such as Liveliness, Deadline, Lifespan and DataReader Lifetime QoS.
PREFIX	Prefix cache entry to associate IP addresses to GUID prefixes for locality determination.
NOTIFICATION	Decoupled listener notification context. Used when any listener needs to be called.
WS_DEFERED	Deferred Waitset notification context. Used when a Waitset is blocked and needs to be woken up.
CFG_UPDATE	When users update configuration parameters dynamically, this context is used to queue the update to the core DDS thread.
SAMPLE_INFO	The SampleInfo structure is allocated when using any of the various DDS_DataReader_read/take() functions, until a DDS_DataReader_return_loan() is called.
WAITSET	A WaitSet context.
STATUS_COND	A StatusCondition descriptor.
READ_COND	A ReadCondition descriptor.
QUERY_COND	A QueryCondition descriptor.
GUARD_COND	A GuardCondition descriptor.
TOPIC_WAIT	Whenever the application needs to wait on a remotely defined topic via DDS_DomainParticipant_find_topic(), this context is used.
RTPS::READER	When a local DataReader is matched for the first time with a remote DataWriter, the RTPS Reader context is added to the DataReader in order to manage the RTPS Proxy Writers.
RTPS::WRITER	When a local DataWriter is matched for the first time with a remote DataReader, the RTPS Reader context is added to the DataWriter in order to manage the RTPS Proxy Readers.
REM_READER	An RTPS Proxy Reader descriptor is used for each matched remote DataWriter.
REM_WRITER	An RTPS Proxy Writer descriptor is used for each matched remote DataReader.
RTPS::CCREF	When data samples are queued in RTPS, whether for a Proxy Reader or for a Proxy Writer, they are queued using Cache Change Reference descriptors, that point to cloned Change descriptors.
MSG_BUF	An RTPS Message descriptor is used, both when RTPS messages are received, and when they are under construction, i.e. when RTPS Submessages are added to them.
MSG_ELEM_BUF	An RTPS Submessage descriptor is used both for Submessage headers, as well as for extra data that needs to be concatenated to a submessage.
IP_CX	An IP Socket descriptor. These are used for both IPv4 and IPv6 UDP, TCP, DTLS/UDP and TLS/TCP connections.
DYNTYPE	Dynamic type descriptor context.
DYNDREF	Reference to a Dynamic data sample.
DATA_HOLDER	A DataHolder context for security handshakes and crypto tokens.
PROPERTY	A string property as used in security handshakes.
BIN_PROPERTY	A binary string property as used in security handshakes.
HOLDER_SEQ	A sequence of DataHolders.

This is then followed by a number of columns specifying either a number of bytes or a counter:

Table 40: spool and spoola counters

Column	Description
Addr	The memory address of the preallocated pool.
PSize	The size of the preallocated pool.
BSize	Block size of each element in the pool.
Rsrvd	Number of reserved elements.
Max	Maximum number of blocks that may be allocated of the pool (or * if unlimited).
MPUse	Maximum number of blocks ever in use during the program lifetime.
CPUse	Current number of blocks in use by the program.
MXUse	Maximum number of blocks ever dynamically allocated.
CXUse	Current number of blocks that are dynamically allocated.
Alloc	Total number of dynamic allocations.
Nomem	Total number of unsuccessful allocations (i.e. out-of-memory conditions).
Block	First available pool block.



Note: If the DDS library was compiled with the `FORCE_MALLOC` define, which disables preallocated pools, then a number of fields will not be significant, such as `Addr`, `PSize`, `Rsrvd`, `MPUse`, `CPUse` and `Block` and these will then be set to 0 (or `NULL`).



Note: The `Addr` and `Block` columns are only shown with the `spoola` command.

As can be seen from this example, not a lot of memory was ever used, but the preallocated pools consume almost 400 KBytes, mainly due to the number of preallocated DDS IP contexts (currently set to 1024). The reason for this is that this pool can currently not be extended, so when a huge number of devices are present in a domain, all having a large set of IP addresses, this might overflow the IP connections table. In practice, setting this maximum to something more reasonable like 64 shouldn't give problems.

Optimizing memory usage can be done very easily by looking at both the `MPUse` and `MXUse` fields, which, when added together, indicate the total number of blocks ever used in the program. Setting the preallocated pool requirements then allows to allocate everything from that pool and will no longer use `malloc()` for those pool elements.

Preallocated pools have a number of advantages compared to pure `malloc()` and `free()` functions:

1. Allocation and release functions are at least twice as fast compared to dynamic memory allocation functions.
2. There is no memory fragmentation overhead (saves ~20%).
3. Pool usage is deterministic and never leads to unexpected out-of-memory conditions while using less elements than the pool sizes permit, since all memory blocks will be allocated at startup.
4. Pool usage can be restricted in order to limit the number of memory resources that a DDS program is allowed to use.

Using the `FORCE_MALLOC` define is alternative way of handling memory, which results in ALL memory being allocated via `malloc()`, effectively bypassing the preallocated pool mechanism.

Sockets

The `scx` command shows the IP resources:

```
>scx
# of IP receive events   = 67988
# of DTLS server events  = 0
```

```
# of DTLS receive events = 0
Sending UDP socket:
  UDP 0.0.0.0:0 fd:4 id:0
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:887928/0 packets Tx/Rx:13865/0
Locators:
  UDP 10.0.0.29:7667 fd:7 id:1 USER UCAST SRC_MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:3613328/1293220 packets Tx/Rx:20308/20197
  UDP 239.255.0.1:7651 fd:9 id:1 USER MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/7444856 packets Tx/Rx:0/47578
  UDP 10.0.0.29:7666 fd:8 id:1 META UCAST SRC_MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:20644/4348 packets Tx/Rx:69/26
  UDP 239.255.0.1:7650 fd:10 id:1 META MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/56988 packets Tx/Rx:0/187
```

Where errors, empty, too_short and no_buffers indicate respectively the number of errors that occurred, the number of unexpted empty reads, too short messages and times there was not enough memory to handle reception properly.

The octets and packets statistics can be useful to determine how much traffic was present for individual peers.

Following shows a more complex set of connections when the TLS/TCP protocol is used as well as IPv6, using the spoola command:

```
scxa
# of IP receive events = 33
# of DTLS server events = 0
# of DTLS receive events = 0
Sending UDP socket:
  (0) UDP 0.0.0.0:0 fd:4 id:0
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/0 packets Tx/Rx:0/0
    self:0x8ce9b0
Sending UDPv6 socket:
  (0) UDP :::0 fd:5
    id:0
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/0 packets Tx/Rx:0/0
    self:0x8d9e30
TCP client:
  (0) TLS.C 0.0.0.0:0 fd:21 id:0 SECURE
    mode:client state:Established/C:Control
    dest:127.0.0.1:5400* prefix:92767f10:00005c42:03e80000
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:636/792 packets Tx/Rx:13/13
    self:0x8cd4e0 clients:0x7fe474076ae0
Locators:
  (1) UDP 10.0.0.29:9175 fd:8 id:1 USER UCAST SRC_MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/0 packets Tx/Rx:0/0
    self:0x8e2cf0
  (2) UDP 2a02:a03f:669:3100:1c3c:6a97:d469:750:9175 fd:9
    id:1 USER UCAST SRC_MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/0 packets Tx/Rx:0/0
    self:0x8e2df0
  (3) UDP 2a02:a03f:669:3100:92e6:baff:fe10:7817:9175 fd:10
    id:1 USER UCAST SRC_MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/0 packets Tx/Rx:0/0
    self:0x8e7bf0
  (4) TCP.D 10.0.0.29:9175 (fd:0) id:1 USER UCAST RX
    state:Closed/D:Idle
    prefix:1752a037:2c14ee02:1f1980b8
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/1180 packets Tx/Rx:0/11
    self:0x8e7cf0
  (5) UDP 239.255.0.1:9151 fd:11 id:1 USER MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/0 packets Tx/Rx:0/0
    self:0x8e7df0
  (6) UDP ff03::80:9151 fd:12
    id:1 USER MCAST
    r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
    octets Tx/Rx:0/0 packets Tx/Rx:0/0
```



```

self:0x8e7ef0
(7) TCP.D 239.255.0.1:9150 (fd:0)      id:1  META MCAST SECURE RX
state:Closed/D:Idle
prefix:1752a037:2c14ee02:1f1980b8
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:0/9868 packets Tx/Rx:0/9
self:0x8e4b10
(8) UDP 10.0.0.29:9174 fd:13      id:1  META UCAST SRC_MCAST
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:11924/0 packets Tx/Rx:11/0
self:0x8e4c10
(9) UDP 2a02:a03f:669:3100:1c3c:6a97:d469:750:9174 fd:14
id:1  META UCAST SRC_MCAST
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:11924/0 packets Tx/Rx:11/0
self:0x8e4d10
(10) UDP 2a02:a03f:669:3100:92e6:baff:fe10:7817:9174 fd:15
id:1  META UCAST SRC_MCAST
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:11924/0 packets Tx/Rx:11/0
self:0x8e4e10
(11) TCP.D 10.0.0.29:9174 (fd:0)      id:1  META UCAST RX
state:Closed/D:Idle
prefix:1752a037:2c14ee02:1f1980b8
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:0/10792 packets Tx/Rx:0/40
self:0x8e4f10
(12) UDP 239.255.0.1:9150 fd:16      id:1  META MCAST
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:0/11924 packets Tx/Rx:0/11
self:0x8e5010
(13) UDP ff03::80:9150 fd:17
id:1  META MCAST
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:0/23848 packets Tx/Rx:0/22
self:0x8e3000
(14) DTLS.S 10.0.0.29:4635 fd:18      id:1  USER META UCAST SECURE SERVER
mode:server
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:0/0 packets Tx/Rx:0/0
self:0x8e31b0 parent:0x8e31b0
(15) DTLS.S 2a02:a03f:669:3100:1c3c:6a97:d469:750:4635 fd:19
id:1  USER META UCAST SECURE SERVER
mode:server
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:0/0 packets Tx/Rx:0/0
self:0x8eb780 parent:0x8eb780
(16) DTLS.S 2a02:a03f:669:3100:92e6:baff:fe10:7817:4635 fd:20
id:1  USER META UCAST SECURE SERVER
mode:server
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:0/0 packets Tx/Rx:0/0
self:0x8eb880 parent:0x8eb880 next:0x8eb780
(17) TLS.D 0.0.0.0:9174 fd:23      id:1  META UCAST SECURE RX
mode:client state:Established/D:Data dest:127.0.0.1:5400*
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:40/10976 packets Tx/Rx:1/41
self:0x7fe47401d550 parent:0x8cd4e0 next:(19) group:(11) paired:(19)
(18) TLS.D 0.0.0.0:0 fd:24      id:1  USER UCAST SECURE TX
mode:client state:Established/D:Data
prefix:1752a037:2c14ac84:e9ddb9c7
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:1332/1204 packets Tx/Rx:13/12
self:0x7fe47403c7c0 parent:0x8cd4e0 next:(17) paired:(20)
(19) TLS.D 0.0.0.0:0 fd:22      id:1  META UCAST SECURE TX
mode:client state:Established/D:Data
prefix:1752a037:2c14ac84:e9ddb9c7
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:9860/10816 packets Tx/Rx:37/41
self:0x7fe474042540 parent:0x8cd4e0 paired:(17)
(20) TLS.D 0.0.0.0:9175 fd:25      id:1  USER UCAST SECURE RX
mode:client state:Established/D:Data dest:127.0.0.1:5400*
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:40/1248 packets Tx/Rx:1/12
self:0x7fe47405a8f0 parent:0x8cd4e0 next:(18) group:(4) paired:(18)
(21) TLS.D 0.0.0.0:9150 fd:26      id:1  META MCAST SECURE RX
mode:client state:Established/D:Data dest:127.0.0.1:5400*
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:40/2256 packets Tx/Rx:1/4
self:0x7fe474068600 parent:0x8cd4e0 next:(20) group:(7) paired:(22)

```

```
(22) TLS.D 0.0.0.0:0 fd:27 id:1 META MCAST SECURE TX
mode:client state:Established/D:Data
dest:239.255.0.1:9150* prefix:1752a037:2c14ac84:e9ddb9c7
r.errors:0 w.errors:0 empty:0 too_short:0 no buffers:0
octets Tx/Rx:3304/2244 packets Tx/Rx:4/4
self:0x7fe474076ae0 parent:0x8cd4e0 next:(21) paired:(21)
```



Note: In this example, the `dcps` test program was used in domain 7 with DDS security enabled, and a TCP/TLS connection to a remote secure TLS/TCP server. In practice, this would lead to double encryption, i.e. the payload is encrypted by the DDS Security mechanism and the RTPS messages are encrypted again over the TLS/TCP connections.



Note: The 3 DTLS.S connections are secure DTLS server connections. The peer doesn't use UDP, and even if it would, the native DDS Security still would take priority over DTLS/UDP if both have native DDS security enabled as in this case. So these connections would only be used between UDP entities using DTLS/UDP transport security.

The `scx` and `scxa` commands give the socket connections that DDS is effectively using. It is sometimes interesting to be able to see other information, such as the IP addresses that are available in the system and the file descriptors that are effectively used by DDS.

The first can be shown with the `sdip` command:

```
>sdip
lo:      ?      IPv6  127.0.0.1, ::1
eth0:    up     IP IPv6  10.0.0.29, fe80::92e6:baff:fe10:7817,
2a02:a03f:669:3100:92e6:baff:fe10:7817, 2a02:a03f:669:3100:1c3c:6a97:d469:750
eth1:    down   IPv6
```

As can be seen from this output, there are two IP interfaces, apart from the loopback interface.

The first one, i.e. `eth0`, is up and has both a local IPv4 address and a number of IPv6 addresses enabled.

The second one, `eth1`, is down.

The `sfd` command then shows the actual file descriptors in use:

```
>sfd
0: [0] DDS.debug { POLLIN POLLPRI POLLERR POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLERR POLLHUP POLLNVAL} Rxfct=0x49b16e,
U=0x8b9aa8
1: [3] DDS.netlink { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x4b67fc, U=(nil)
2: [6] DDS.wakeup { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x40a654, U=(nil)
3: [8] DDS.UDPv4 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e2cf0
4: [9] DDS.UDPv6 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e2df0
5: [10] DDS.UDPv6 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e7bf0
6: [11] DDS.UDPv4 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e7df0
7: [12] DDS.UDPv6 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e7ef0
8: [13] DDS.UDPv4 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e4c10
9: [14] DDS.UDPv6 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e4d10
10: [15] DDS.UDPv6 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e4e10
11: [16] DDS.UDPv4 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e5010
12: [17] DDS.UDPv6 { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x5055fa, U=0x8e3000
13: [18] DDS.DTLS-S { POLLIN POLLPRI} -> { POLLIN POLLPRI} Rxfct=0x50d7fb, U=0x8e31b0
14: [19] DDS.DTLS-S { POLLIN POLLPRI} -> { POLLIN POLLPRI} Rxfct=0x50d7fb, U=0x8eb780
15: [20] DDS.DTLS-S { POLLIN POLLPRI} -> { POLLIN POLLPRI} Rxfct=0x50d7fb, U=0x8eb880
16: [21] DDS.TLS-C { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x51d041, U=0x8cd4e0
17: [22] DDS.TLS-C { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x51d041, U=0x7fe474042540
18: [23] DDS.TLS-C { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x51d041, U=0x7fe47401d550
19: [24] DDS.TLS-C { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x51d041, U=0x7fe47403c7c0
20: [25] DDS.TLS-C { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x51d041, U=0x7fe47405a8f0
21: [26] DDS.TLS-C { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x51d041, U=0x7fe474068600
22: [27] DDS.TLS-C { POLLIN POLLPRI POLLHUP POLLNVAL} -> { POLLIN POLLPRI POLLHUP POLLNVAL} Rxfct=0x51d041, U=0x7fe474076ae0
```

Locators

The `sloc` command gives the following output:

```
>sloc
UDP:10.0.0.29:4635-DTLS(S)(14)*2
UDP:239.255.0.1:9150(12)*25
UDP:239.255.0.1:9151(5)*4
UDP:10.0.0.29:9174(8)*22
UDP:10.0.0.29:9175(1)*2
UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:4635-DTLS(S)(15)*2
UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:4635-DTLS(S)(16)*2
UDPv6:ff03::80:9150(13)*25
UDPv6:ff03::80:9151(6)*4
UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:9174(9)*22
UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:9174(10)*22
UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:9175(2)*2
```

```

UDpv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:9175(3)*2
TCP:10.0.0.29:5400-TLS(S)*1
TCP:239.255.0.1:9150-(null)(C)(7)*5
TCP:10.0.0.29:9166(19)*17
TCP:10.0.0.29:9167(18)*3
TCP:10.0.0.29:9174(11)*22
TCP:10.0.0.29:9175(4)*2

```

Each locator as displayed as:

- Protocol type.
- The protocol specific address (unicast or multicast IP address).
- A protocol specific port number (:port).
- If transport security is used, this is indicated (-DTLS(<mode>) or -TLS(<mode>)), where <mode> is either C: client, S: server, or H: handled by server.
- If the locator is associated with an IP connection (see scx and scxa), the connection handle is given as (<handle>).
- How many times the locator is referenced (the *n).

As can be seen from the display, locators are reused a lot.

Quality of Service

Concept definition.

The `sqos` command shows the following output:

```

>sqos
5: 0x9b865b0*2
7: 0x9b86660*5
23: 0x9b86608*22
33: 0x9b866b8*2

```

The QoS display is somewhat simplistic because it simply displays the address of the in-memory unified QoS data container, following the hash key and followed by the reference counter.

It does however show that reusing QoS parameters is very advantageous in order to reduce the overall memory footprint.

Displaying specific entity QoS parameters is much more appropriate, as can be done with the `seqos` command.

In order to do this, the user needs to specify the handle of the entity that needs to be displayed:

```

>seq 42
Reader QoS:
  USER_DATA: <empty>
  DURABILITY: TRANSIENT_LOCAL
  DEADLINE: <inf>
  LATENCY_BUDGET: 0s
  OWNERSHIP: SHARED
  LIVELINESS: AUTOMATIC, lease_duration=<inf>
  RELIABILITY: RELIABLE, max_blocking_time=0.100000000s
  DESTINATION_ORDER: BY_RECEPTION_TIMESTAMP
  HISTORY: KEEP_LAST, depth=1
  RESOURCE_LIMITS: max_samples=<inf>, max_inst=<inf>, max_samples_per_inst=<inf>
  TIME_BASED_FILTER: minimum_separation=0s
  READER_DATA_LIFECYCLE: autopurge_nowriter=<inf>, autopurge_disposed=<inf>
G: PARTITION: <none>
  GROUP_DATA: <empty>
T: TOPIC_DATA: <empty>
>seq 45
Writer QoS:
  USER_DATA: <empty>
  DURABILITY: TRANSIENT_LOCAL
  DURABILITY_SERVICE:
  DEADLINE: <inf>
  LATENCY_BUDGET: 0s
  OWNERSHIP: SHARED
  OWNERSHIP_STRENGTH: 0
  LIVELINESS: AUTOMATIC, lease_duration=<inf>
  RELIABILITY: RELIABLE, max_blocking_time=0.100000000s
  TRANSPORT_PRIORITY: 0
  LIFESPAN: <inf>
  DESTINATION_ORDER: BY_RECEPTION_TIMESTAMP

```

```

HISTORY: KEEP_LAST, depth=1
RESOURCE_LIMITS: max_samples=<inf>, max_inst=<inf>, max_samples_per_inst=<inf>
WRITER_DATA_LIFECYCLE: autodispose_unregistered_instances=1
G: PARTITION: <none>
  GROUP_DATA: <empty>
T: TOPIC_DATA: <empty>

```



Note: The G: and T: references in the display refer to data that is effectively present in the Group (Partition QoS and Group Data) and Topic (Topic Data) respectively.



Note: This command is not restricted to DataReader and DataWriter QoS. It can be used to display the QoS parameters of any entity in the DDS domains that has QoS parameters. It is just a matter of giving the correct instance handle.

Domain data

There are three commands for displaying domain data:

- The `sdomain` command is the most versatile, but requires some arguments.
- The `sdisc` and `sdisca` commands don't expect arguments. The `sdisc` command shows minimal info, the `sdisca` command shows most available info.

The `sdomain` command accepts 3 arguments:

- The domain identifier.
- The local information to display as a hexadecimal bitmap.
- The discovered information to display as a hexadecimal bitmap.

The bitmap has following values:

Table 41: Domain Information bitmap flags

Flag	Description
1	Include Locators
2	Include Builtin endpoints
4	Include Types.
8	Include Topics.
16	Include Guard records.
32	Include Peer participants.

Some examples without security enabled, using domain id 1:

```

>sdisc
Domain 1 (pid=10): {1}
  GUID prefix: 92767f10:010a6e9e:03e80000
  RTPS Protocol version: v2.1
  Vendor Id: 1.14 - Technicolor, Inc. - Qeo
  Technicolor DDS version: 4.0-0, Forward: 0
  SecureTransport: none
  Authorisation: Authenticated
  Entity name: Technicolor Shapes
  Flags: Enabled
  Type: C
  Meta Unicast:
    UDP:10.0.0.29:7680(3) {MD,UC} H:3
  Meta Multicast:
    UDP:239.255.0.1:7650(4) {MD,MC} H:4
  Default Unicast:
    UDP:10.0.0.29:7681(1) {UD,UC} H:1
  Default Multicast:
    UDP:239.255.0.1:7651(2) {UD,MC} H:2
  Manual Liveliness: 0
  Lease duration: 50.000000000s
  Endpoints: 11 entries (6 readers, 5 writers).

```

```

Security: level=Unclassified, access=any, RTPS=clear
Resend period: 10.000000000s
Destination Locators:
  UDP:239.255.0.1:7650(4) {MD,MC} H:4
  TCP:239.255.0.1:7650 {MD,MC}
Discovered participants:
Peer #0: {20} - Local activity: 3.07s
  GUID prefix: 92767f10:01086e8c:03e80000
  RTPS Protocol version: v2.1
  Vendor Id: 1.14 - Technicolor, Inc. - Qeo
  Technicolor DDS version: 4.0-0, Forward: 0
  SecureTransport: remote: DDSSEC, local: DDSSEC
  Authorisation: Authenticated
  Entity name: Technicolor Shapes
  Meta Unicast:
    UDP:10.0.0.29:7676 {MD,UC}
    UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:7676 {MD,UC}
    UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:7676 {MD,UC}
  Meta Multicast:
    UDP:239.255.0.1:7650(4) {MD,MC} H:4
    UDPv6:ff03::80:7650 {MD,MC}
  Default Unicast:
    UDP:10.0.0.29:7677 {UD,UC}
    UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:7677 {UD,UC}
    UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:7677 {UD,UC}
  Default Multicast:
    UDP:239.255.0.1:7651(2) {UD,MC} H:2
    UDPv6:ff03::80:7651 {UD,MC}
  Manual Liveliness: 0
  Lease duration: 50.000000000s
  Endpoints: 9 entries (5 readers, 4 writers).
  Source:
    UDP:10.0.0.29:7676 {MD,UC}
  Timer = 46.24s
Peer #1: {27} - Local activity: 3.07s
  GUID prefix: 92767f10:01096e8f:03e80000
  RTPS Protocol version: v2.1
  Vendor Id: 1.14 - Technicolor, Inc. - Qeo
  Technicolor DDS version: 4.0-0, Forward: 0
  SecureTransport: remote: DDSSEC, local: DDSSEC
  Authorisation: Authenticated
  Entity name: Technicolor Shapes
  Meta Unicast:
    UDP:10.0.0.29:7678 {MD,UC}
    UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:7678 {MD,UC}
    UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:7678 {MD,UC}
  Meta Multicast:
    UDP:239.255.0.1:7650(4) {MD,MC} H:4
    UDPv6:ff03::80:7650 {MD,MC}
  Default Unicast:
    UDP:10.0.0.29:7679 {UD,UC}
    UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:7679 {UD,UC}
    UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:7679 {UD,UC}
  Default Multicast:
    UDP:239.255.0.1:7651(2) {UD,MC} H:2
    UDPv6:ff03::80:7651 {UD,MC}
  Manual Liveliness: 0
  Lease duration: 50.000000000s
  Endpoints: 8 entries (4 readers, 4 writers).
  Source:
    UDP:10.0.0.29:7678 {MD,UC}
  Timer = 46.24s

>sdisca
Domain 1 (pid=10): {1}
  GUID prefix: 92767f10:010a6e9e:03e80000
  RTPS Protocol version: v2.1
  Vendor Id: 1.14 - Technicolor, Inc. - Qeo
  Technicolor DDS version: 4.0-0, Forward: 0
  SecureTransport: none
  Authorisation: Authenticated
  Entity name: Technicolor Shapes
  Flags: Enabled
  Type: C
  Meta Unicast:
    UDP:10.0.0.29:7680(3) {MD,UC} H:3
  Meta Multicast:
    UDP:239.255.0.1:7650(4) {MD,MC} H:4
  Default Unicast:
    UDP:10.0.0.29:7681(1) {UD,UC} H:1

```

```

Default Multicast:
  UDP:239.255.0.1:7651(2) {UD,MC} H:2
Manual Liveliness: 0
Lease duration: 50.000000000s
Endpoints: 11 entries (6 readers, 5 writers).
  000001-7, {36}, InlineQoS: No, Reader, Square/ShapeType
  000002-7, {38}, InlineQoS: No, Reader, Triangle/ShapeType
  000003-2, {40}, InlineQoS: No, Writer, Square/ShapeType
Topics:
  BuiltinParticipantMessageReader/ParticipantMessageData
  BuiltinParticipantMessageWriter/ParticipantMessageData
  SEDPBuiltinPublicationsReader/PublicationBuiltinTopicData
  SEDPBuiltinPublicationsWriter/PublicationBuiltinTopicData
  SEDPBuiltinSubscriptionsReader/SubscriptionBuiltinTopicData
  SEDPBuiltinSubscriptionsWriter/SubscriptionBuiltinTopicData
  SPDPBuiltinParticipantReader/ParticipantBuiltinTopicData
  SPDPBuiltinParticipantWriter/ParticipantBuiltinTopicData
  Square/ShapeType Triangle/ShapeType
Security: level=Unclassified, access=any, RTPS=clear
Resend period: 10.000000000s
Destination Locators:
  UDP:239.255.0.1:7650(4) {MD,MC} H:4
  TCP:239.255.0.1:7650 {MD,MC}
Discovered participants:
  Peer #0: {20} - Local activity: 3.02s
    GUID prefix: 92767f10:01086e8c:03e80000
    RTPS Protocol version: v2.1
    Vendor Id: 1.14 - Technicolor, Inc. - Qeo
    Technicolor DDS version: 4.0-0, Forward: 0
    SecureTransport: remote: DDSSEC, local: DDSSEC
    Authorisation: Authenticated
    Entity name: Technicolor Shapes
    Meta Unicast:
      UDP:10.0.0.29:7676 {MD,UC}
      UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:7676 {MD,UC}
      UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:7676 {MD,UC}
    Meta Multicast:
      UDP:239.255.0.1:7650(4) {MD,MC} H:4
      UDPv6:ff03::80:7650 {MD,MC}
    Default Unicast:
      UDP:10.0.0.29:7677 {UD,UC}
      UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:7677 {UD,UC}
      UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:7677 {UD,UC}
    Default Multicast:
      UDP:239.255.0.1:7651(2) {UD,MC} H:2
      UDPv6:ff03::80:7651 {UD,MC}
    Manual Liveliness: 0
    Lease duration: 50.000000000s
    Endpoints: 9 entries (5 readers, 4 writers).
      000001-7, {43}, InlineQoS: No, Reader, Square/ShapeType
      000002-7, {44}, InlineQoS: No, Reader, Triangle/ShapeType
      000003-2, {45}, InlineQoS: No, Writer, Triangle/ShapeType
    Topics:
      Square/ShapeType Triangle/ShapeType
    Source:
      UDP:10.0.0.29:7676 {MD,UC}
    Timer = 46.76s
  Peer #1: {27} - Local activity: 3.02s
    GUID prefix: 92767f10:01096e8f:03e80000
    RTPS Protocol version: v2.1
    Vendor Id: 1.14 - Technicolor, Inc. - Qeo
    Technicolor DDS version: 4.0-0, Forward: 0
    SecureTransport: remote: DDSSEC, local: DDSSEC
    Authorisation: Authenticated
    Entity name: Technicolor Shapes
    Meta Unicast:
      UDP:10.0.0.29:7678 {MD,UC}
      UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:7678 {MD,UC}
      UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:7678 {MD,UC}
    Meta Multicast:
      UDP:239.255.0.1:7650(4) {MD,MC} H:4
      UDPv6:ff03::80:7650 {MD,MC}
    Default Unicast:
      UDP:10.0.0.29:7679 {UD,UC}
      UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:7679 {UD,UC}
      UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:7679 {UD,UC}
    Default Multicast:
      UDP:239.255.0.1:7651(2) {UD,MC} H:2
      UDPv6:ff03::80:7651 {UD,MC}
    Manual Liveliness: 0

```

```

Lease duration: 50.000000000s
Endpoints: 8 entries (4 readers, 4 writers).
    000001-2, {41}, InlineQoS: No, Writer, Square/ShapeType
    000002-7, {42}, InlineQoS: No, Reader, Square/ShapeType
Topics:
    Square/ShapeType
Source:
    UDP:10.0.0.29:7678 {MD,UC}
Timer = 46.76s

```



Note: Instance handles are displayed as {n}, where n is a decimal number. This number can be referenced when using other commands such as seqos and sdcache (see further).

As can be seen in this example, the shapes program has discovered two remote participants, with the first having a square writer and a square reader, and the other having both square and triangle readers and a triangle writer.

Another example showing the two DCPS test programs with security enabled and connected over TLS/TCP:

```

>sdisca
Domain 7 (pid=7): {1}
  GUID prefix: 1752a037:2c14ee02:1f1980b8
  RTPS Protocol version: v2.1
  Vendor Id: 1.14 - Technicolor, Inc. - Qeo
  Technicolor DDS version: 4.0-0, Forward: 0
  SecureTransport: remote: DTLS+TLS+DDSSEC, local: DTLS+TLS+DDSSEC
  Authorisation: Authenticated
  Identity tokens:
    Auth:DTLS-Name:
      4443505320546573742070726f6772616d00
    DDS:Auth:X.509-PEM-SHA256:
      e860bfff08c8eal7172adc7bcb2e6672cc1ddf4ab703074e8f4f933b1fb7976a
    Qeo:Auth:X.509-PEM-SHA256:
      e860bfff08c8eal7172adc7bcb2e6672cc1ddf4ab703074e8f4f933b1fb7976a
  Permission tokens:
    Access:DTLS-Name-MD5:
      ce9423c1697e44dc749b6249209a5a27
    DDS:Access:PKI-Signed-XML-Permissions-SHA256:
      3e5fd7cf259ff0282a998286549ed44cc3f2a991412f9d804dfd477103ff3134
    Qeo:Access:PKI-Signed-Policy-SHA256:
      3e5fd7cf259ff0282a998286549ed44cc3f2a991412f9d804dfd477103ff3134
  Entity name: Technicolor DCPS Tester
  Flags: NoMCast, Enabled
  Type: C
  Meta Unicast:
    UDP:10.0.0.29:9174(8) {MD,UC} H:8
    UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:9174(9) {MD,UC} H:9
    UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:9174(10) {MD,UC} H:10
    TCP:10.0.0.29:9174(11) {MD,UC} H:11
  Meta Multicast:
    UDP:239.255.0.1:9150(12) {MD,MC} H:12
    UDPv6:ff03::80:9150(13) {MD,MC} H:13
  Default Unicast:
    UDP:10.0.0.29:9175(1) {UD,UC} H:1
    UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:9175(2) {UD,UC} H:2
    UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:9175(3) {UD,UC} H:3
    TCP:10.0.0.29:9175(4) {UD,UC} H:4
  Default Multicast:
    UDP:239.255.0.1:9151(5) {UD,MC} H:5
    UDPv6:ff03::80:9151(6) {UD,MC} H:6
  Secure Tunnel:
    UDP:10.0.0.29:4635-DTLS(S)(14) {UD,MD,UC,E,S} H:14 DTLS
    UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:4635-DTLS(S)(15) {UD,MD,UC,E,S} H:15 DTLS
    UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:4635-DTLS(S)(16) {UD,MD,UC,E,S} H:16 DTLS
  Manual Liveliness: 0
  Lease duration: 50.000000000s
  Endpoints: 22 entries (11 readers, 11 writers).
    000001-2, {46}, InlineQoS: No, Writer, HelloWorld/HelloWorldData
    000002-7, {48}, InlineQoS: No, Reader, HelloWorld/HelloWorldData
  Topics:
    BuiltinParticipantMessageReader/ParticipantMessageData
    BuiltinParticipantMessageSecureReader/ParticipantMessageData
    BuiltinParticipantMessageSecureWriter/ParticipantMessageData
    BuiltinParticipantMessageWriter/ParticipantMessageData
    BuiltinParticipantStatelessMessageReader/ParticipantStatelessMessage
    BuiltinParticipantStatelessMessageWriter/ParticipantStatelessMessage
    BuiltinParticipantVolatileMessageSecureReader/ParticipantVolatileSecureMessage
    BuiltinParticipantVolatileMessageSecureWriter/ParticipantVolatileSecureMessage
    HelloWorld/HelloWorldData

```

```

QeobuiltinPolicyStateReader/PolicyUpdaterMessageData
QeobuiltinPolicyStateWriter/PolicyUpdaterMessageData
SEDPbuiltinPublicationsReader/PublicationBuiltinTopicData
SEDPbuiltinPublicationsSecureReader/PublicationBuiltinTopicDataSecure
SEDPbuiltinPublicationsSecureWriter/PublicationBuiltinTopicDataSecure
SEDPbuiltinPublicationsWriter/PublicationBuiltinTopicData
SEDPbuiltinSubscriptionsReader/SubscriptionBuiltinTopicData
SEDPbuiltinSubscriptionsSecureReader/SubscriptionBuiltinTopicDataSecure
SEDPbuiltinSubscriptionsSecureWriter/SubscriptionBuiltinTopicDataSecure
SEDPbuiltinSubscriptionsWriter/SubscriptionBuiltinTopicData
SPDPbuiltinParticipantReader/ParticipantBuiltinTopicDataSecure
SPDPbuiltinParticipantWriter/ParticipantBuiltinTopicDataSecure
Security: level=Secret, access=enforce, RTPS=clear
Resend period: 10.000000000s
Destination Locators:
  UDP:239.255.0.1:9150(12) {MD,MC} H:12
  UDPv6:ff03::80:9150(13) {MD,MC} H:13
  TCP:239.255.0.1:9150-(null)(C)(7) {MD,MC,E} H:7
Discovered participants:
Peer #0: {49}
  GUID prefix: 1752a037:2c14ac84:e9ddb9c7
  RTPS Protocol version: v2.1
  Vendor Id: 1.14 - Technicolor, Inc. - Qeo
  Technicolor DDS version: 4.0-0, Forward: 0
  SecureTransport: remote: DTLS+TLS+DDSSEC, local: DTLS+TLS+DDSSEC
  Authorisation: Authenticated
  Identity tokens:
    Qeo:Auth:X.509-PEM-SHA256 (*2):
      e860bfff08c8eal7172adc7bcb2e6672cc1ddf4ab703074e8f4f933b1fb7976a
    DDS:Auth:X.509-PEM-SHA256:
      e860bfff08c8eal7172adc7bcb2e6672cc1ddf4ab703074e8f4f933b1fb7976a
    Auth:DTLS-Name:
      4443505320546573742070726f6772616d00
  Permission tokens:
    Qeo:Access:PKI-Signed-Policy-SHA256 (*2):
      3e5fd7cf259ff0282a998286549ed44cc3f2a991412f9d804dfd477103ff3134
    DDS:Access:PKI-Signed-XML-Permissions-SHA256:
      3e5fd7cf259ff0282a998286549ed44cc3f2a991412f9d804dfd477103ff3134
    Access:DTLS-Name-MD5:
      ce9423c1697e44dc749b6249209a5a27
  Entity name: Technicolor DCPS Tester
  Flags: NoMCast
  Meta Unicast:
    TCP:10.0.0.29:9166(19) {MD,UC} H:19
  Default Unicast:
    TCP:10.0.0.29:9167(18) {UD,UC} H:18
  Secure Tunnel:
    TCP:10.0.0.29:5400-TLS(S) {UD,MD,UC,E,S} TLS
  Manual Liveliness: 0
  Lease duration: 50.000000000s
  Endpoints: 18 entries (9 readers, 9 writers).
    000001-2, {67}, InlineQoS: No, Writer, HelloWorld/HelloWorldData
    000002-7, {66}, InlineQoS: No, Reader, HelloWorld/HelloWorldData
  Topics:
    HelloWorld/HelloWorldData
  Source:
    TCP:239.255.0.1:9150-(null)(C)(7) {MD,MC,E} H:7
  Timer = 46.62s

```



Note: The output shows that there are 3 sets of security mechanisms that could be used (3 security token types). The first two are the Qeo security plugin and the standard DDS security plugin, both using native DDS security. The third uses transport security over DTLS/UDP. Since both peers support the Qeo handshake in combination with native DDS security, that one is selected.

Topics

Topics can either be displayed as the list of Topics in the domain, or more specific Topic information can be shown.

The `stopic` command expects a Domain Identifier, and then displays all Topics in the Domain as follows:

```

>stopic 1
BuiltinParticipantMessageReader/ParticipantMessageData:
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers:
  DW: {29} 92767f10:01096e8f:03e80000-000200-c2
  DW: {22} 92767f10:01086e8c:03e80000-000200-c2

```



```

Readers:
R: {17} 92767f10:010a6e9e:03e80000-000200-c7
BuiltinParticipantMessageWriter/ParticipantMessageData:
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers:
W: {19} 92767f10:010a6e9e:03e80000-000200-c2
Readers:
DR: {28} 92767f10:01096e8f:03e80000-000200-c7
DR: {21} 92767f10:01086e8c:03e80000-000200-c7
SEDPBuiltinPublicationsReader/PublicationBuiltinTopicData:
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers:
DW: {31} 92767f10:01096e8f:03e80000-000003-c2
DW: {24} 92767f10:01086e8c:03e80000-000003-c2
Readers:
R: {9} 92767f10:010a6e9e:03e80000-000003-c7
SEDPBuiltinPublicationsWriter/PublicationBuiltinTopicData:
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers:
W: {13} 92767f10:010a6e9e:03e80000-000003-c2
Readers:
DR: {30} 92767f10:01096e8f:03e80000-000003-c7
DR: {23} 92767f10:01086e8c:03e80000-000003-c7
SEDPBuiltinSubscriptionsReader/SubscriptionBuiltinTopicData:
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers:
DW: {33} 92767f10:01096e8f:03e80000-000004-c2
DW: {26} 92767f10:01086e8c:03e80000-000004-c2
Readers:
R: {11} 92767f10:010a6e9e:03e80000-000004-c7
SEDPBuiltinSubscriptionsWriter/SubscriptionBuiltinTopicData:
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers:
W: {15} 92767f10:010a6e9e:03e80000-000004-c2
Readers:
DR: {32} 92767f10:01096e8f:03e80000-000004-c7
DR: {25} 92767f10:01086e8c:03e80000-000004-c7
SPDPBuiltinParticipantReader/ParticipantBuiltinTopicData:
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers: <none>
Readers:
R: {7} 92767f10:010a6e9e:03e80000-000100-c7
SPDPBuiltinParticipantWriter/ParticipantBuiltinTopicData:
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers:
W: {5} 92767f10:010a6e9e:03e80000-000100-c2
Readers: <none>
Square/ShapeType:
# of local create/find_topic() calls = 1, # of discoveries = 3
Writers:
DW: {41} 92767f10:01096e8f:03e80000-000001-2
W: {40} 92767f10:010a6e9e:03e80000-000003-2
Readers:
DR: {43} 92767f10:01086e8c:03e80000-000001-7
DR: {42} 92767f10:01096e8f:03e80000-000002-7
R: {36} 92767f10:010a6e9e:03e80000-000001-7
Triangle/ShapeType:
# of local create/find_topic() calls = 1, # of discoveries = 2
Writers:
DW: {45} 92767f10:01086e8c:03e80000-000003-2
Readers:
DR: {44} 92767f10:01086e8c:03e80000-000002-7
R: {38} 92767f10:010a6e9e:03e80000-000002-7

```

The following information is displayed for each Topic:

- Topic name / type.
- Some Topic parameters needed to do correct Topic accounting.
- The list of Writers (local and discovered) associated with the Topic.
- The list of Readers (local and discovered) associated with the Topic.

The following information is displayed for each Writer or Reader:

- The type of endpoint (W = local Writer, DW = discovered Writer, R = local Reader and DR = discovered Reader).
- The handle of the endpoint (between '{' and '}').
- The unique GUID.

When the Topic name is specified, a lot more information is given, as shown below:

```
>stopic 0 Square
Square/ShapeType:
  Preferred: CDR, dynamic: 0, length: 140, keys: 1, mkeysize: 133
  struct ShapeType {
/*0*/   string<128> /*ESize:1*/ color;  //@Key
/*132*/ long x;
/*136*/ long y;
/*140*/ long shapesize;
  };
  # of local create/find_topic() calls = 1, # of discoveries = 3
  Writers:
    W: {45}      c0a84080:000008cc:03e80000-000003-2
    DW: {36}     c0a84080:0001091e:03e80000-000002-2
  Readers:
    R: {42}      c0a84080:000008cc:03e80000-000001-7
    DR: {37}     c0a84080:00020973:03e80000-000001-7
    DR: {34}     c0a84080:0001091e:03e80000-000001-7
```

As can be seen from this example, for individual topics, extra Topic type information is given:

- The on-the-wire encoding/decoding type (CDR in this case).
- Whether the type is dynamic or fixed-size.
- The size of a sample in bytes.
- Whether or not key information is present, and the total size of the key fields.
- The type itself as originally registered as an IDL type definition, giving extra information in comments like the field offset and the element size for arrays.

For filtered topics, if displayed with this command, even more data is displayed, e.g. the expression string, expression parameters and the generated bytecode program that is used for filtering data.

In the following example, a filtered topic is shown, as part of a topics list:

```
Triangle/ShapeType:
  Filter name: Filter_2
  Class name: DDSSQL
  Expression: x > %0 and x < %1 and y > %2 and y < %3
  Parameters:
    %0 = "50"
    %1 = "150"
    %2 = "30"
    %3 = "90"
  Bytecode program size: 38 bytes.
  # of local create/find_topic() calls = 1, # of discoveries = 0
  Writers: <none>
  Readers:
    R: {56} 0a000044:000c38a9:03e80000-000002-7
```

To show the filter topic in detail, the filter name needs to be specified as the topic name:

```
>stopic 0 Filter_2
Triangle/ShapeType:
  Filter name: Filter_2
  Class name: DDSSQL
  Expression: x > %0 and x < %1 and y > %2 and y < %3
  Parameters:
    %0 = "50"
    %1 = "150"
    %2 = "30"
    %3 = "90"
  Bytecode program:
0000      LDWS 1
0003      LPWS 0
0005      CMPWS
0006      BLE 37
0009      LDWS 1
0012      LPWS 1
0014      CMPWS
0015      BGE 37
0018      LDWS 2
0021      LPWS 2
0023      CMPWS
0024      BLE 37
0027      LDWS 2
```

```

0030      LPWS 3
0032      CMPWS
0033      BGE 37
0036      RETT
0037      RETF
Preferred: CDR, dynamic: 0, length: 140, keys: 1, mkeysize: 133
struct ShapeType {
    string<128> color;  //@key
    long x;
    long y;
    long shapesize;
};
# of local create/find_topic() calls = 1, # of discoveries = 0
Writers: <none>
Readers:
R: {56} 0a000044:000c38a9:03e80000-000002-7

```



Note: For discovered filtered topics, where the filter is specified by a remote Reader, the same information will also be present locally, and can also be displayed in this manner.

Endpoints

A list of endpoints and some condensed endpoint information can be displayed with the `sendpoints` and `sep` commands as shown here:

```

>sep
DCPS:
0/000001-07 R{42} 0/2 Square/ShapeType
0/000002-07 R{43} 0/1 Triangle/ShapeType
0/000003-02 W{45} 1/1 Square/ShapeType
0/000003-c2 W{13} 1/1 SEDPbuiltinPublicationsWriter/discoveredWriterData
0/000003-c7 R{9} 0/0 SEDPbuiltinPublicationsReader/discoveredWriterData
0/000004-c2 W{15} 2/2 SEDPbuiltinSubscriptionsWriter/discoveredReaderData
0/000004-c7 R{11} 0/0 SEDPbuiltinSubscriptionsReader/discoveredReaderData
0/000100-c2 W{5} 1/1 SPDPbuiltinParticipantWriter/SPDPdiscoveredParticipantData
0/000100-c7 R{7} 0/0 SPDPbuiltinParticipantReader/SPDPdiscoveredParticipantData
0/000200-c2 W{19} 0/0 BuiltinParticipantMessageWriter/ParticipantMessageData
0/000200-c7 R{17} 0/0 BuiltinParticipantMessageReader/ParticipantMessageData
RTPS:
0/000001-07 R{42} 0/2 Square/ShapeType
-> GUID: c0a84080:00021ed3:03e80000-000001-2
0/000002-07 R{43} 0/1 Triangle/ShapeType
-> GUID: c0a84080:00062810:03e80000-000003-2
0/000003-02 W{45} 0/1 Square/ShapeType
-> GUID: c0a84080:00021ed3:03e80000-000002-7
-> GUID: c0a84080:00062810:03e80000-000001-7
0/000003-c2 W{13} 0/1 SEDPbuiltinPublicationsWriter/discoveredWriterData
-> GUID: c0a84080:00062810:03e80000-000003-c7
-> GUID: c0a84080:00021ed3:03e80000-000003-c7
0/000003-c7 R{9} 0/0 SEDPbuiltinPublicationsReader/discoveredWriterData
-> GUID: c0a84080:00062810:03e80000-000003-c2
-> GUID: c0a84080:00021ed3:03e80000-000003-c2
0/000004-c2 W{15} 0/2 SEDPbuiltinSubscriptionsWriter/discoveredReaderData
-> GUID: c0a84080:00062810:03e80000-000004-c7
-> GUID: c0a84080:00021ed3:03e80000-000004-c7
0/000004-c7 R{11} 0/0 SEDPbuiltinSubscriptionsReader/discoveredReaderData
-> GUID: c0a84080:00062810:03e80000-000004-c2
-> GUID: c0a84080:00021ed3:03e80000-000004-c2
0/000100-c2 W{5} 1/1 SPDPbuiltinParticipantWriter/SPDPdiscoveredParticipantData
-> Locator: UDP:239.255.0.1:7400
0/000100-c7 R{7} 0/0 SPDPbuiltinParticipantReader/SPDPdiscoveredParticipantData
0/000200-c2 W{19} 0/0 BuiltinParticipantMessageWriter/ParticipantMessageData
-> GUID: c0a84080:00062810:03e80000-000200-c7
-> GUID: c0a84080:00021ed3:03e80000-000200-c7
0/000200-c7 R{17} 0/0 BuiltinParticipantMessageReader/ParticipantMessageData
-> GUID: c0a84080:00062810:03e80000-000200-c2
-> GUID: c0a84080:00021ed3:03e80000-000200-c2

```

Endpoints are displayed as two separate groups of endpoints:

- The first group lists the endpoints as known on DCPS level. This typically includes all endpoints that are defined in the domain without further information.
- The second group lists the endpoints that have RTPS connectivity with discovered endpoints. For each of those, either the ReaderLocator instance (SPDPbuiltinParticipantWriter) or the Proxy Reader or Writer contexts are shown as GUIDs.



Note: Not all endpoints are expected to be connected on RTPS. Technicolor DDS only creates RTPS Readers and RTPS Writers when remote Writers and Readers are actually discovered.

Endpoints are displayed as:

- A Domain Id/Entity Id pair, followed by
- the endpoint type (R or W) and handle,
- summary history cache information as *<number_of_samples> / <number_of_instances>*, followed by
- Topic name / Type name strings.

In the RTPS group, each matched endpoint is listed explicitly on separate lines.

The previous example listed all the endpoints within an open domain in which security was not enabled. Secure endpoints in a secure domain are shown with extra security parameters as an extra line specifying the security attributes, as shown in the following example:

```
>sep
DCPS:
7/000001-02 W{46} 3/3 HelloWorld/HelloWorldData
7/000002-07 R{48} 6/6 HelloWorld/HelloWorldData
7/000003-c2 W{13} 0/0 SEDPbuiltinPublicationsWriter/PublicationBuiltinTopicData
7/000003-c7 R{9} 0/0 SEDPbuiltinPublicationsReader/PublicationBuiltinTopicData
7/000004-c2 W{15} 0/0 SEDPbuiltinSubscriptionsWriter/SubscriptionBuiltinTopicData
7/000004-c7 R{11} 0/0 SEDPbuiltinSubscriptionsReader/SubscriptionBuiltinTopicData
7/000100-c2 W{5} 1/1 SPDPbuiltinParticipantWriter/ParticipantBuiltinTopicDataSecure
7/000100-c7 R{7} 0/0 SPDPbuiltinParticipantReader/ParticipantBuiltinTopicDataSecure
7/000200-c2 W{27} 1/1 BuiltinParticipantMessageWriter/ParticipantMessageData
7/000200-c7 R{25} 0/0 BuiltinParticipantMessageReader/ParticipantMessageData
7/000201-c3 W{33} 0 BuiltinParticipantStatelessMessageWriter/ParticipantStatelessMessage
7/000201-c4 R{35} 0 BuiltinParticipantStatelessMessageReader/ParticipantStatelessMessage
7/ff0003-c2 W{21} 1/1 SEDPbuiltinPublicationsSecureWriter/PublicationBuiltinTopicDataSecure
7/ff0003-c7 R{17} 0/0 SEDPbuiltinPublicationsSecureReader/PublicationBuiltinTopicDataSecure
7/ff0004-c2 W{23} 1/1 SEDPbuiltinSubscriptionsSecureWriter/SubscriptionBuiltinTopicDataSecure
7/ff0004-c7 R{19} 0/0 SEDPbuiltinSubscriptionsSecureReader/SubscriptionBuiltinTopicDataSecure
7/ff0200-c2 W{31} 0/0 BuiltinParticipantMessageSecureWriter/ParticipantMessageData
7/ff0200-c7 R{29} 0/0 BuiltinParticipantMessageSecureReader/ParticipantMessageData
7/ff0202-c3 W{37} 0 BuiltinParticipantVolatileMessageSecureWriter/ParticipantVolatileSecureMessage
7/ff0202-c4 R{39} 0 BuiltinParticipantVolatileMessageSecureReader/ParticipantVolatileSecureMessage
7/ff0401-c2 W{43} 1/1 QeobuiltinPolicyStateWriter/PolicyUpdaterMessageData
7/ff0401-c7 R{41} 0/1 QeobuiltinPolicyStateReader/PolicyUpdaterMessageData

RTPS:
7/000001-02 W{46} 0/3 HelloWorld/HelloWorldData
  Access control, Secure discovery, Encryption: payload(AES128_HMAC_SHA1)
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-000002-7 {66}
7/000002-07 R{48} 0/6 HelloWorld/HelloWorldData
  Access control, Secure discovery, Encryption: payload(AES128_HMAC_SHA1)
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-000001-2 {67}
7/000003-c2 W{13} 0/0 SEDPbuiltinPublicationsWriter/PublicationBuiltinTopicData
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-000003-c7 {58}
7/000003-c7 R{9} 0/0 SEDPbuiltinPublicationsReader/PublicationBuiltinTopicData
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-000003-c2 {59}
7/000004-c2 W{15} 0/0 SEDPbuiltinSubscriptionsWriter/SubscriptionBuiltinTopicData
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-000004-c7 {60}
7/000004-c7 R{11} 0/0 SEDPbuiltinSubscriptionsReader/SubscriptionBuiltinTopicData
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-000004-c2 {61}
7/000100-c2 W{5} 3/1 SPDPbuiltinParticipantWriter/ParticipantBuiltinTopicDataSecure
  -> Locator: UDP:239.255.0.1:9150(12) {5}
  -> Locator: UDPv6:ff03::80:9150(13) {5}
  -> Locator: TCP:239.255.0.1:9150-(null)(C)(7) {5}
7/000100-c7 R{7} 0/0 SPDPbuiltinParticipantReader/ParticipantBuiltinTopicDataSecure
7/000200-c2 W{27} 0/1 BuiltinParticipantMessageWriter/ParticipantMessageData
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-000200-c7 {54}
7/000200-c7 R{25} 0/0 BuiltinParticipantMessageReader/ParticipantMessageData
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-000200-c2 {55}
7/000201-c3 W{33} 0 BuiltinParticipantStatelessMessageWriter/ParticipantStatelessMessage
  -> Locator: UDP:239.255.0.1:9150(12) {33}
  -> Locator: UDPv6:ff03::80:9150(13) {33}
  -> Locator: TCP:239.255.0.1:9150-(null)(C)(7) {33}
7/000201-c4 R{35} 0 BuiltinParticipantStatelessMessageReader/ParticipantStatelessMessage
7/ff0003-c2 W{21} 0/1 SEDPbuiltinPublicationsSecureWriter/PublicationBuiltinTopicDataSecure
  No Access control, Normal discovery, Encryption: payload(AES128_HMAC_SHA1)
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0003-c7 {62}
7/ff0003-c7 R{17} 0/0 SEDPbuiltinPublicationsSecureReader/PublicationBuiltinTopicDataSecure
  No Access control, Normal discovery, Encryption: payload(AES128_HMAC_SHA1)
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0003-c2 {63}
7/ff0004-c2 W{23} 0/1 SEDPbuiltinSubscriptionsSecureWriter/SubscriptionBuiltinTopicDataSecure
  No Access control, Normal discovery, Encryption: payload(AES128_HMAC_SHA1)
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0004-c2 {65}
7/ff0004-c7 R{19} 0/0 SEDPbuiltinSubscriptionsSecureReader/SubscriptionBuiltinTopicDataSecure
  No Access control, Normal discovery, Encryption: payload(AES128_HMAC_SHA1)
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0004-c7 {64}
7/ff0200-c2 W{31} 0/0 BuiltinParticipantMessageSecureWriter/ParticipantMessageData
  No Access control, Normal discovery, Encryption: payload(AES128_HMAC_SHA1)
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0200-c7 {56}
7/ff0200-c7 R{29} 0/0 BuiltinParticipantMessageSecureReader/ParticipantMessageData
  No Access control, Normal discovery, Encryption: payload(AES128_HMAC_SHA1)
  -> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0200-c2 {57}
7/ff0202-c3 W{37} 0 BuiltinParticipantVolatileMessageSecureWriter/ParticipantVolatileSecureMessage
```

```

-> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0202-c4 {51}
7/ff0202-c4 R{39} 0 BuiltinParticipantVolatileMessageSecureReader/ParticipantVolatileSecureMessage
-> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0202-c3 {50}
7/ff0401-c2 W{43} 0/1 QeobuiltinPolicyStateWriter/PolicyUpdaterMessageData
No Access control, Normal discovery, Encryption: payload(AES128_HMAC_SHA1)
-> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0401-c7 {52}
7/ff0401-c7 R{41} 0/1 QeobuiltinPolicyStateReader/PolicyUpdaterMessageData
No Access control, Normal discovery, Encryption: payload(AES128_HMAC_SHA1)
-> GUID: 1752a037:2c14ac84:e9ddb9c7-ff0401-c2 {53}

```

As can be seen from this example, there are a lot more topics present when DDS security is used.

History Cache

The history cache contents can be shown with the `scache` or `sdcache` commands as shown below:

```

>sdsc 36
Topic/Type: Square/ShapeType
GUID: 92767f10:010a6e9e:03e80000-000001-7
Max. depth = 1, Flags = Transient-Local/Auto-dispose/Inst-ordered/Keep-last
Limits: * samples, * instances, 1 samples/instance
Hashed: 0, monitor: 0, inform: 1, unacked: 0, ownership: shared
Changes: {2}
  3295.625815887s - [0x2278090*2:ALIVE:0:h=2,R,NN]->0x227469c*1,148 bytes
  3295.723990600s - [0x2278110*2:ALIVE:0:h=1,R,NN]->0x2274440*3,148 bytes
Last handle = 2, Maximum key size = 133
  1: 30727bee:ff213066:4d6cfc73:0a6c72c8, ALIVE, NOT_NEW, D/NW=0/1, W:40 {1}
    Statistics:
      # of cache bytes added: 3876416
      # of cache samples added: 26192
      # of cache registers: 0
      # of cache disposes: 0
      # of cache unregisters: 0
      # of cache notifications: 26192
      # of cache gets: 26192
    Key:
      {color='RED'}
    ->: (1) {
      3295.723990600s - [0x2278110*2:ALIVE:0:h=1,R,NN]->0x2274440*3,148 bytes
    }
  2: 262cfcbbc:42b9f5a5:a1186f2c:5cde78d5, ALIVE, NOT_NEW, D/NW=0/1, I, W:41 {1}
    Statistics:
      # of cache bytes added: 3860432
      # of cache samples added: 26084
      # of cache registers: 0
      # of cache disposes: 0
      # of cache unregisters: 0
      # of cache notifications: 26084
      # of cache gets: 26084
    Key:
      {color='GREEN'}
    ->: (1) {
      3295.625815887s - [0x2278090*2:ALIVE:0:h=2,R,NN]->0x227469c*1,148 bytes
    }

```

In this particular example, where we display the contents of the Square DataReader, samples remain in the history cache, since we always read() from the cache while processing the data instead of doing a take().

The following information is effectively displayed:

- The Topic/Type names.
- The GUID of the DataReader.
- The maximum depth for *KEEP_LAST* caches or the *max_samples_per_instance* limit for *KEEP_ALL* caches.
- Cache flags that are derived from QoS parameters and which control cache behaviour.
- The resource limits for this cache.
- Whether the cache is hashed, actively being monitored.
- The number of as yet unacknowledged samples.
- The type of ownership.
- The global list of changes in the cache, sorted in arrival order (2 samples in this cache).
- If the cache has multiple instances, i.e. is keyed, the last handle and the key size are displayed, as well as the instance list.
- Each active instance is listed as:

1. The instance handle.
2. The hash value of the key as an MD5 checksum if the key is larger than 16 bytes, or the key itself¹³.
3. The last instance and view states.
4. The number of disposed/no_writer state changes.
5. The handle of the writers actively writing data to this cache.
6. Instance statistics if enabled (-DEXTRA_STATS).
7. The key value if the key is hashed.
8. The arrival time of the last sample or
9. The list of samples queued in that instance.



Note: The instance list for a multi-instance cache is either a linear list, if small enough, or a dual skiplist, where the first skiplist has as key the keyhash value, and the second skiplist has as key the instance handle.



Note: Multi-instance caches are automatically converted from one format to the other, when items are added or removed from.

For each sample, the following information is displayed:

- The timestamp for when the sample was generated, relative since the time the program was started, followed by "
- ["
- The address of the Change descriptor and the number of references to it.
- If the change is **ALIVE** or **DISPOSED** or **UNREGISTERED** or **ZOMBIE** (both disposed and unregistered).
- The number of outstanding acknowledgements.
- The writer handle.
- If a writer cache, the sequence number as a *high.low* tuple.
- If the sample was due to a directed write, the list of destinations (after a '=' symbol).
- If a reader cache (in this case), the sample state (NR=Not Read, R=Read) and the view state (NN=Not New, N=New), followed by ']'
- If the sample contains data (ALIVE sample), this is indicated with a "-><addr>*<nrefs>,<len> bytes" string where:
 1. <addr> is the memory address of the sample data.
 2. <nrefs> is the number of references to this sample data.
 3. <len> is the total length of the sample data in bytes.

The cache contents can be queried with the `qcache` command, as follows:

```
>qcache 42 x < 150 and y > 10
1: {color="GREEN", x=104, y=83, shapesize=30}
```

Proxy Endpoints

The `sproxy` command can be used to display RTPS proxy endpoint contexts:

```
>spro 40
1/000003-02 W{40} - Square/ShapeType:
  GUID: 92767f10:01096e8f:03e80000-000002-7 {42} -
  InlineQoS=0, Reliable=1, Active=0, Marshall=0, MCast=1, Tunnel=0
  States (Control/Tx/Ack): READY/ANNOUNCING/WAITING
  Changes: : (1)
             {UNDERWAY:1          [0x2278190*2:ALIVE:h=1,0.37223]}
  LastAck=37126
  Unicast locators:
    UDP:10.0.0.0.29:7679
  Multicast locators:
    UDP:239.255.0.1:7651(2)
  Direct reply locator:
    UDP:10.0.0.0.29:7679
  Last: 0.37222, local=7, remote=0
```

¹³ Secure topics always use an MD5 checksum.

```

# of DATA messages sent:      34576
# of GAPS sent:                0
# of HEARTBEATs sent:         282
# of ACKNACKs received:       37084
# of DATAFRAG messages sent: 0
# of HEARTBEATFRAGs sent:     0
# of NACKFRAGs received:      0
Next GUID: (nil)
GUID: 92767f10:01086e8c:03e80000-000001-7 {43} -
  InlineQoS=0, Reliable=1, Active=0, Marshall=0, MCast=1, Tunnel=0
  States (Control/Tx/Ack): READY/ANNOUNCING/WAITING
  Changes: : (1)
    {UNDERWAY:1                [0x2278190*2:ALIVE:h=1,0.37223]}
  LastAck=74034
  Unicast locators:
    UDP:10.0.0.29:7677
  Multicast locators:
    UDP:239.255.0.1:7651(2)
  Direct reply locator:
    UDP:10.0.0.29:7677
  Last: 0.37222, local=8, remote=0
# of DATA messages sent:      2509
# of GAPS sent:                0
# of HEARTBEATs sent:         37000
# of ACKNACKs received:       37000
# of DATAFRAG messages sent: 0
# of HEARTBEATFRAGs sent:     0
# of NACKFRAGs received:      0
Next GUID: (nil)

```

As shown above, for the local DataWriter with handle 40, there are two proxy Reader contexts.

For each proxy context, the following information is displayed:

- The GUID
- The type of endpoint, e.g. reliable or best-effort.
- The proxy state machine states (see RTPS specification for details).
- The queued changes.
- The last acknowledgement number.
- The unicast locators to use.
- The multicast locators to use.
- The determined locator to use for replies.
- The last sent sequence number.
- The proxy instance ids.
- Optional statistics counters (if enabled with -DEXTRA_STATS).
- The next proxy in the proxy list of a discovered endpoint.

In the example above, there is one sample that has not been acknowledged yet, as indicated with UNDERWAY. Refer to the RTPS specification as to which states can be used for samples in stateful reliable Readers/Writers.

Security data

Introduction

In secure domains, DDS operation becomes a lot more complex. Not only are there additional security-related discovery topics, but there is an extra 3-way handshake needed in order to authenticate remote participants in which user certificates and security policies are verified.

Once successfully authenticated, a shared secret is created which is known only by the two participants that did the handshake.

Using this shared secret, and for each secure topic endpoint, crypto tokens are then transferred between the participants in order to setup encryption (AES128 or AES256) and/or signing (SHA1 or SHA256) parameters for each individual topic Reader/Writer.

Note that there are 3 ways in which security can be used in DDS:

- Full RTPS message encryption/signing.
- Encryption/signing of individual Reader/Writer submessages.
- Payload encryption.

Technicolor DDS only supports payload encryption currently, but all the hooks are present to support the other types in the future should the need arise.

The authentication and access control parameters are handled using an external security plugin module (as recommended by the DDS standard).

The various security parameters can be displayed with the `scrypto`, `spdb` and `spv` commands.

Crypto contexts

In secure domains, crypto contexts are created on various levels and the `scrypto` command can be used on any of these levels to display the parameters that apply there.

Following different levels can be distinguished:

1. Local DomainParticipant.
2. Discovered Participant.
3. Local Writer.
4. Local Reader.
5. Discovered Writer.
6. Discovered Reader.

For each individual level, an example will be shown.

Local DomainParticipant crypto data.

Using the `sdisc` command, we find out what the handle is of the local DomainParticipant. Using the `scrypto` command on that handle, we get the following information:

```
>scrypto 1
Txdata: Crypto not enabled.
Total crypto data: 6052 bytes.
```

This may seem surprising, but it's not, actually. On this level, the crypto data for RTPS message encryption would be stored, and since we don't keep do RTPS message encryption for native DDS security yet, there are no parameters filled in (we do allow transport level encryption which has a similar effect).

The second line gives the total amount of memory that is currently used for all crypto contexts within this domain.

Discovered Participant crypto data.

Here things become a lot more interesting, as can be seen with the output of the `scrypto` command (note that the handle of the discovered participant is shown in the output of the `sdisc` command):

```
>scrypto 49
RxData: Crypto not enabled.
TxSign: Hash not enabled.
RxSign: Hash not enabled.
KxKeys: M.Key: 4c505a48d92134ec2f689c885b39a8b88f975d3863d2e9e13e479df7b979e5af
        M.Mac: a6b6e33266bce3f31e3ab0ecb5dad4d796d8abd06aeb761255a16afec5288615
TToken: AES256/HMAC-SHA256 encrypted/signed data.
        S.Id: 189802
        S.Key: bb1e12f7467e5ed38d7aff8d4b95192d6347921b8c7273c249f30f6ab0f6b72f
        S.Salt: d7e5662b04848642764e9d4bef94ab80
        S.HMAC: c7b471e475aaf15ef05e7371d68618446531e11c66a66f3d05d880d143fa28a4
        S.Count: 94
RToken: AES256/HMAC-SHA256 encrypted/signed data.
        S.Id: 118363
        S.Key: abe9ald764d02322930c93f6de5ebb3f6b87cace0581539990230180a5dd2814
        S.Salt: 4c9ff83fd3e86caa139d0d4fea4a2bbe
        S.HMAC: cfa6e02d637e5d38a4478d5d0d252d0d611e46de36b71e44e703f8036bbbfelc
RxKeys:
[000001-2]:
```



```

Data: AES128/HMAC-SHA1 encrypted/signed data.
  M.Id: 97383,
  M.Key: d1d27c3d5e87c6d9d9f3423817ab60b7
  M.IV: dcef9b829d5608d24af83d87411652e5
  M.HMAC: e27c4cc2ff02b6cafd89466e7c1c329b
  S.Id: 97384
  S.Key: 9b909f210a5155f4ea744394679f72d4
  S.Salt: 41866d7dc02898da570c40535d6af843
  S.HMAC: be9aa60600ffb449cb6dcb04705666d6
Sign: HMAC-SHA1 signed data.
  M.Id: 118382,
  M.Key: e037b50a4296cbfa30b9a37092638d85
  M.HMAC: f13a0137d6c95c547c028e4cf90fc8d5
  S.Id: 0
  S.HMAC: 00000000000000000000000000000000
[000002-7]:
Data: AES128/HMAC-SHA1 encrypted/signed data.
  M.Id: 118361,
  M.Key: e6d43dcc09247eacc47751dfd6310d3b
  M.IV: f3cb6b286ea36c4bf366d32926b72c7b
  M.HMAC: ea1513a4220c52bf6dd01e1c16dab587
  S.Id: 0
  S.Key: 00000000000000000000000000000000
  S.Salt: 00000000000000000000000000000000
  S.HMAC: 00000000000000000000000000000000
Sign: HMAC-SHA1 signed data.
  M.Id: 118380,
  M.Key: d3664ad2d9cf1975029fe4e619add789
  M.HMAC: cedacd0d28277e1ce7a7af0ald8ff7c5
  S.Id: 0
  S.HMAC: 00000000000000000000000000000000
[ff0003-c2]:
Data: AES128/HMAC-SHA1 encrypted/signed data.
  M.Id: 51868,
  M.Key: e76da98037d912251c496ed66adf80eb
  M.IV: f6a4d5194a543095544a1a674019b45
  M.HMAC: dc258b5d781506bb7b27f5631e21709b
  S.Id: 51869
  S.Key: 378bdfd7f5645e6e15c4eb78effff77b
  S.Salt: fc7c7ecd3e6b463f4d4a5ac81488e464
  S.HMAC: deb251e46404984296d3edce0ea7ecc5
Sign: HMAC-SHA1 signed data.
  M.Id: 118372,
  M.Key: d9eeb8eale8f8f2336748fc2cb984349
  M.HMAC: fa6857ca0df43e1e0b3d8465562ed9b9
  S.Id: 0
  S.HMAC: 00000000000000000000000000000000
[ff0003-c7]:
Data: AES128/HMAC-SHA1 encrypted/signed data.
  M.Id: 10802,
  M.Key: c589caa7ae555e687fad7b274f75759f7
  M.IV: e5236a2fc3bd593ff2d401e641703cdb
  M.HMAC: e488a0alb22b4a396e63be54c0397eab
  S.Id: 0
  S.Key: 00000000000000000000000000000000
  S.Salt: 00000000000000000000000000000000
  S.HMAC: 00000000000000000000000000000000
Sign: HMAC-SHA1 signed data.
  M.Id: 118374,
  M.Key: d34e34c96d70e2b9069c665ef2838fc9
  M.HMAC: d2c5b421b73760c2ba4756d357f4cf39
  S.Id: 0
  S.HMAC: 00000000000000000000000000000000
.
.
[ff0401-c7]:
Data: AES128/HMAC-SHA1 encrypted/signed data.
  M.Id: 68097,
  M.Key: fb3a2efba938d4d9d0313e1197d51d21
  M.IV: f3619f99a52014ade4a94fa03d6841b3
  M.HMAC: c5062345ab0fea7beal7a2ba0e7fd6a1
  S.Id: 0
  S.Key: 00000000000000000000000000000000
  S.Salt: 00000000000000000000000000000000
  S.HMAC: 00000000000000000000000000000000
Sign: HMAC-SHA1 signed data.
  M.Id: 118366,
  M.Key: ef0c241b03df678fd21e3254db7134f9

```

```

M.HMAC: f7abd1d25328df498e95c7828979e303
S.Id: 0
S.HMAC: 00000000000000000000000000000000
Keysize: 2000 bytes.

```

The first part of the output gives the crypto parameters that are used on Participant level, the next part (after RxKeys:) lists all the received crypto tokens from the remote participant each one for a specific Entity Id (as shown: [*<entity_id>*] followed by the key data.). The last information shown is the amount of crypto data bytes that are needed for this Participant.

The following data is used on Participant level:

Table 42: Discovered Participant crypto data

Item	Description
RxData	If RTPS message encryption is used: the crypto keys (master/session) for AES decryption of messages arriving from the discovered participant.
TxSign	If RTPS message signing is used: the crypto keys (master/session) for SHA1 or SHA256 signing of messages sent to the discovered participant.
RxSign	If RTPS message signing is used: the crypto keys (master/session) for SHA1 or SHA256 sign verification of messages received from the discovered participant.
KxKeys	The Kx keys are derived from the shared secret that was created during the authentication handshake and are the master keys for encryption and decryption of crypto tokens transferred between the participants. From these parameters, the TToken and RToken crypto session data is derived.
TToken	The session data needed for encryption (AES256) and signing (SHA256) of transmitted crypto tokens to the discovered participant.
RToken	The session data needed for decryption (AES256) and sign verification (SHA256) of received crypto tokens from the discovered participant.

Local Writer crypto data.

The `scrypto` command can be used to display the crypto information of a local DataWriter:

```

>scrypto 46
TxData: AES128/HMAC-SHA1 encrypted/signed data.
  M.Id: 168143,
  M.Key: cbae65f2680f76173fe64c951497597f
  M.IV: e043d43779c20a577dfa5ec96b07a825
  M.HMAC: fdcfda6de6300d0ff6cf996a710fa53d
  S.Id: 168144
  S.Key: 668e456d3ba728948b74aaa332fee5ec
  S.Salt: 86ba9b7fa178c402135d36545409413c
  S.HMAC: ee717f11686e6d845ffd98fb2aa13875
  S.Count: 23

```

This displays the encryption and/or signing crypto data for either payload or submessage encryption of data sent to any discovered DataReader. Both master and session information is displayed.

Note that the information is stored on DataWriter level, since data can be multicasted in DDS. It therefore needs to be encrypted/signed similarly for all allowed DataReaders.

Discovered Reader crypto data.

The `scrypto` command on Discovered Reader level shows a lot more information:

```

>scrypto 66
RxData: AES128/HMAC-SHA1 encrypted/signed data.
  M.Id: 118361,
  M.Key: e6d43dcc09247eacc47751dfd6310d3b
  M.IV: f3cb6b286ea36c4bf366d32926b72c7b

```

```

M.HMAC: ea1513a4220c52bf6dd01e1c16dab587
S.Id: 0
S.Key: 00000000000000000000000000000000
S.Salt: 00000000000000000000000000000000
S.HMAC: 00000000000000000000000000000000
TxSign: HMAC-SHA1 signed data.
M.Id: 189819,
M.Key: e0810e7625af8f0fb9b0b0eac433a483
M.HMAC: e6cf0f8039546efb1b2a0685087550a3
S.Id: 189820
S.HMAC: f77be0e8f8a0241dd9763d2f9fda65d1
RxSign: Hash not enabled.

```

Except for TxSign, most of the information is there for Submessage encryption/signing.

Following information can be distinguished:

Table 43: Discovered Reader crypto data

Item	Description
RxData	If Submessage encryption is used: the crypto keys (master/session) for AES decryption of encrypted submessages coming from the discovered DataReader.
TxSign	If per destination Submessage or Payload signing is used: the crypto keys (master/session) for signing of encrypted data for this specific discovered DataReader.
RxSign	If per destination Submessage or Payload signing is used: the crypto keys (master/session) for verification of signed data from this specific discovered DataReader.

Local Reader crypto data.

The `scripto` command on a local DataReader gives the following information:

```

>scripto 48
TxData: AES128/HMAC-SHA1 encrypted/signed data.
M.Id: 189800,
M.Key: e54445e4bb153e8b34a68f981587e5e5
M.IV: f6dead294778887a00fa6314480f3c9d
M.HMAC: d2772a735195c5aa56a00ffc5848249d
S.Id: 189801
S.Key: 40269e61d35173fa796bc7565d49d63a
S.Salt: 0fdd3ec62593b684b84f1c5229f0a573
S.HMAC: d33f8ed8070876649e89024b17d016cb
S.Count: 0

```

The TxData crypto is used for encrypting/signing Submessages from the DataReader to all the reliable DataWriters.

Discovered Writer crypto data.

The `scripto` command on a discovered DataWriter gives the following info:

```

>scripto 67
RxData: AES128/HMAC-SHA1 encrypted/signed data.
M.Id: 97383,
M.Key: d1d27c3d5e87c6d9d9f3423817ab60b7
M.IV: dcef9b829d5608d24af83d87411652e5
M.HMAC: e27c4cc2ff02b6cafd89466e7c1c329b
S.Id: 97384
S.Key: 9b909f210a5155f4ea744394679f72d4
S.Salt: 41866d7dc02898da570c40535d6af843
S.HMAC: be9aa60600ffb449cb6dcb04705666d6
TxSign: HMAC-SHA1 signed data.
M.Id: 189821,
M.Key: cb7e8b682955783817d85069dc2f0eeb
M.HMAC: d75953c43b7b1bf0e864a45139e27b5b
S.Id: 189822
S.HMAC: 341f9edc4a503a9b72d3519af516391d
RxSign: Hash not enabled.

```

Again a lot of information:

Table 44: Discovered Writer crypto data

Item	Description
RxData	If submessage or payload encryption is used: the crypto keys (master/session) for AES decryption of encrypted payload data coming from the discovered DataWriter.
TxSign	If Submessages need to be signed, the crypto keys (master/session) for signing of encrypted data for this specific discovered DataWriter.
RxSign	If Payload or submessages are signed individually for each DataReader, the crypto keys for sign verification.

DTLS contexts dump command.

The DTLS context dump command can be used when DTLS/UDP based transport security is used with peer-to-peer DTLS/UDP connections.

Since this quickly leads to scaling issues if more than 5 or so peers are used, this is no longer used, especially since native DDS security performs a lot better in this case.

```
>dtls
DTLS connection attempts:
    <No connections>
DTLS v4 Servers:
    DTLS.S(14) UDP:10.0.0.29:4635-DTLS(S)(14) [18] p:0x8e31b0
    <No connections>
DTLS v6 Servers:
    DTLS.S(16) UDPv6:2a02:a03f:669:3100:92e6:baff:fe10:7817:4635-DTLS(S)(16) [20] p:0x8eb880
    <No connections>
    DTLS.S(15) UDPv6:2a02:a03f:669:3100:1c3c:6a97:d469:750:4635-DTLS(S)(15) [19] p:0x8eb780
    <No connections>
```

Security Policy database.

The current security policy database can be queried with the `spdb` command (which stands for "show policy database") can be used:

```
>spdb
Security policy database rules:
5) Participant [2]: '7fef45661ffa779f 9050 21a5' | access: Unclassified | blacklist: 0 | clone from 3
4) Participant [1]: 'DCPS Test program' | access: Unclassified | blacklist: 0 | clone from 1
1) Participant: 'DCPS Test program' | access: Secret | blacklist: 0
|----> Topic: HelloWorld | mode: * | controlled: 1 | secure_disc: 1 | encrypt: payload(AES128_HMAC_SHA1) | domId: 7 | blacklist: 0
|----> Partition: * | mode: * | domId: * | blacklist: 0
2) Participant: '7fef45661ffa779f 9283 21a6' | access: Secret | blacklist: 0
|----> Topic: HelloWorld | mode: * | controlled: 0 | secure_disc: 0 | encrypt: none | domId: 7 | blacklist: 0
|----> Partition: * | mode: * | domId: * | blacklist: 0
3) Participant: '7fef45661ffa779f 9050 21a5' | access: Secret | blacklist: 0
|----> Topic: HelloWorld | mode: * | controlled: 0 | secure_disc: 0 | encrypt: none | domId: 7 | blacklist: 0
|----> Partition: * | mode: * | domId: * | blacklist: 0
1) Domain 7: access: Secret | exclusive: 1 | controlled: 1 | encrypt: none | secure transport: Local=any/Remote=any | blacklist: 0
|----> Topic: * | mode: none | controlled: 0 | secure_disc: 1 | encrypt: payload(HMAC_SHA256) | domId: 7 | blacklist: 0
|----> Partition: * | mode: * | domId: 7 | blacklist: 0
2) Domain 2: access: Top-Secret | exclusive: 1 | controlled: 1 | encrypt: message(AES256_HMAC_SHA256) | secure transport:
Local=any/Remote=any | blacklist: 0
|----> Topic: * | mode: none | controlled: 0 | secure_disc: 1 | encrypt: submessage(AES128_HMAC_SHA1) | domId: 2 | blacklist: 0
|----> Partition: * | mode: * | domId: 2 | blacklist: 0
3) Domain *: access: Unclassified | exclusive: 0 | controlled: 0 | encrypt: none | secure transport: Local=no/Remote=no |
blacklist: 0
|----> Topic: * | mode: * | controlled: 0 | secure_disc: 0 | encrypt: none | domId: * | blacklist: 0
|----> Partition: * | mode: * | domId: * | blacklist: 0
```

The information shown here is actually identical to the content of the `security.xml` file shown before.

RTPS Send/Receive contexts

The RTPS Send and Receive contexts can be displayed with the `stx` and `srx` commands respectively.

Following examples should demonstrate this:

```
>stx
# of locator transmits: 30761
# of missing locator events: 0
# of out-of-memory conditions: 0
```

```

>srx
Protocol version: {2, 1}
Src. Vendor Id: {0x1, 0xe}
Src. GUID Prefix: c0a84080:00062810:03e80000
Dst. GUID Prefix: c0a84080:000d3c04:03e80000
Domain: 0x9644b00
# of invalid submessages: 0
# of too short submessages: 0
# of invalid QoS errors: 0
# of out-of-memory errors: 0
# of unknown destination errors: 15108
# of invalid marshalling data: 0
Last error: Unknown destination
Last MEP:
    00 00 00 00 d8 0e 64 09 - d0 0e 64 09 bc 00 04 00      .....d...d....
    15 07 bc 00 00 00 02 07 - 00 00 03 02 00 00 00 00      .....
    9e 1d 00 00 00 00 00 00 - 9e 1d 00 00 6e 1e 00 00      .....n....
    42 b9 f5 a5 a1 18 6f 2c - 5c de 78 d5 01 00 00 00      B.....O,\.x....
    00 00 00 00 e8 10 64 09 - 00 00 00 00 90 00 0a 00      .....d.....
    00 00 00 00 00 00 00 00 - c8 37 64 09 00 00 00 00      .....7d.....
    00 00 00 00                                ....
Last domain: 0
Last msg:
    00 00 10 00 00 00 00 00 - 00 00 03 02 00 00 00 00      .....
    9e 1d 00 00                                ....

```

Memory dumps

A number of commands are available for dumping memory regions in a variety of formats.

The following commands can be used for this:

- `d` or `dm` display memory in mixed hexadecimal and ASCII characters.
- `db` displays memory as a number of bytes in rows of 16 bytes.
- `ds` displays memory as a number of short words (16-bits) in rows of 8 words.
- `dl` displays memory as a number of long words (32-bits) in rows of 4 long words.

If no arguments are given with a command, the dump continues from the memory address that follows the previous dump command, and dumps 64 bytes, e.g. 4 rows by default.



Note: Make sure that the first dump command that is given specifies a valid address, since the default memory location is set to 0 if no dump command was specified yet!

If an address is specified, dumping starts from that location.

If a number follows the address, it indicates the length (in bytes) of the memory region that needs to be dumped, otherwise the default 64-byte range is dumped.



Note: Be very careful with these commands. If an inappropriate address is specified, segfaults might occur!!

Control commands

The DDS Debug Shell provides a number of commands that can be used to control application programs directly:

- The `pause`, `resume`, `delay`, `menu` and `quit` commands can be used, for example, to control the `dcps` test program and the `shapes` demo program. This is described in the chapters on these programs. Other application programs can register themselves in the Shell to enable some of these commands selectively (see `dds/dds_debug.h`).
- The `server` command can be used to start the DDS Debug server in order to provide remote access to the shell via the `telnet` protocol.
- Remotely connected `telnet` sessions should be closed via the `exit` command.
- The `ase` and `asp` commands can be used to assert Topic and Participant Liveliness, but should only be used in the context of the `dcps` test program.

- RTPS proxy contexts can be reset with the `rproxy` command.
- Security rehandshaking can be forced with the `rehs` command.
- The suspend state can be entered and exited with `suspend` and `activate` commands.
- Configuration parameters can be added, changed and removed with the `set` and `unset` commands.



Note: A number of these commands (`ase`, `asp`, `rproxy`, `rehs`, `suspend` and `activate`, `set` and `unset`) are very dangerous and should not be used on deployed systems (for obvious reasons) unless you know exactly what you are doing. They are mainly present for testing purposes.

Tracing and Profiling commands

Various classes of tracing and profiling are supported in Technicolor DDS:

- RTPS Messages, Finite State Machine state changes, Control and Timer tracing can be used while debugging (DDS Debugging must be enabled for this) in order to verify correct operation of selective endpoints. The `dtrace` and `trace <endpoint>` commands are provided for this.
- Cyclic tracing can optionally be enabled via an additional compilation option (`-DCTRACE_USED`) to provide for DDS-wide event tracing via the `cton`, `ctoff`, `ctclr`, `ctmode`, `ctinfo`, `ctdump` and `ctsave` commands.
- Lock tracing can optionally be enabled via another compilation option (`-DLOCK_TRACE`) to provide for tracing of locks, in order to find out the cause of some deadlocks. The `slstat` can then be used to see which locks were taken, by whom and who's waiting for these locks.
- Profiling can be enabled on selective operations within the DDS Middleware via another conditional compilation option (`-DPROFILE`). The commands to control profiling are `profs` and `cprofs`.

We will not go into more detail here on these features, since they are quite advanced and require some deep insights into the working of Technicolor DDS, and are not intended for casual users.